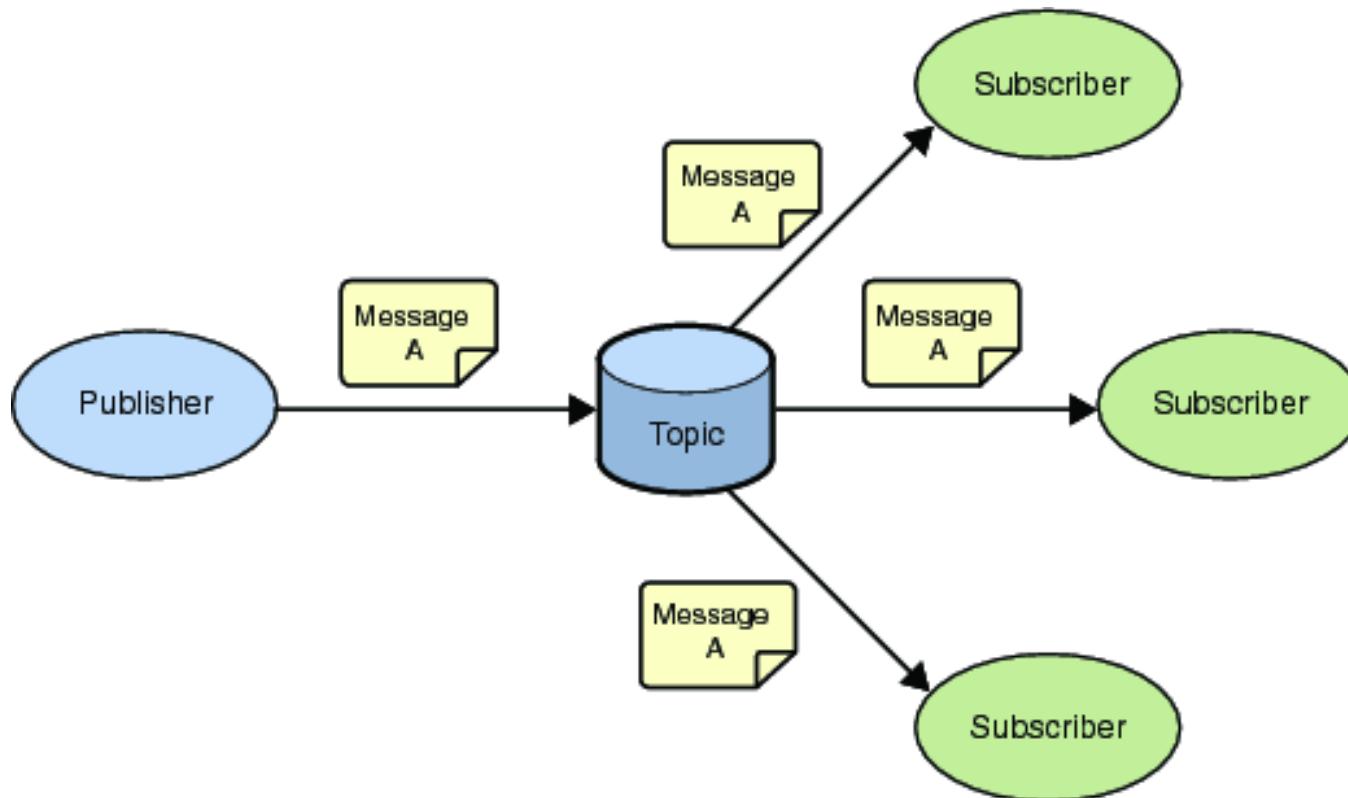


# Publish/Subscribe Systems



# Agenda

- Publish/Subscribe
  - Definition
  - Decoupling
  - Applications
- Matching
  - Topic-based Model
  - Content-based Model
- Routing
  - Rendezvous-based
  - Overlay-based routing

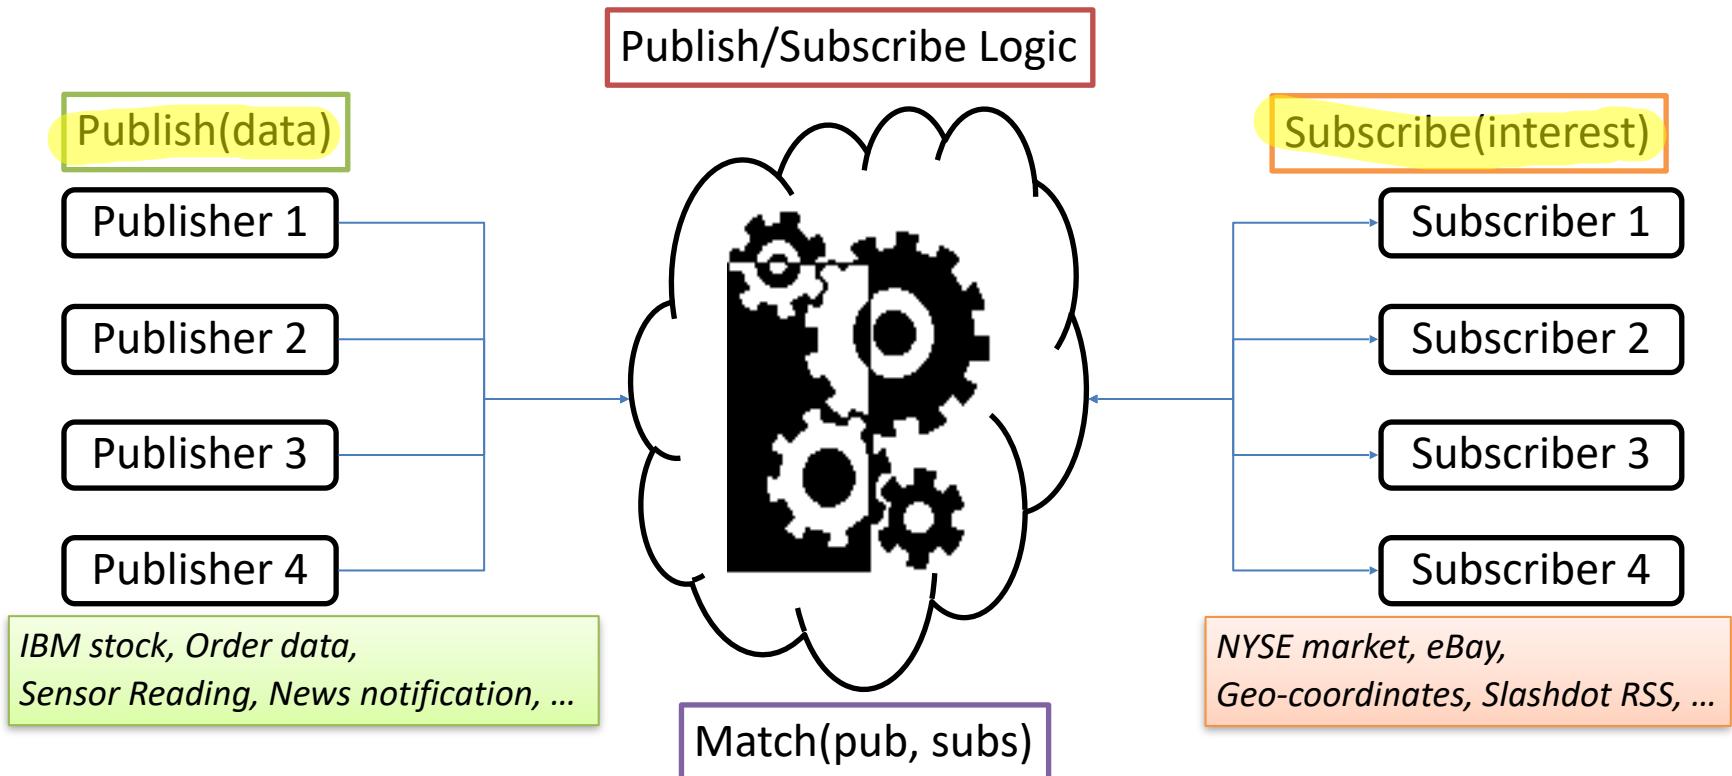


Pixabay.com

# PUBLISH/SUBSCRIBE

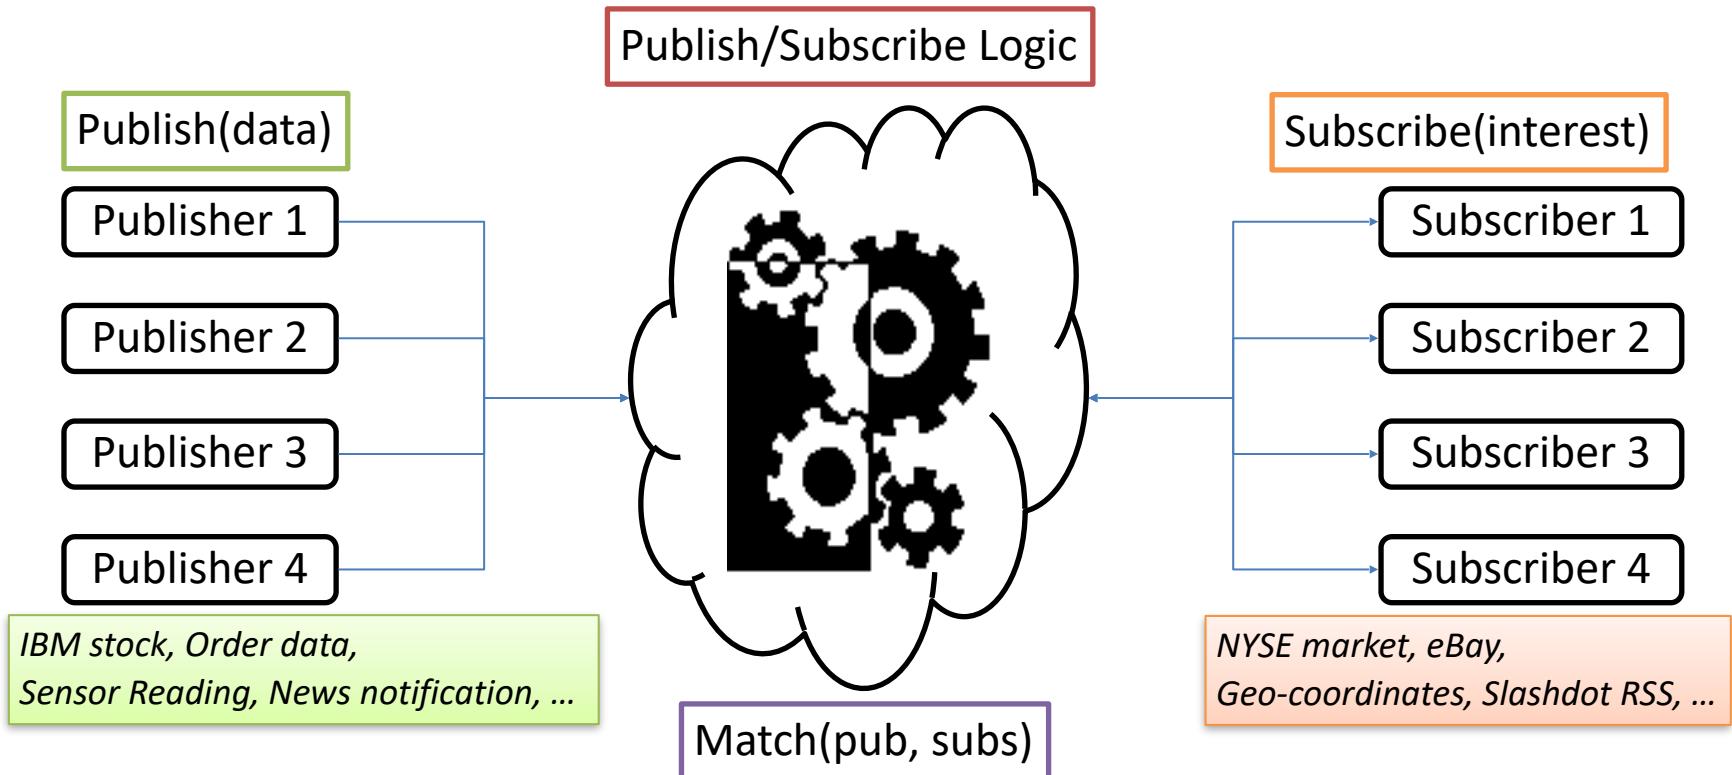
## DEFINITION AND MODEL

# Definitions



Publish/Subscribe is a **messaging pattern** which allows **for many-to-many** communication: Data sources (publishers) publish **publications**, data sinks (subscribers) subscribe via **subscriptions** and receive **notifications**, i.e., publications of interest (publications matching registered subscriptions.)

# Definitions



Publish/Subscribe is a **messaging pattern** which allows for **many-to-many** communication: Data sources (publishers) publish **publications**, data sinks (subscribers) subscribe via **subscriptions** and receive **notifications**, i.e., publications of interest (publications matching registered subscriptions.)

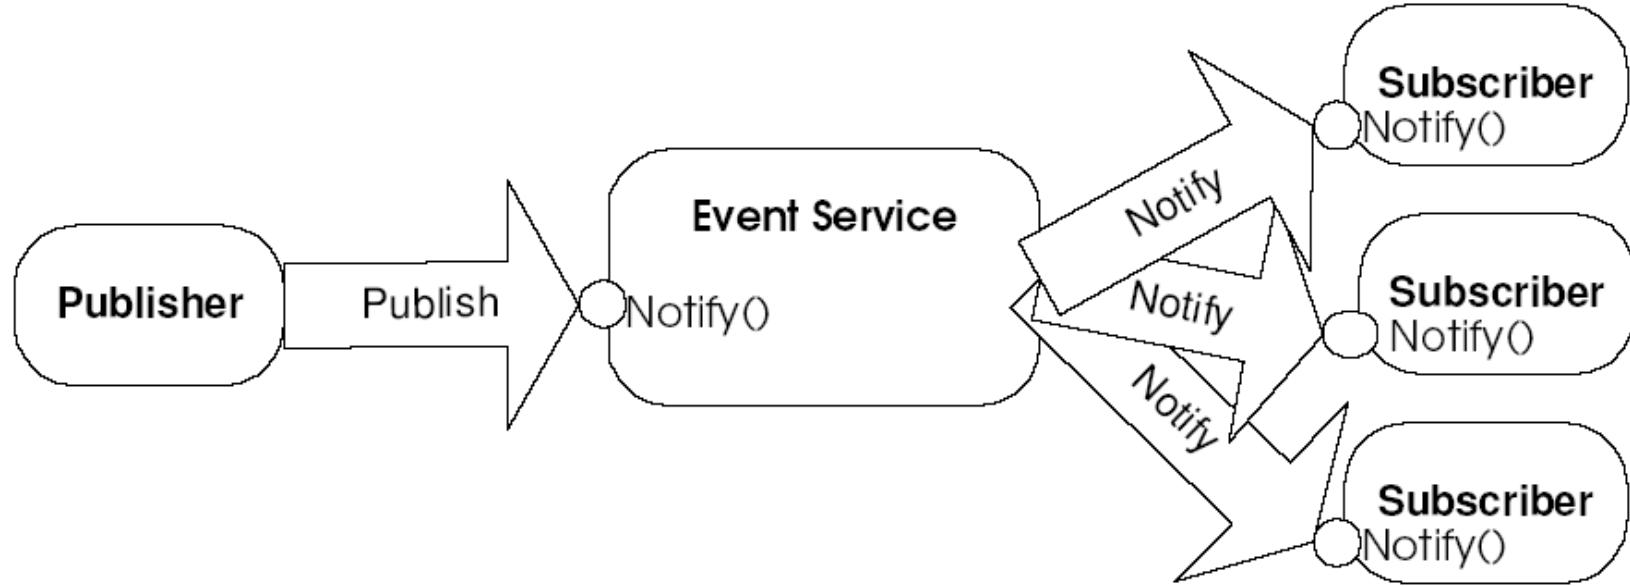
# Pub/Sub API

- Publish(data, metadata) – publication metadata depends on the language (cf. Matching model)
- Subscribe(interest) – subscription interest is defined depending on the language (cf. Matching model)
- Unsubscribe(interest) – remove a subscription
- (Un-)Advertise(data type/domain) – publisher advertises type of content it will publish

# Pub/Sub System Design Dimensions

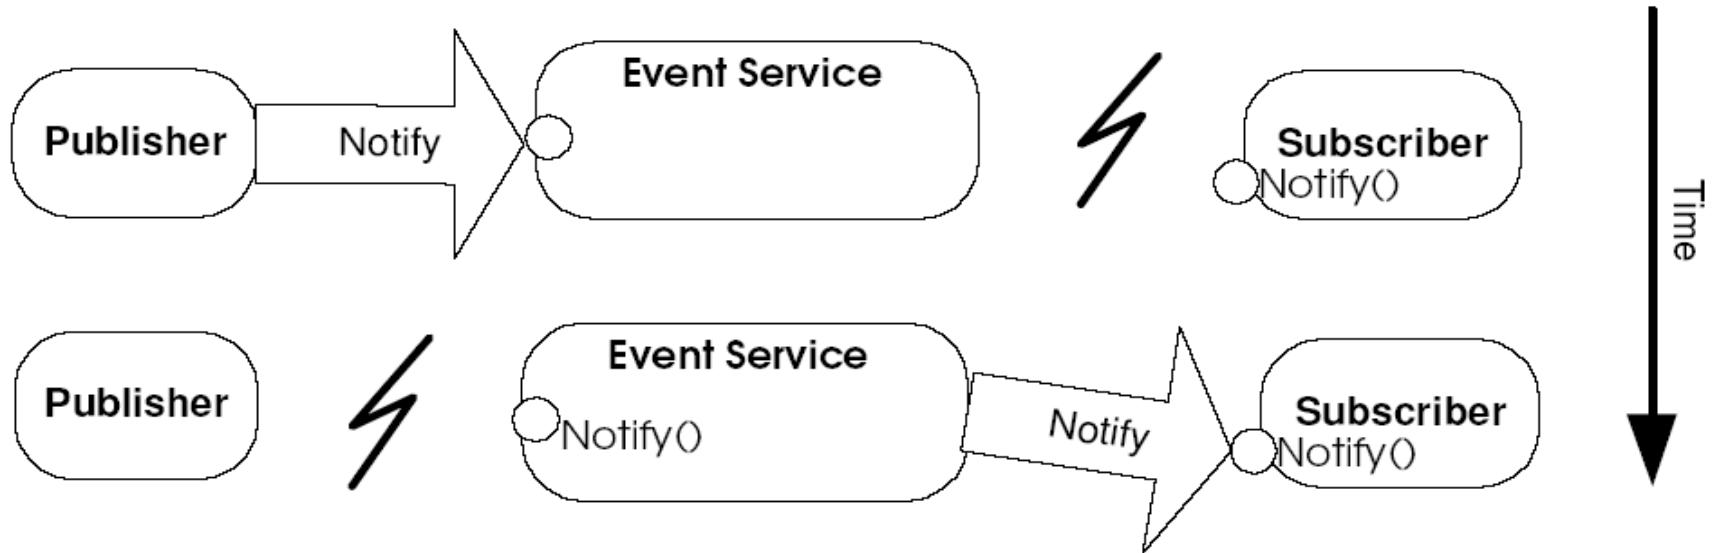
- Matching-based communication
  - Not host-to-host communication
  - Communication is based on the **matching model**
- Optimized for **scalability** and **performance**
  - Large number of publishers and subscribers
  - High rate of publications
  - Fast matching: usually **stateless** (does not consider sequences or composite patterns)
  - Simple matching: commonly limited to **topics**
- **Decoupling properties** (most popular):
  - Time
  - Space
  - Synchronization

# Space Decoupling



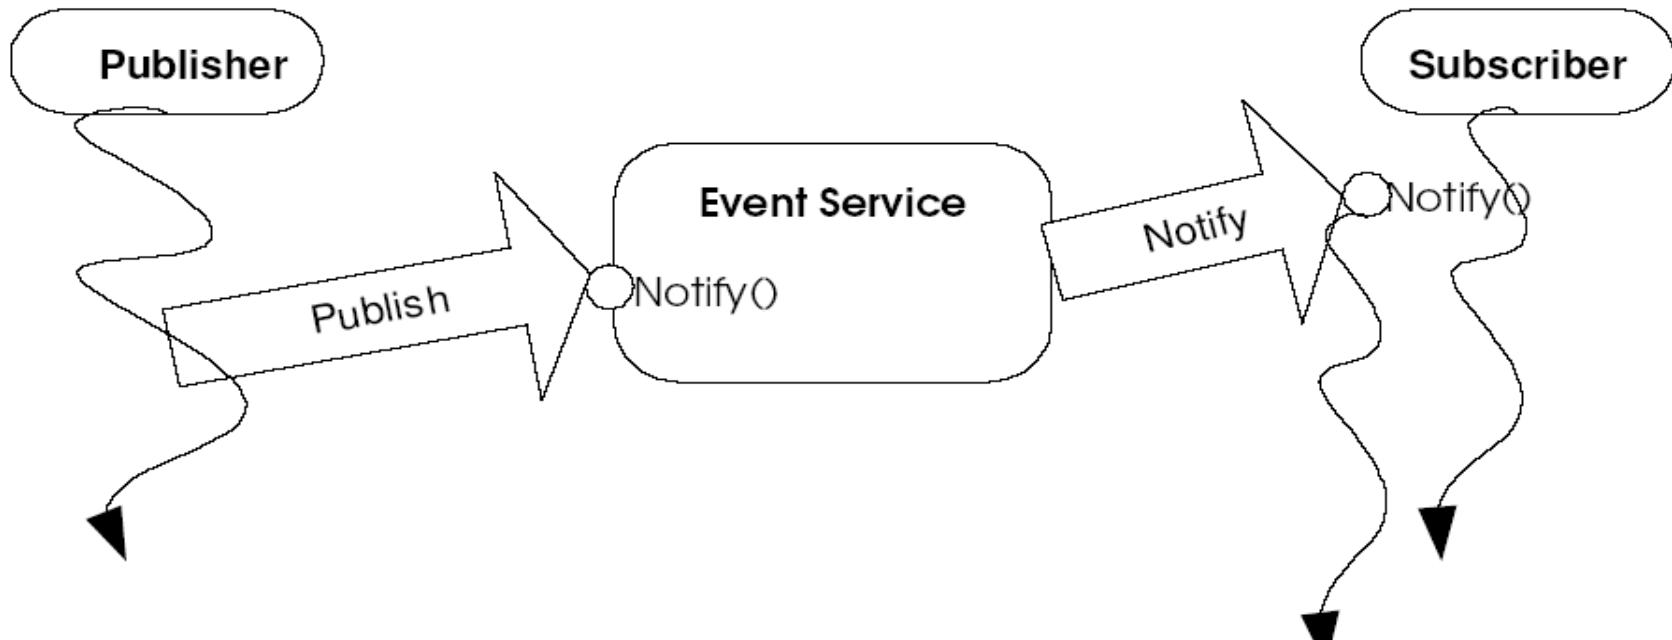
- No need for clients (publishers & subscribers) to hold references to or even know about each other
- Clients may be physically distributed

# Time Decoupling



- No need for clients to be available at the same time

# Synchronization Decoupling



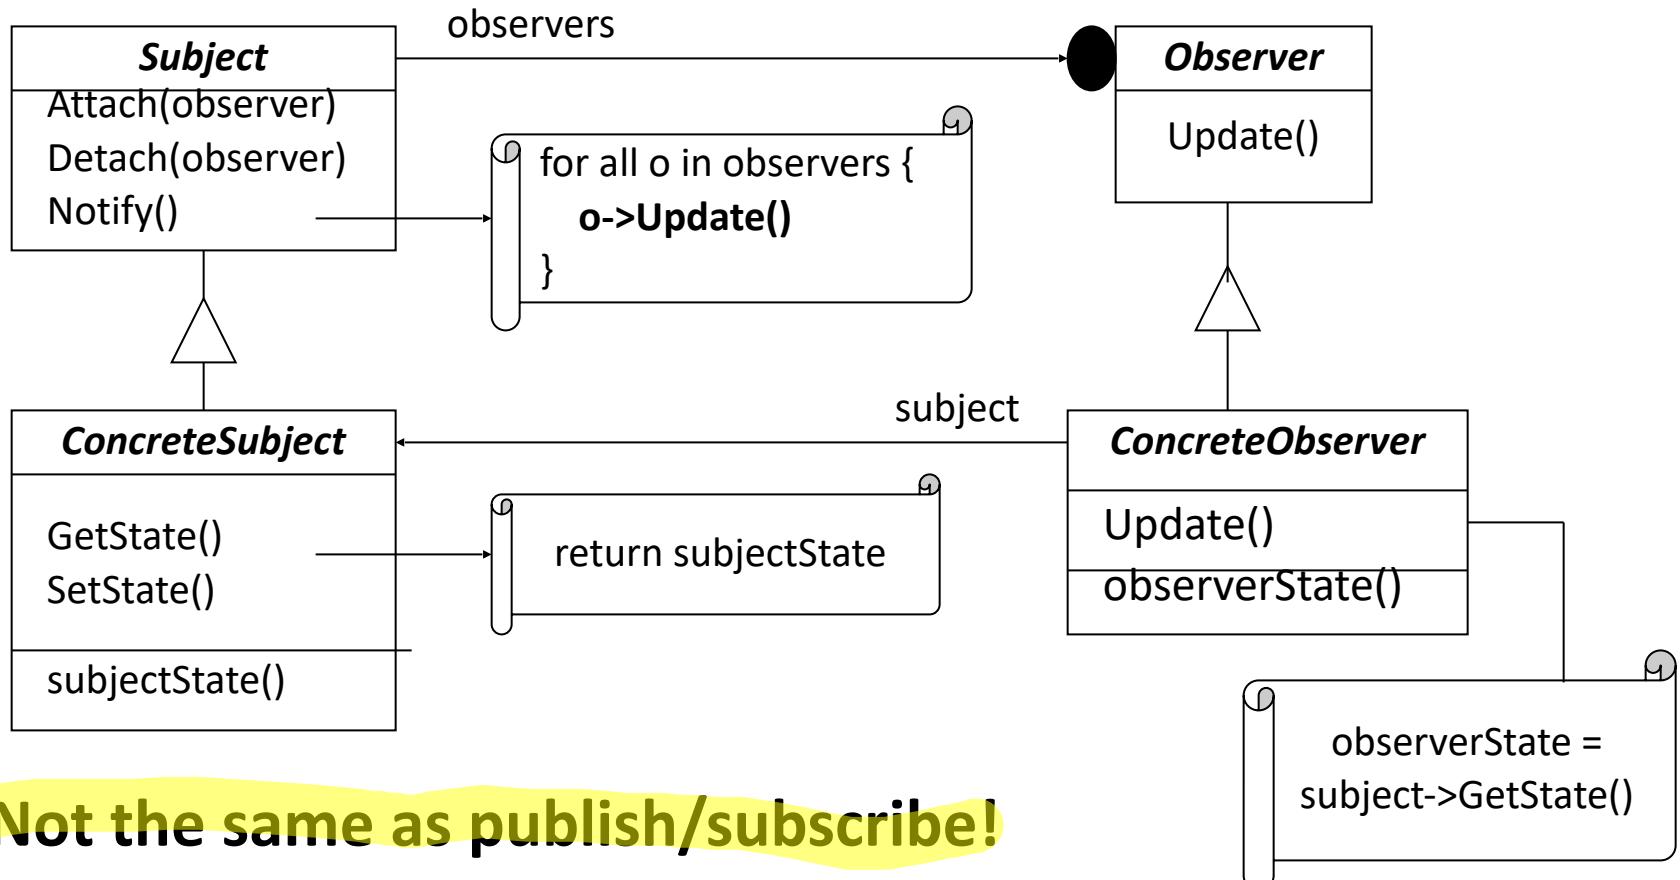
- Control flow is not blocked by the interaction

# *Why Decouple?*

- Decouple publishing and subscribing end-points
  - Removes explicit dependencies
  - Reduces coordination
  - Reduces synchronization
- Creates highly dynamic systems
  - Data sources and sinks come and go
  - Supports continuous operation
- Enables event-based systems, increases scalability of system
- Makes debugging system more difficult

# Pub/Sub vs. Observer Design Pattern

## a.k.a. Listener



- **Not the same as publish/subscribe!**
- Publisher (subject) knows subscribers (observer) and vice versa (tightly coupled)

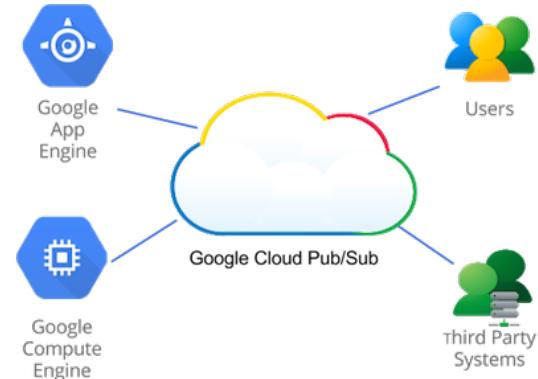
# Applications



*News Syndication*



*Stock Tracker*



*App Notifications*



*Network Engine for Games*

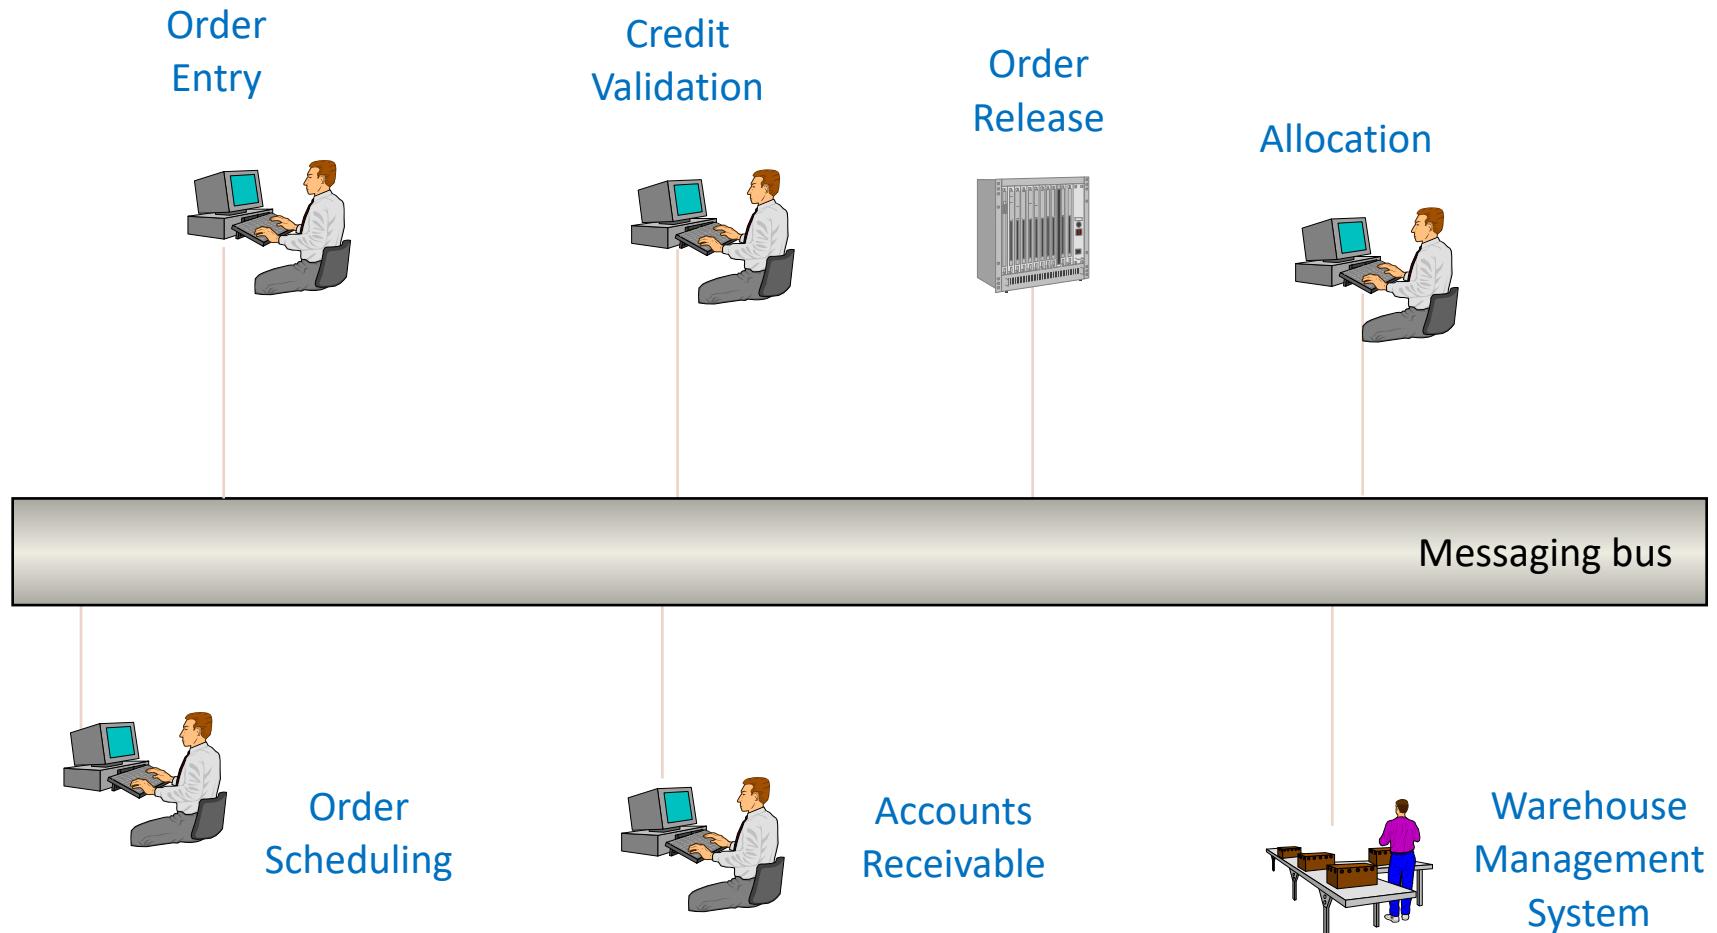


*Social Networks*

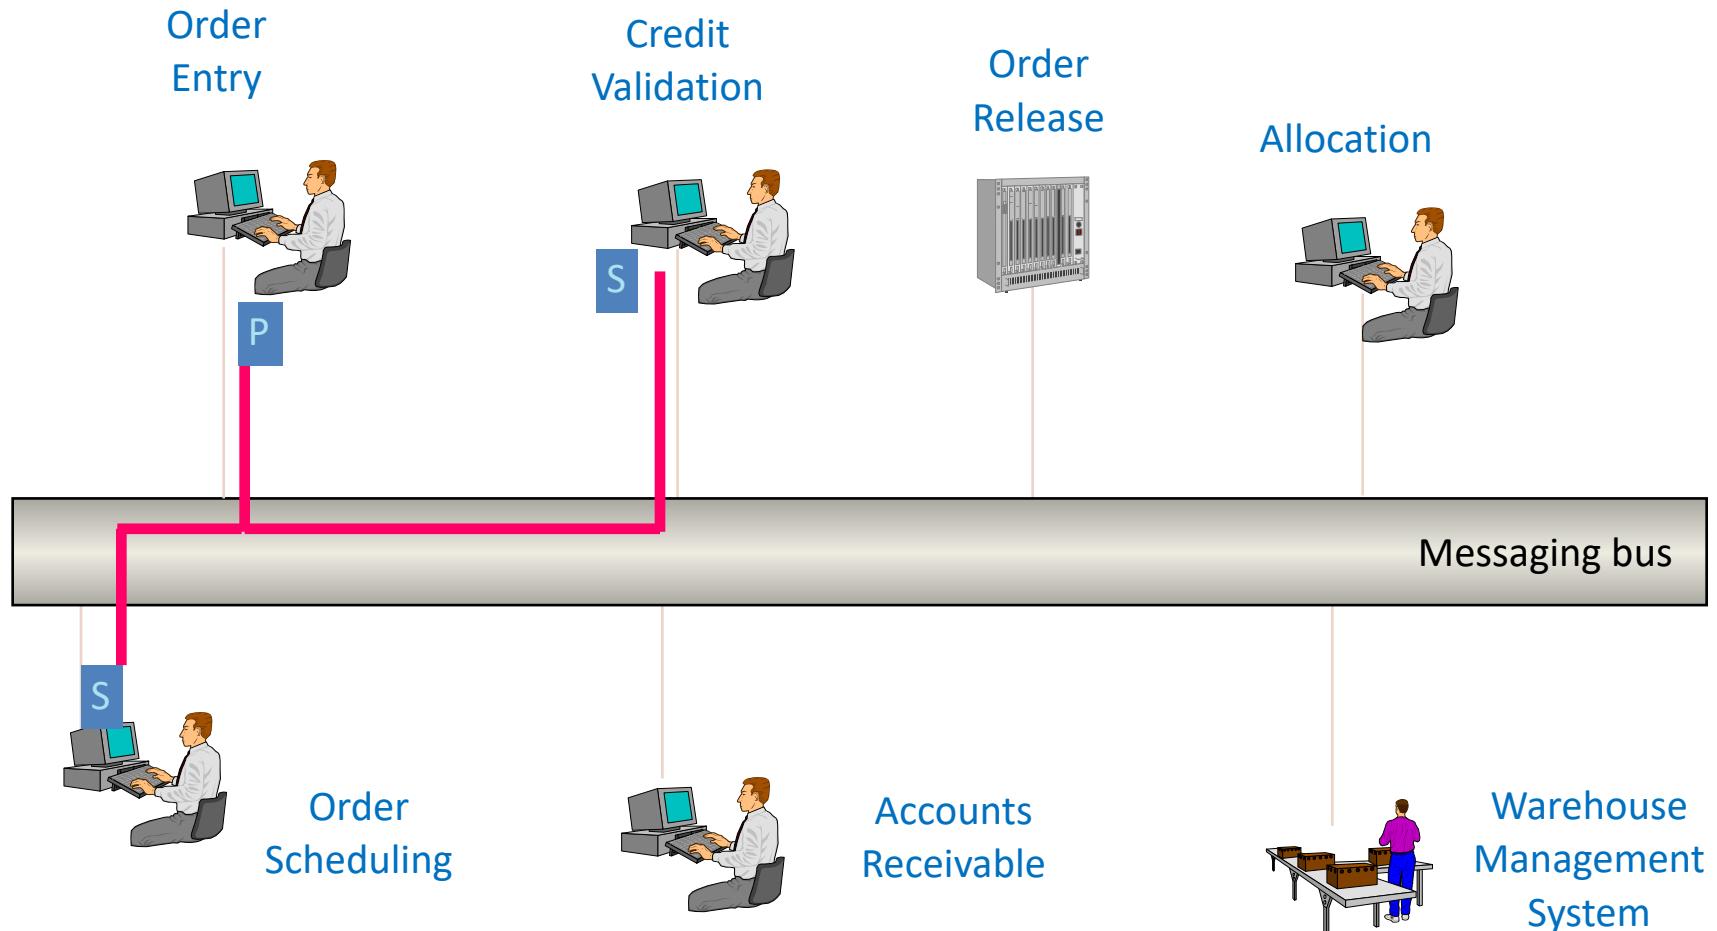


*Traffic Monitoring*

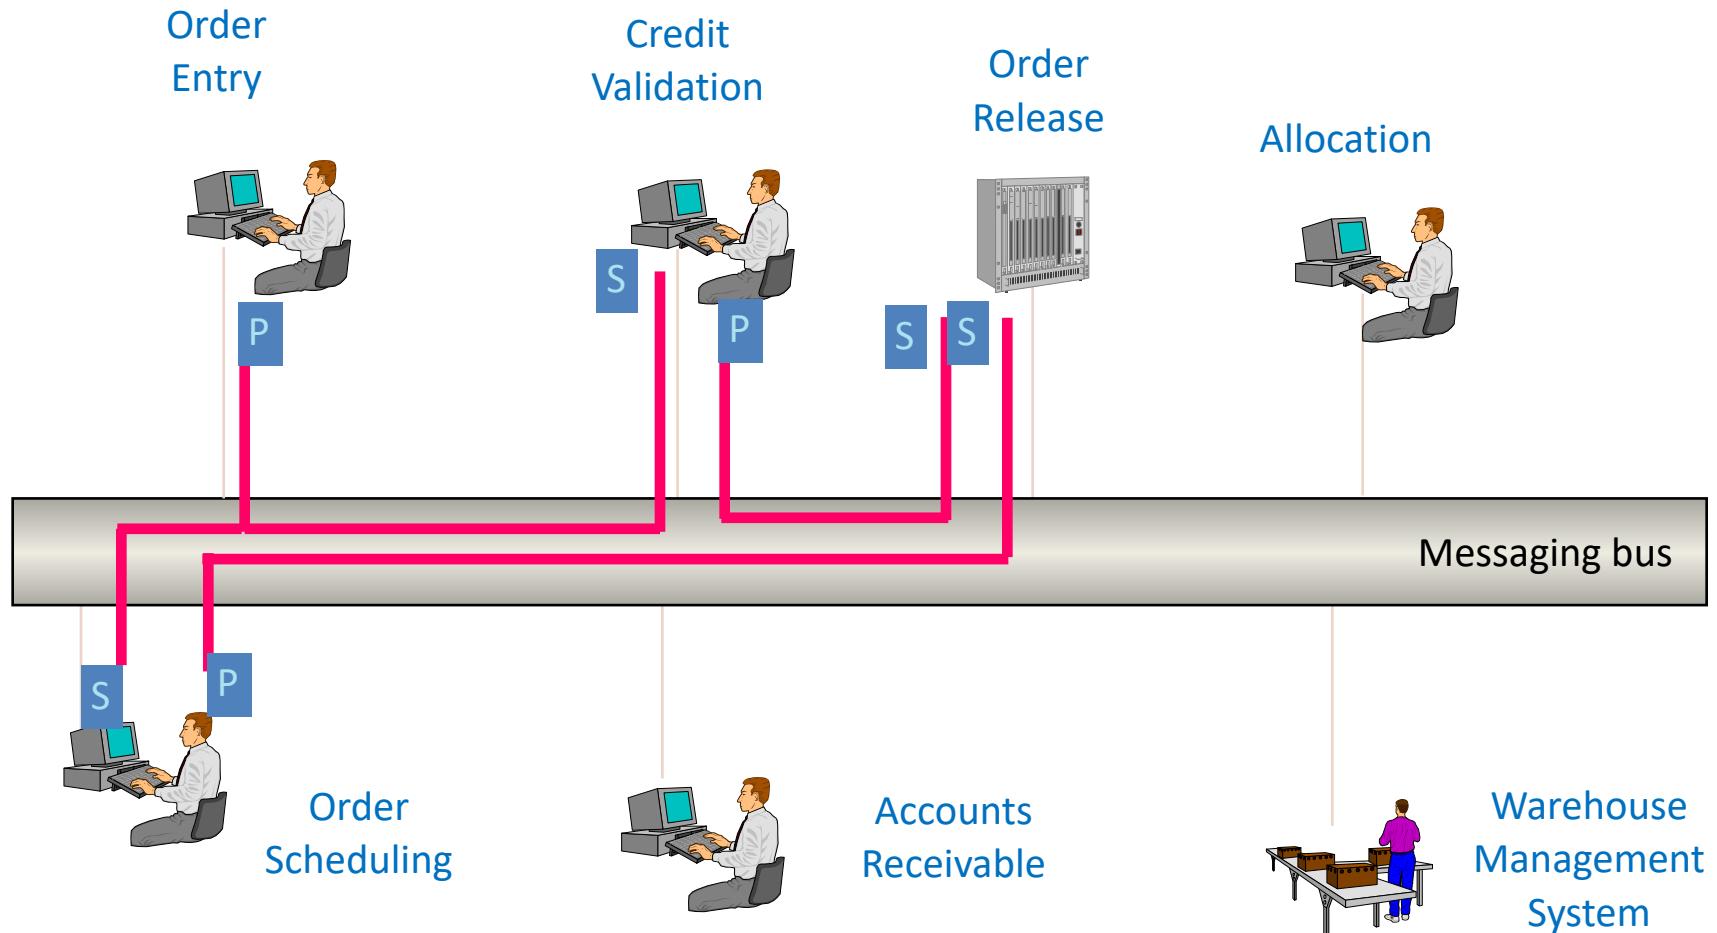
# Enterprise Application Integration



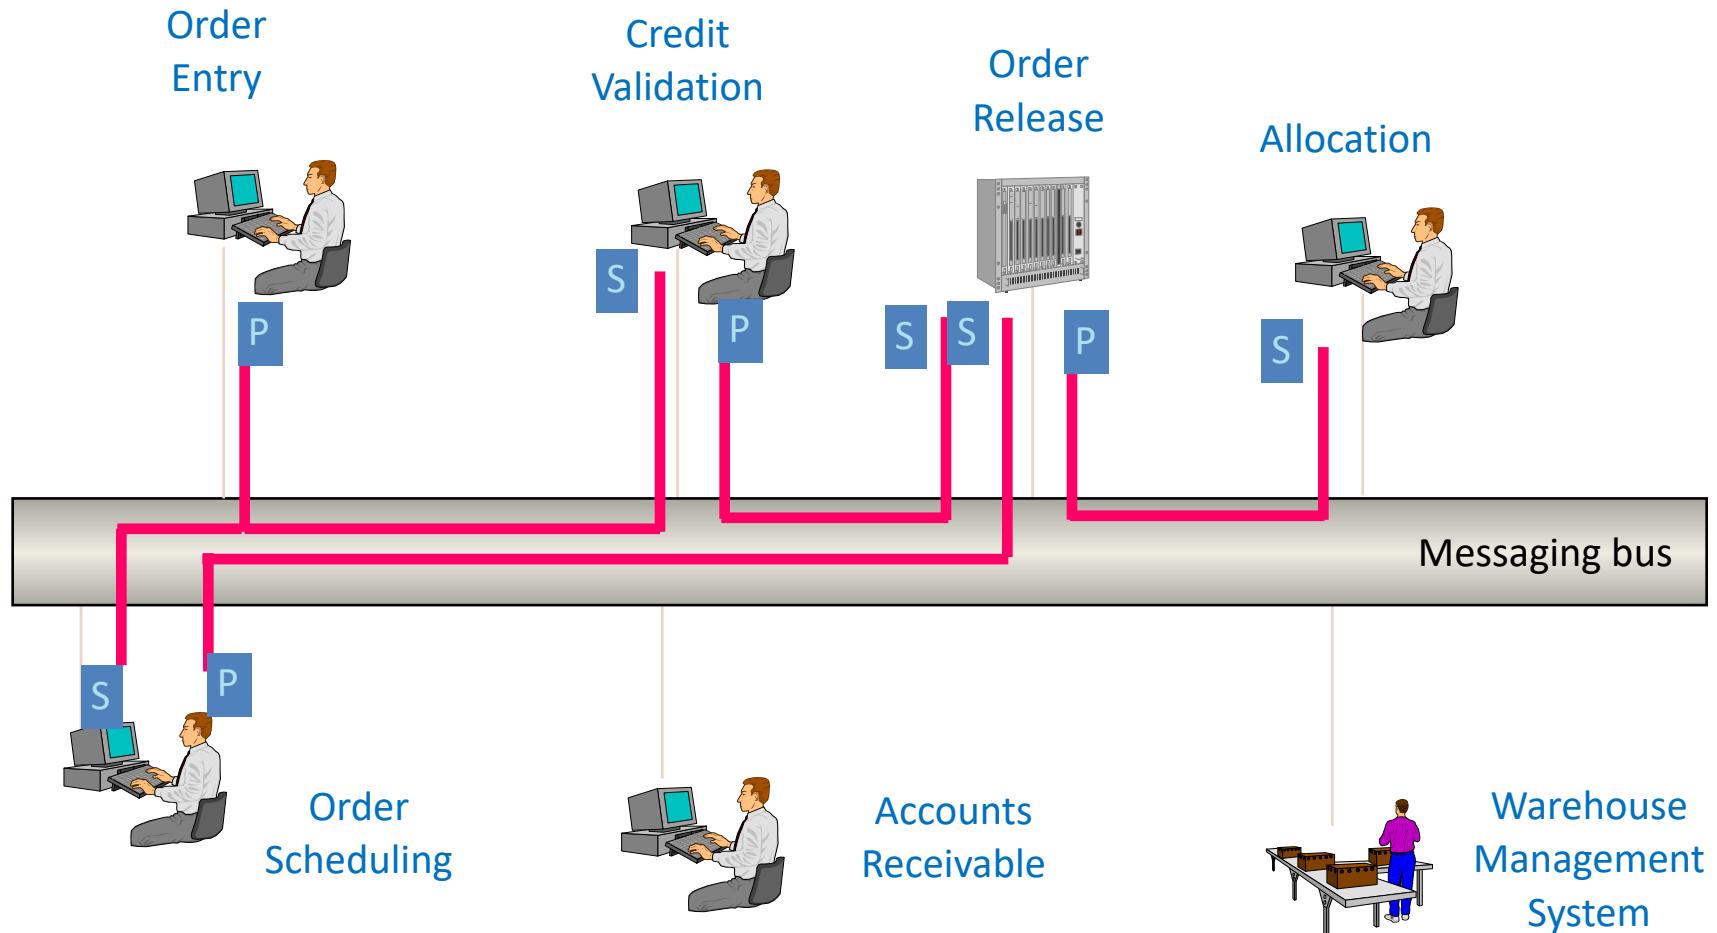
# Enterprise Application Integration



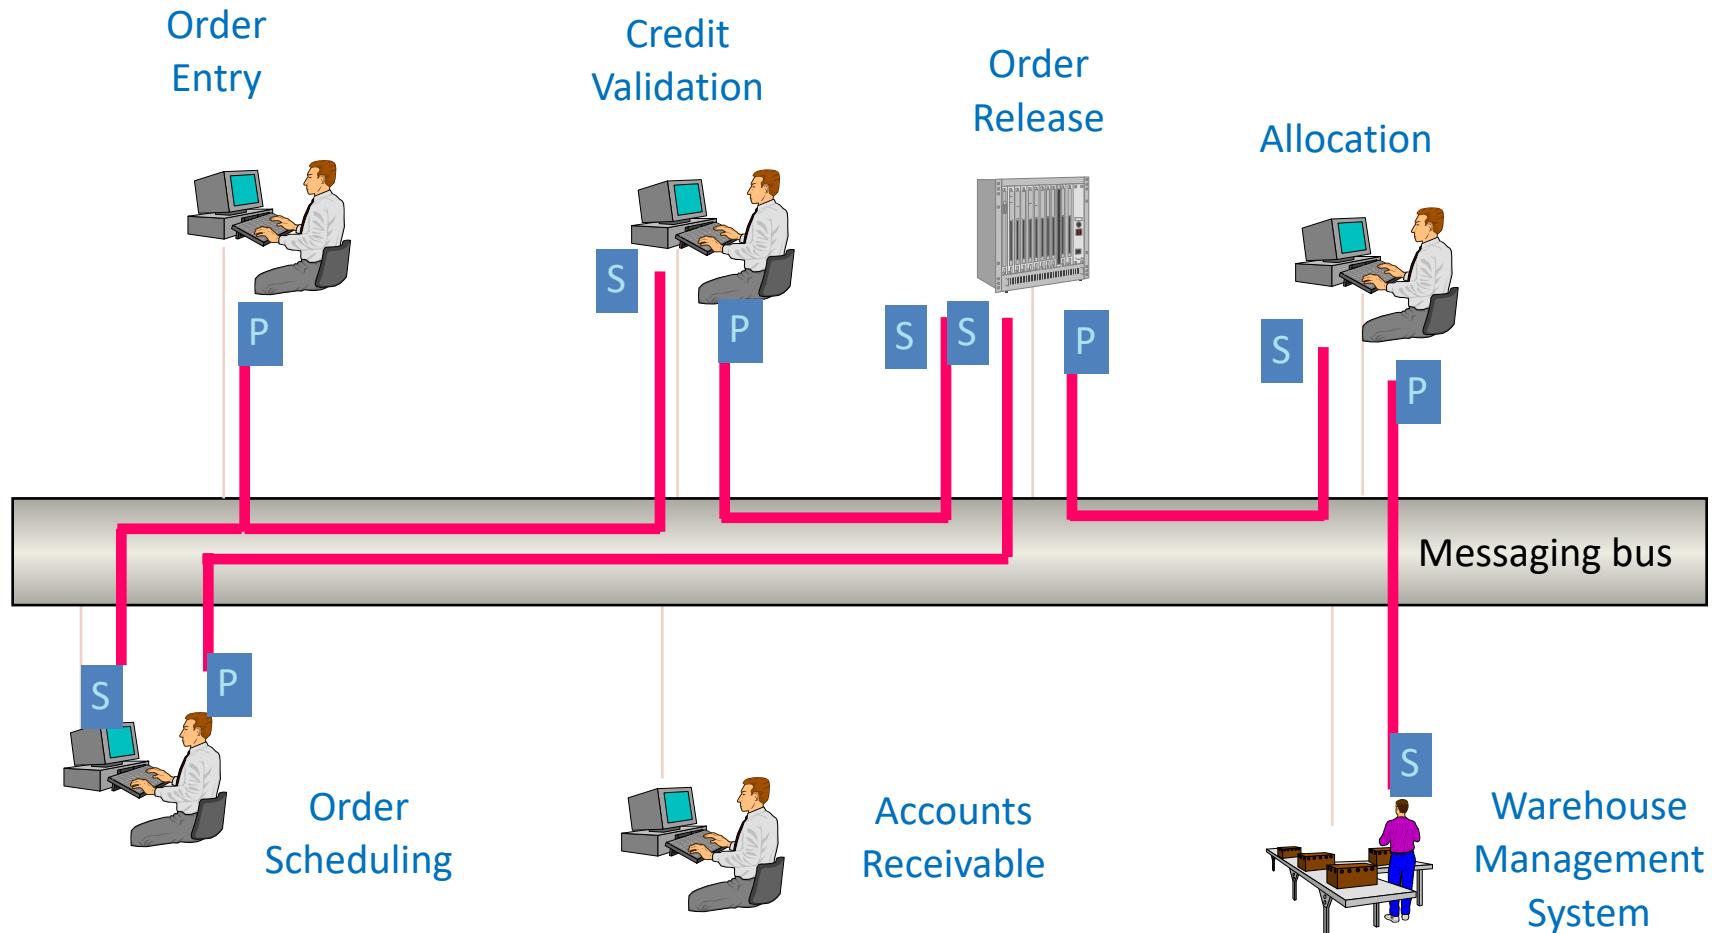
# Enterprise Application Integration



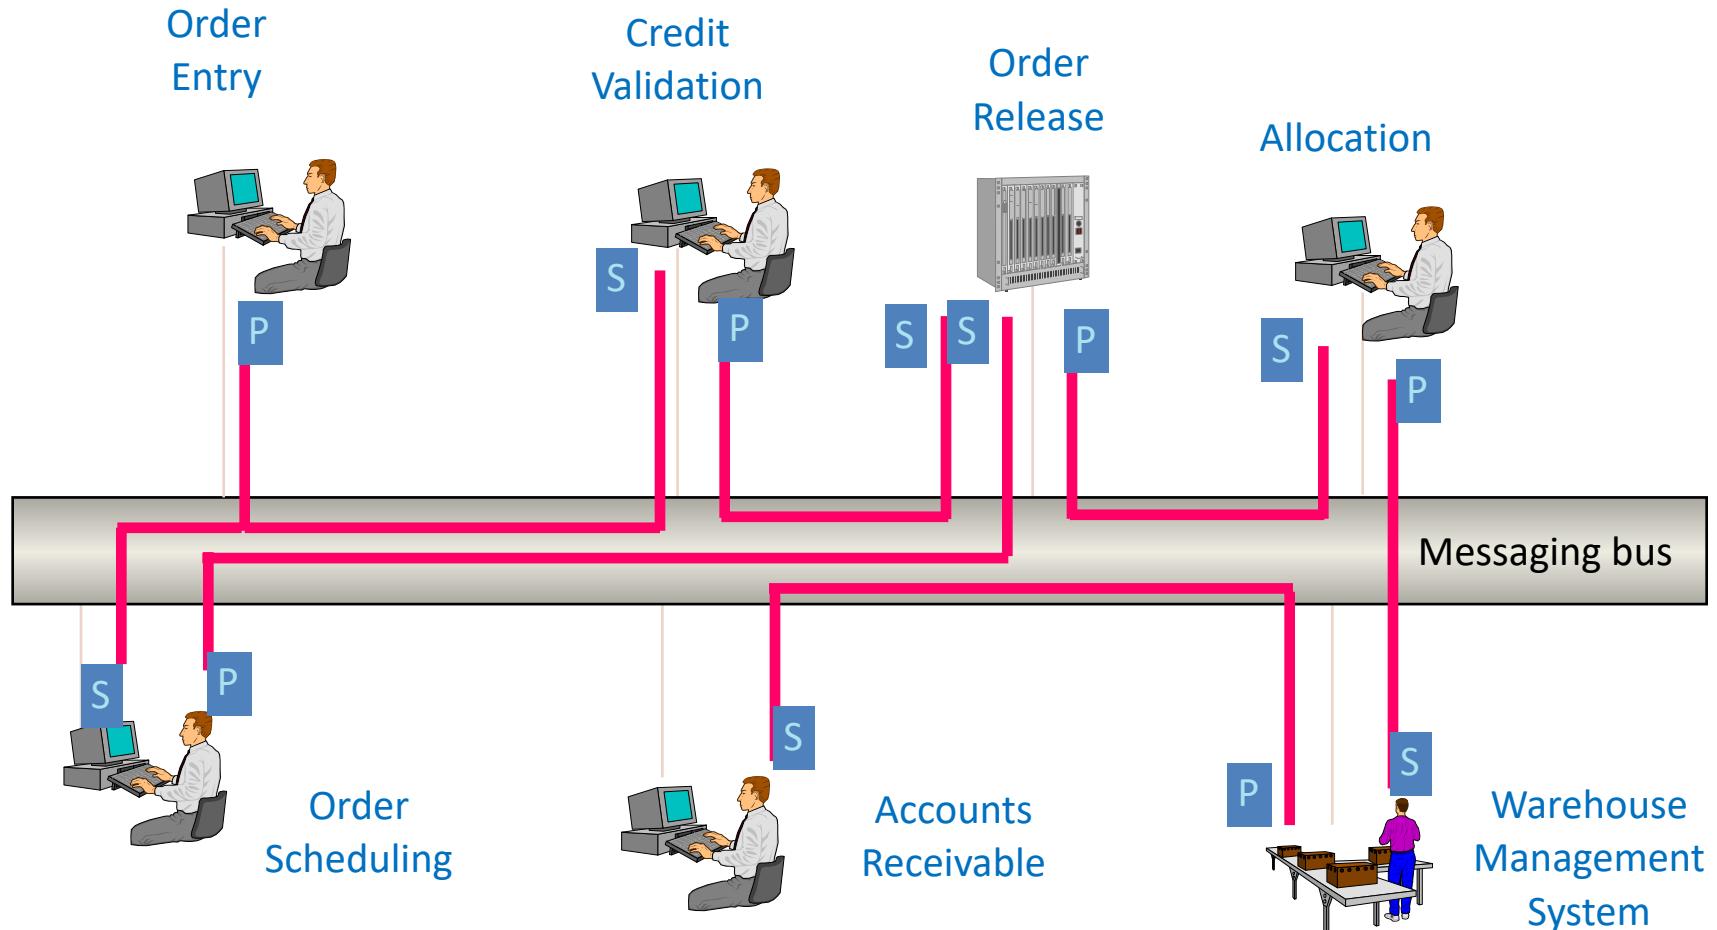
# Enterprise Application Integration



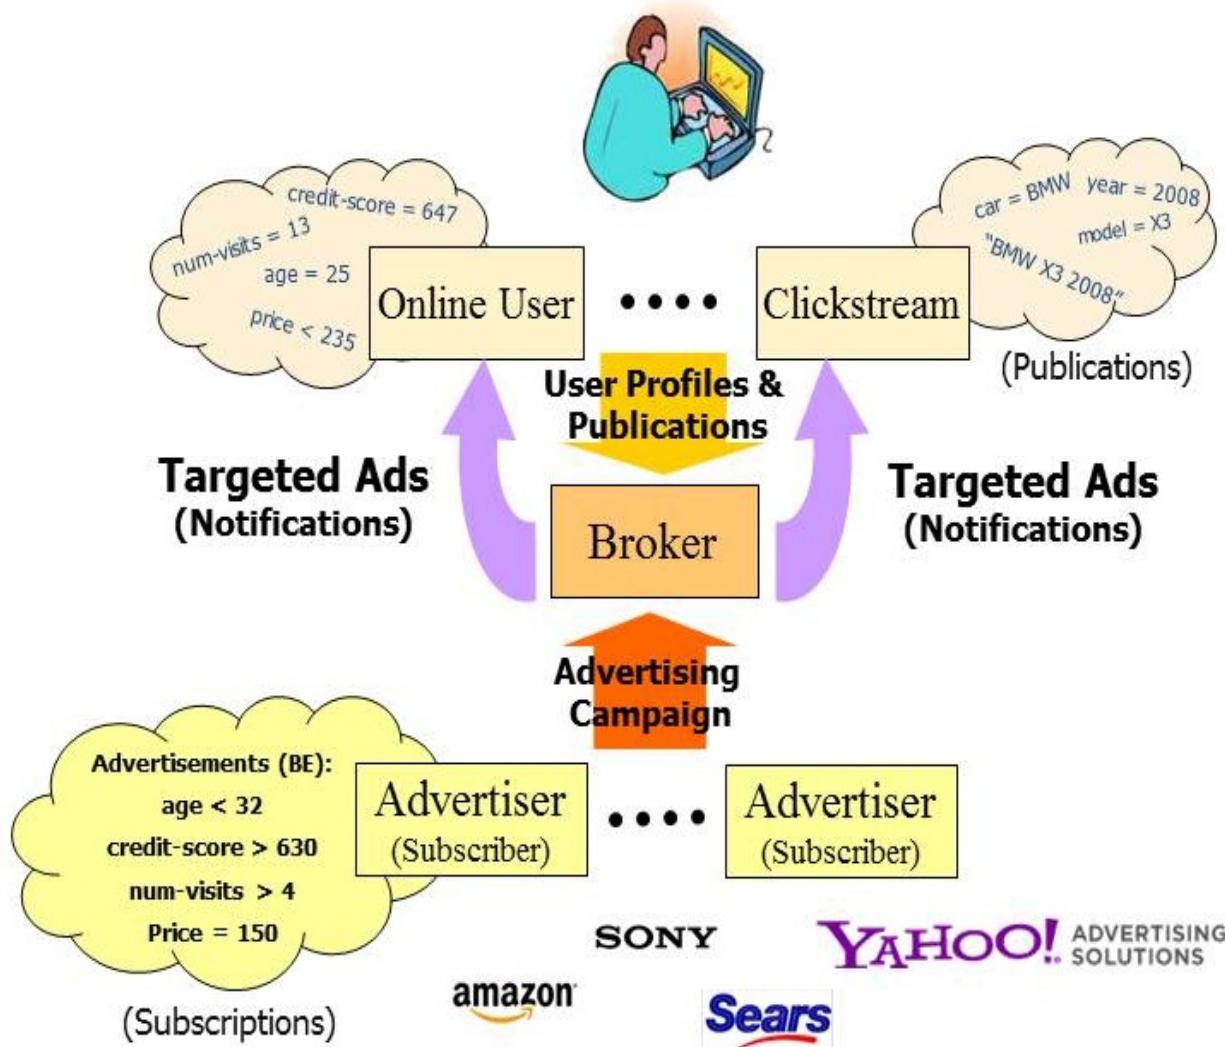
# Enterprise Application Integration



# Enterprise Application Integration



# Computational Advertising: Filtering & Matching



# Pub/Sub Standards

- DCE Event Management Service (1990s)
- OMG Event Channel (1995)
- OMG Notification Service (1997)
- OMG Data Dissemination Service
- MQ Telemetry Transport (1999)
- Java Messaging Service (JMS) (1998, 2002)
- WS Eventing (2004)
- WS Notifications (2004)
- Active Message Queueing Protocol (AMQP)
- OGSI-Notification (6/2003) (Open Grid Services Architecture)



# Popular Systems in Industry



## PubSubHubbub

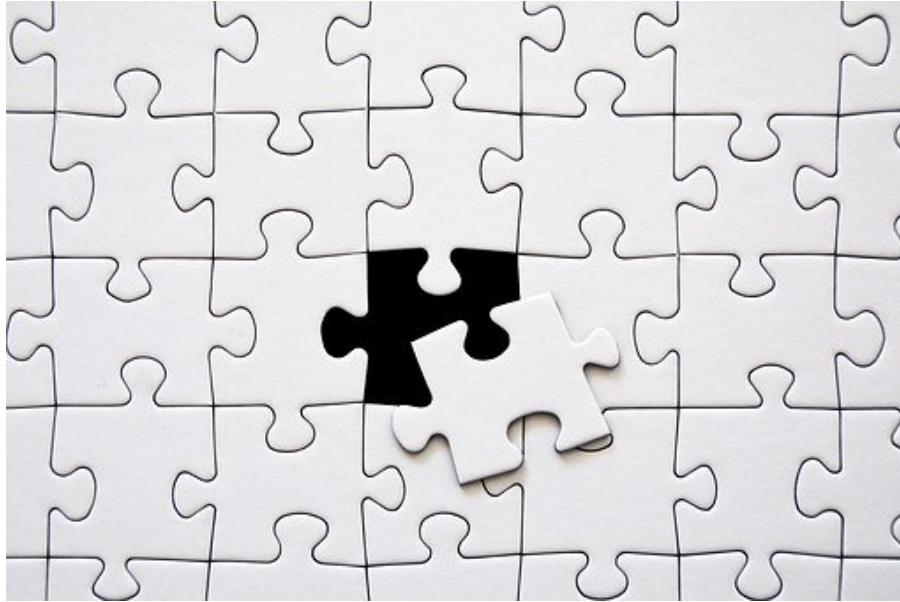
Facebook Wormhole



# Summary of Pub/Sub

- Pub/sub provides **many-to-many** communication between publishers and subscribers
- Communication is decoupled with respect to at least **time, space, and synchronization**
- Many applications are **event-based** or require **event notifications**, which leverage pub/sub
- Many flavors of pub/sub, which differ in **matching** and **routing**



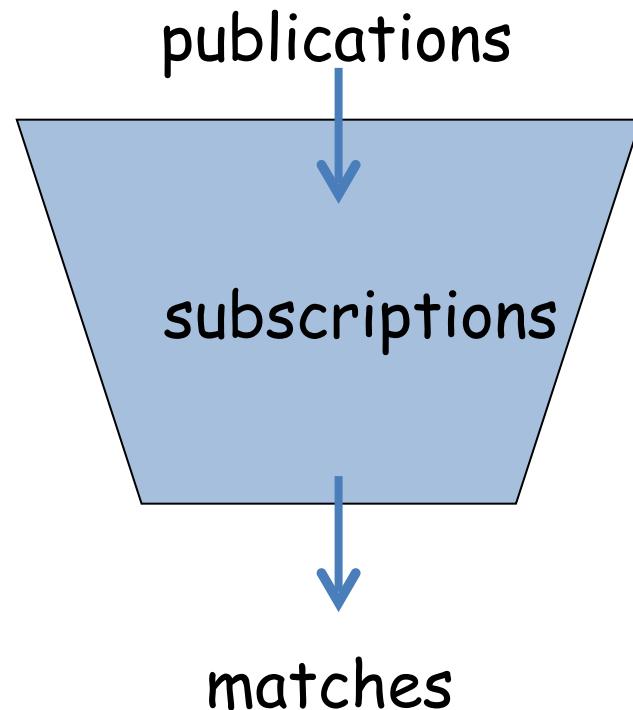


Pixabay.com

# **PUBLISH/SUBSCRIBE MATCHING MODELS**

# Matching & Filtering Problem in Pub/Sub

- **Matching problem:** Given a publication,  $e$ , and a set of subscriptions,  $S$ , determine all subscriptions,  $s \in S$ , that match  $e$ .
- Different for each model
  - Topic-based filtering
  - Content-based filtering
  - Type-based filtering
  - ...



# Matching Model

- The matching model describes **how** subscriptions and publications are **expressed**.
- Each subscription contains **filters** which determine whether a publication matches or not.
- When a publication is sent, it is compared against each subscription to see which ones **match**.
- The publication is then delivered to **each subscriber with a matching subscription**.

# Topic-Based Matching

- A.k.a. subject-based, group-based or channel-based event filtering (yet, there are subtle differences!)
- Publish(data, metadata) → Publish(data, topic)
- Subscribe(interest) → Subscribe(topic)
- A publication **matches** a subscription if it is published on the **same topic** as the subscription.
- Example:  
`subscribe("IBM"),  
publish("price:175.31","IBM")`

# Topic-Based Matching II

- Static publisher & subscriber relationship
  - At publication time, the subscriber recipient set is known
  - For a given topic, unless topics or subscriptions change, every message published on the topic, is delivered to the same recipient set
- Topics can be hierarchically organized
  - E.g., sports/basketball/NBA
  - E.g., publish message to sports/basketball/NBA
- Topic subscriptions may contain wildcards
  - E.g., Subscribe to sports/basketball/\*

# Content-Based Matching

- **Subscription:** Conjunction of **predicates** (i.e., an attribute-operator-value triple)  
$$(\text{subject} = \text{news}) \wedge (\text{topic} = \text{travel}) \wedge (\text{date} > 29.11.2011)$$
- **Publication (a.k.a. event):** Sets of attribute-value pairs  
$$(\text{subject}, \text{news}), (\text{topic}, \text{travel}), (\text{date}, 21.2.2011), \dots$$

# Content-Based Matching II

- Publish(data, metadata) → Publish(data, list(attribute-value))
- Subscribe(interest) → Subscribe(predicates)
- A pub matches a sub if all the predicates are satisfied (conjunction).
- Example:

subscribe("stock = IBM", "price > 175.0", "date > 1-1-1970")

publish("stock = IBM", "price = 175.31", "date = 31-12-1969")

publish("stock = IBM", "price = 175.5", "date = 1-1-2017")

publish("stock = IBM", "price = 180.0", "market = NYSE")

publish("stock = IBM", "price = 177.5", "date = 1-1-2017",  
"market = NYSE")

Example	Tree-structured data	Graph-structured data	Unstructured data	Regular languages	Relational model
<b>Subscription</b>	XPath	RDF Query	Keywords	Regular expressions	SQL
<b>Publication</b>	XML	RSS feeds	Text, documents	Sentences over some alphabet	DBs, i.e., tables

# Content-Based Matching II

- Publish(data, metadata) → Publish(data, list(attribute-value))
- Subscribe(interest) → Subscribe(predicates)
- A pub **matches** a sub if **all the predicates** are satisfied (conjunction).
- Example:

subscribe("stock = IBM", "price > 175.0", "date > 1-1-1970")

✗ publish("stock = IBM", "price = 175.31", "date = 31-12-1969")

publish("stock = IBM", "price = 175.5", "date = 1-1-2017")

publish("stock = IBM", "price = 180.0", "market = NYSE")

publish("stock = IBM", "price = 177.5", "date = 1-1-2017",  
"market = NYSE")

Example	Tree-structured data	Graph-structured data	Unstructured data	Regular languages	Relational model
<b>Subscription</b>	XPath	RDF Query	Keywords	Regular expressions	SQL
<b>Publication</b>	XML	RSS feeds	Text, documents	Sentences over some alphabet	DBs, i.e., tables

# Content-Based Matching II

- Publish(data, metadata) → Publish(data, list(attribute-value))
- Subscribe(interest) → Subscribe(predicates)
- A pub **matches** a sub if **all the predicates** are satisfied (conjunction).
- Example:

subscribe("stock = IBM", "price > 175.0", "date > 1-1-1970")

X publish("stock = IBM", "price = 175.31", "date = 31-12-1969")

✓ publish("stock = IBM", "price = 175.5", "date = 1-1-2017")

publish("stock = IBM", "price = 180.0", "market = NYSE")

publish("stock = IBM", "price = 177.5", "date = 1-1-2017",  
"market = NYSE")

Example	Tree-structured data	Graph-structured data	Unstructured data	Regular languages	Relational model
<b>Subscription</b>	XPath	RDF Query	Keywords	Regular expressions	SQL
<b>Publication</b>	XML	RSS feeds	Text, documents	Sentences over some alphabet	DBs, i.e., tables

# Content-Based Matching II

- Publish(data, metadata) → Publish(data, list(attribute-value))
- Subscribe(interest) → Subscribe(predicates)
- A pub **matches** a sub if **all the predicates** are satisfied (conjunction).
- Example:

subscribe("stock = IBM", "price > 175.0", "date > 1-1-1970")

X publish("stock = IBM", "price = 175.31", "date = 31-12-1969")

✓ publish("stock = IBM", "price = 175.5", "date = 1-1-2017")

X publish("stock = IBM", "price = 180.0", "market = NYSE")

publish("stock = IBM", "price = 177.5", "date = 1-1-2017",  
"market = NYSE")

Example	Tree-structured data	Graph-structured data	Unstructured data	Regular languages	Relational model
<b>Subscription</b>	XPath	RDF Query	Keywords	Regular expressions	SQL
<b>Publication</b>	XML	RSS feeds	Text, documents	Sentences over some alphabet	DBs, i.e., tables

# Content-Based Matching II

- Publish(data, metadata) → Publish(data, list(attribute-value))
- Subscribe(interest) → Subscribe(predicates)
- A pub **matches** a sub if **all the predicates** are satisfied (conjunction).
- Example:

subscribe("stock = IBM", "price > 175.0", "date > 1-1-1970")

✗publish("stock = IBM", "price = 175.31", "date = 31-12-1969")

✓publish("stock = IBM", "price = 175.5", "date = 1-1-2017")

✗publish("stock = IBM", "price = 180.0", "market = NYSE")

✓publish("stock = IBM", "price = 177.5", "date = 1-1-2017",  
"market = NYSE")

Example	Tree-structured data	Graph-structured data	Unstructured data	Regular languages	Relational model
<b>Subscription</b>	XPath	RDF Query	Keywords	Regular expressions	SQL
<b>Publication</b>	XML	RSS feeds	Text, documents	Sentences over some alphabet	DBs, i.e., tables

# Language Categories

Simple

- Channel-based
  - OMG CORBA Event Service, group communication ...
- Topic-based (Subject-based)
  - WS Notifications, OMG Data Dissemination Service ...
- Type-based
  - OMG Data Dissemination Service (partially), ...
- Content-based
  - The PADRES ESB (see below), ...

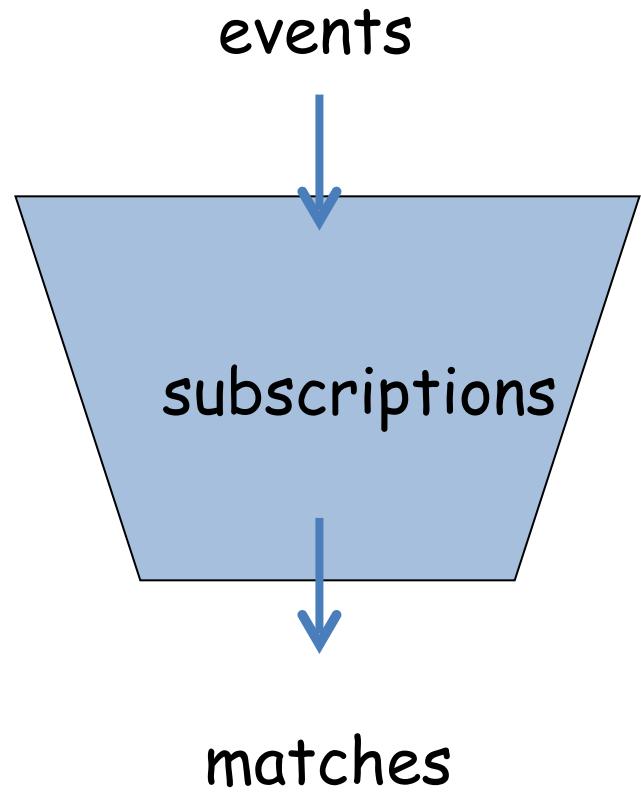
Stateless

- 
- State-based
    - Many rule-based systems & tuple spaces
  - Subject Spaces (state persistent)

Stateful

Complex – Still, a research idea, concept & prototype

**Matching problem:** Given an event,  $e$ , and a set of subscriptions,  $S$ , determine all subscriptions,  $s \in S$ , that match  $e$ .



## CONTENT-BASED MATCHING ALGORITHMS

# Matching Algorithms

- Example algorithms
  - **Counting algorithm** [origins unknown to me; guess: 1970s]
  - *Clustering algorithm* [Fabret, Jacobsen et al., 2001]
- Both are **two-phased** matching algorithms
  1. Match all predicates (**predicate matching phase**)
  2. Match subscriptions from results of Phase 1 (**subscriptions matching phase**)
- Let us look at both phases separately
  - Predicate matching
  - Subscription matching

# Two-Phased Matching

- **Decompose subscriptions into predicates**

s<sub>1</sub> (subject <sup>p<sub>1</sub></sup>= news)  $\wedge$  (topic <sup>p<sub>2</sub></sup>= travel)  $\wedge$  (date > <sup>p<sub>3</sub></sup>29.11.2011)

s<sub>2</sub> (subject = news)  $\wedge$  (topic = stock)  $\wedge$  (date > 30.11.2011)  
                <sup>p<sub>1</sub></sup>                            <sup>p<sub>4</sub></sup>                            <sup>p<sub>5</sub></sup>

- **Convert subscriptions to reference predicates**

S1: P1  $\wedge$  P2  $\wedge$  P3

S2: P1  $\wedge$  P4  $\wedge$  P5

Break down subscriptions into predicates, store and match them individually

# **PHASE 1: PREDICATE MATCHING**

# Predicate Matching Problem

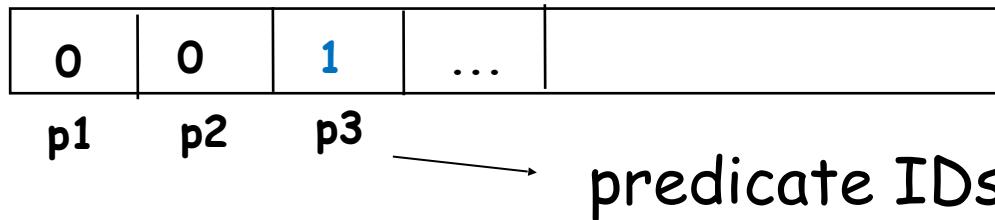
- Given a set  $P$  of predicates and a publication  $e$ , *identify all predicates  $p$  of  $P$  which evaluate to true under  $e$ .*
- Example:

$e = \{\dots, (\text{price}, 5), (\text{color}, \text{white}), \dots\}$

$p_1: \text{price} > 5; \quad p_2: \text{color} = \text{red}; \quad p_3: \text{price} < 7$

$p_1$  is false       $p_2$  is false       $p_3$  is true

Predicate bit vector:

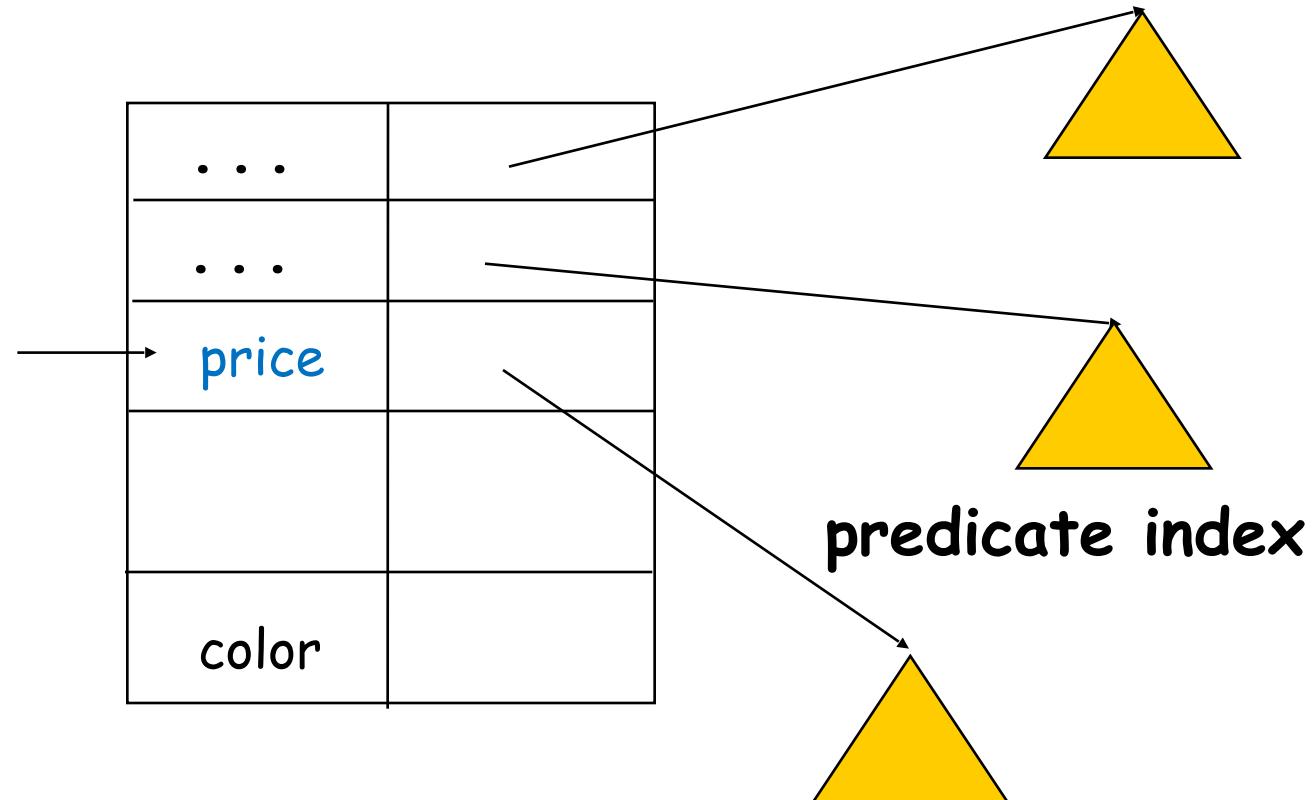


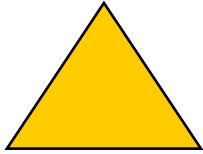
# Predicate Matching:

## Top-level Data Structure

- Hash table on attribute name

$e = \{..., (\text{price}, 5), ...\}$



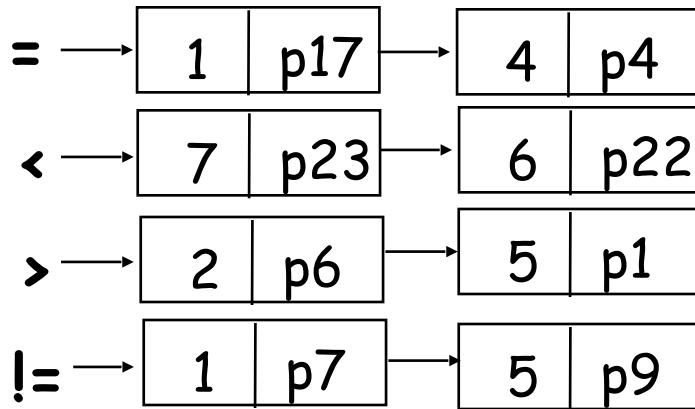


price

## Predicate Index:

# General Purpose Data Structure

Operators



p17: price = 1;

p11: price < 0;

p4: price = 4

p3: price < 4

- One ordered linked list per operator
- Insert, delete, match are  $O(n)$ -operations (per attribute name in  $e$  and per operator)
- Alternatively, use a B-tree or B+-tree etc.

# Finite Predicate Value Domain Types

- Countable domain types with small cardinality
  - Integer intervals
  - Collections (enums)
  - Set of tags (e.g., in XML)
- Examples
  - Price : [0, 1000], models variety of prices
  - Color, city, state, country, size, weight
  - All tags defined in a given XML schema
  - Predicate domain often context dependant, but limited in size
    - Prices of cars vs. prices of groceries

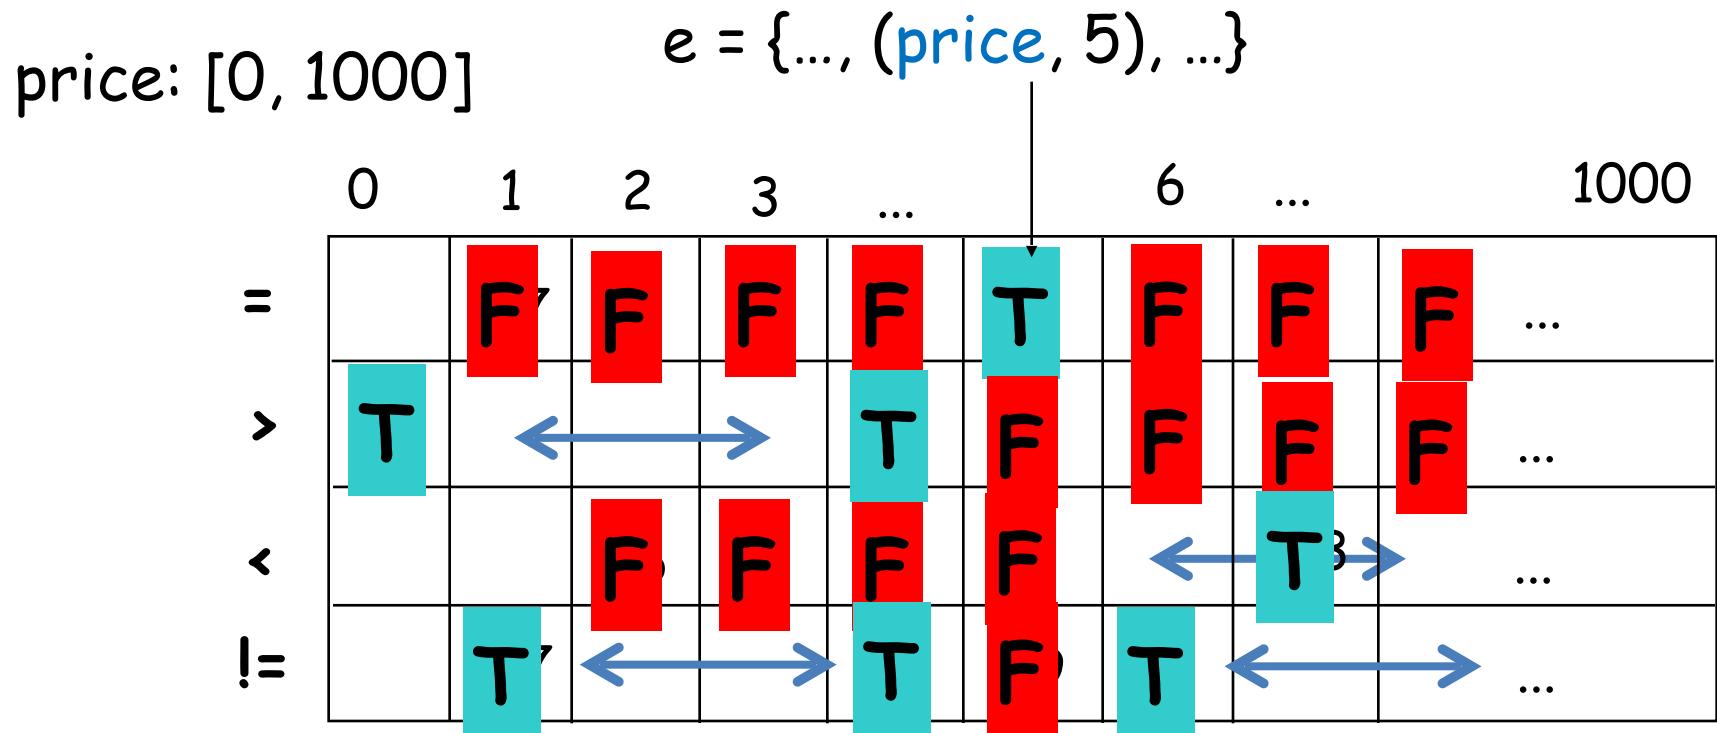
# Predicate Matching For Finite Domains

price: [0, 1000]       $e = \{\dots, (\text{price}, 5), \dots\}$

	0	1	2	3	...	6	...	1000
=		p17				p4		...
>			p6			p1	p13	...
<		p11				p3		...
!=			p7			p9		...

p1: price > 5;    p3: price < 5;    p4: price = 5;    p9: price != 5

# Predicate Matching **Symmetries**

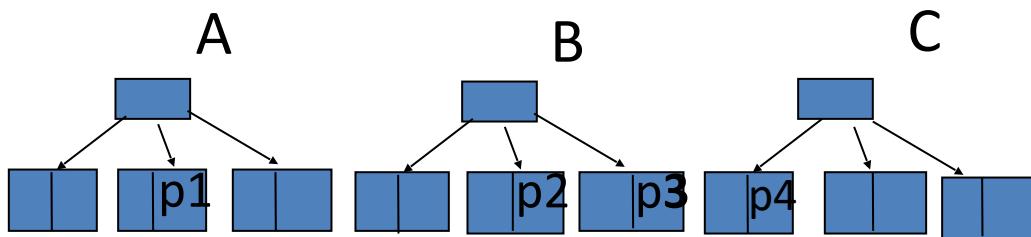


# **PHASE 2: SUBSCRIPTION MATCHING**

# Subscription Matching: Counting Algorithm

- Subscriptions consist of a set of predicates
  - S1:  $(2 < A < 4) \& (B = 6) \& (C > 4)$   $\Rightarrow p1: (2 < A < 4), p2: (B = 6), p3: (C > 4)$
  - S2:  $(2 < A < 4) \& (C = 3)$   $\Rightarrow p1: (2 < A < 4), p4: (C = 3)$
- Subscription matches the event if all its predicates are satisfied
- Idea: Count the number of satisfied predicates per subscription and compare with actual number

# Counting Algorithm: Subscription Insertion



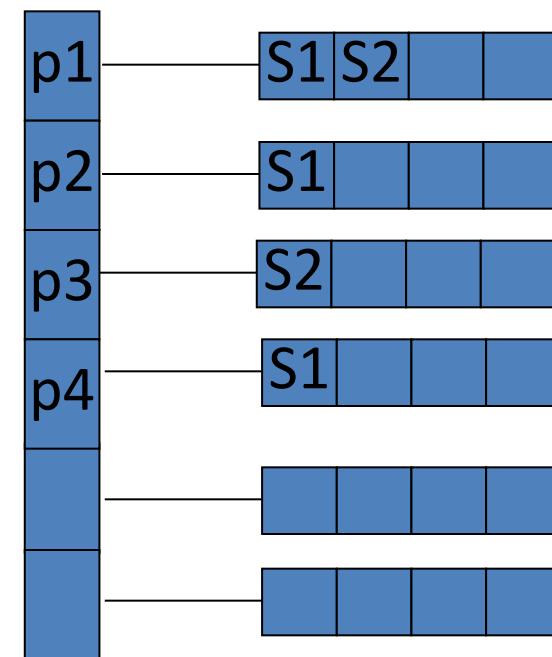
preds-per-sub      hit count

S1	3
S2	2

S1	0
S2	0

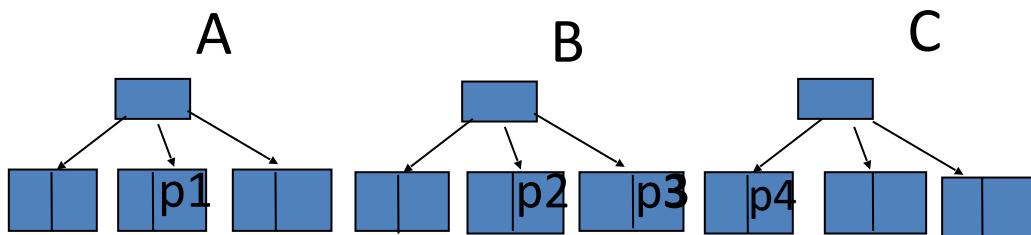
S1: p1, p2, p4  
S2: p1, p3

Predicate vector



predicate-to-subscription association

# Counting Algorithm: Subscription Insertion



preds-per-sub    hit count

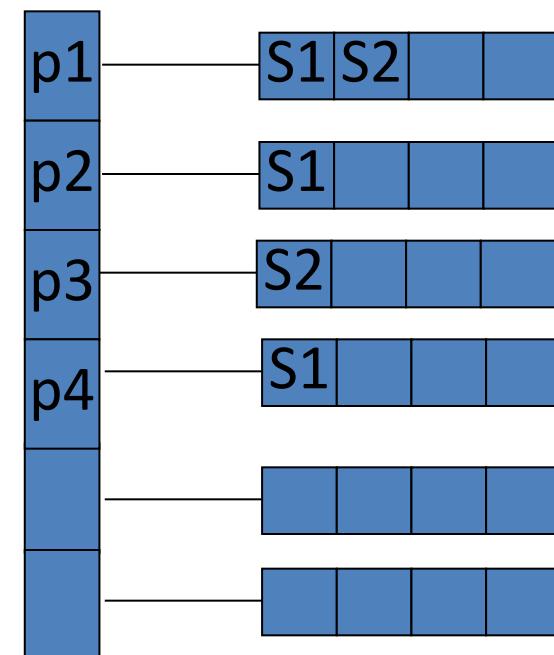
S1	3
S2	2

S1	0
S2	0

S1: p1, p2, p4  
S2: p1, p3

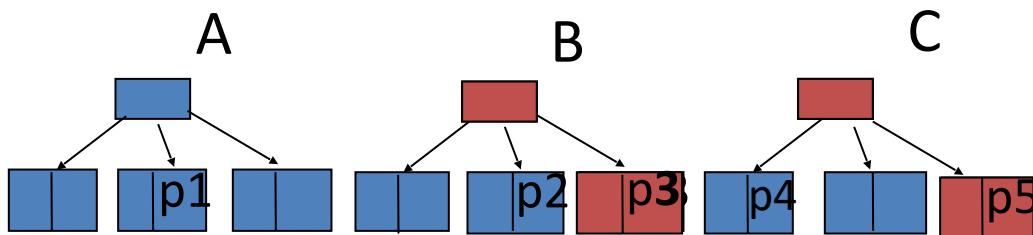
S3: p3, p5

Predicate vector



predicate-to-subscription association

# Counting Algorithm: Subscription Insertion

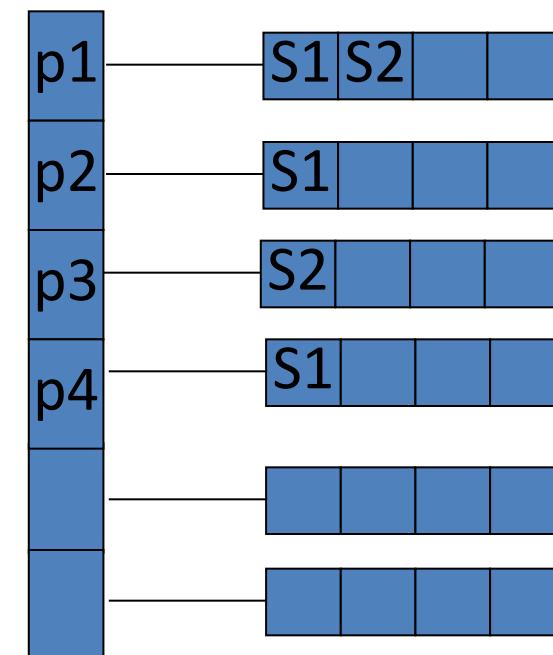


preds-per-sub      hit count

S1: p1, p2, p4  
S2: p1, p3

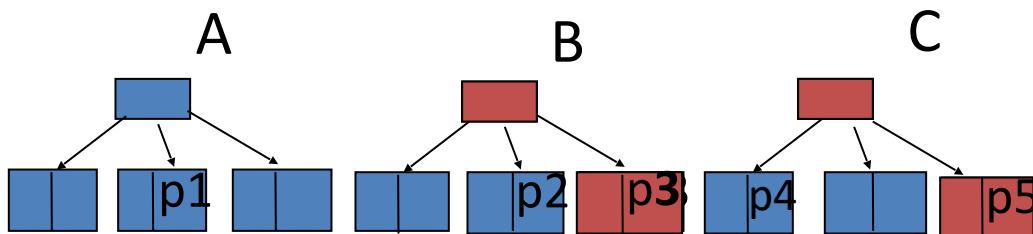
S3: p3, p5

Predicate vector



predicate-to-subscription association

# Counting Algorithm: Subscription Insertion

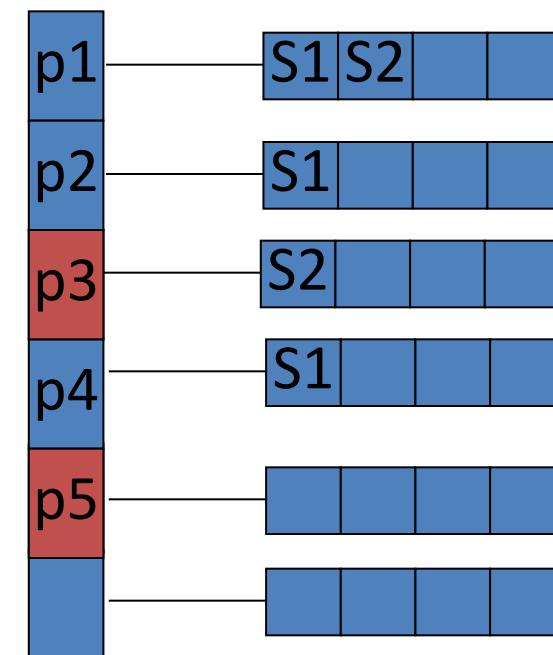


preds-per-sub      hit count

S1: p1, p2, p4  
S2: p1, p3

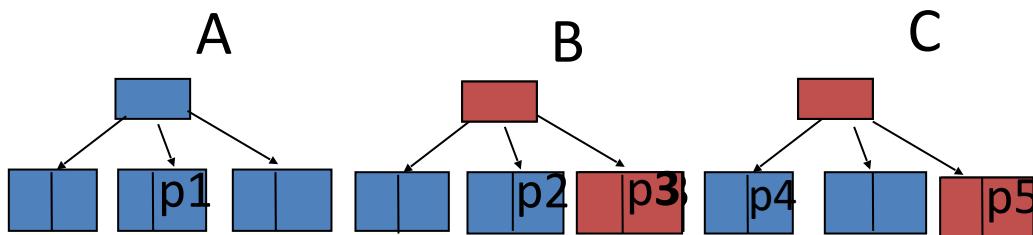
S3: p3, p5

Predicate vector



predicate-to-subscription association

# Counting Algorithm: Subscription Insertion



preds-per-sub    hit count

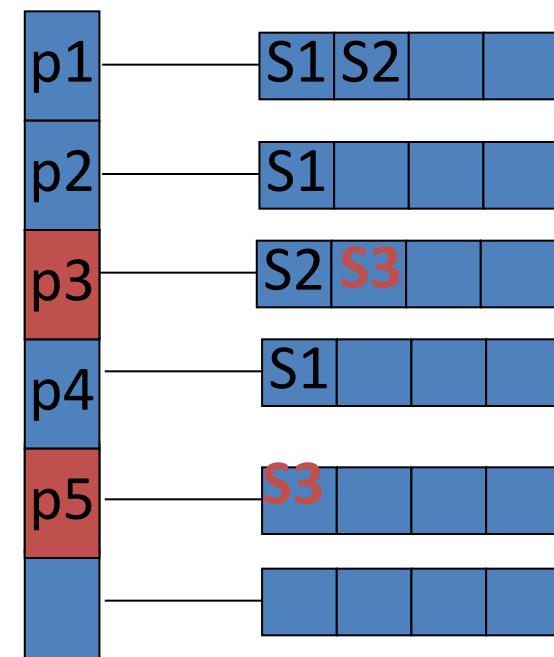
S1	3
S2	2

S1	0
S2	0

S1: p1, p2, p4  
S2: p1, p3

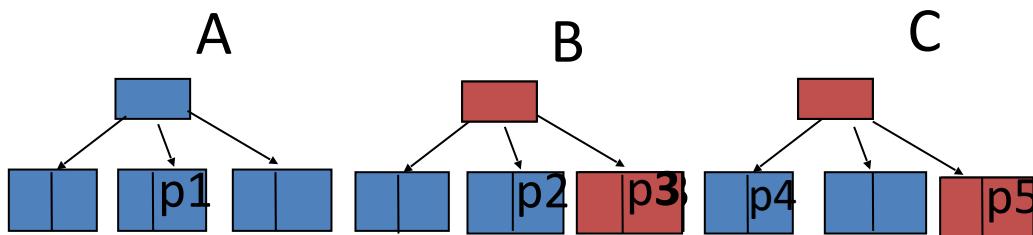
S3: p3, p5

Predicate vector



predicate-to-subscription association

# Counting Algorithm: Subscription Insertion



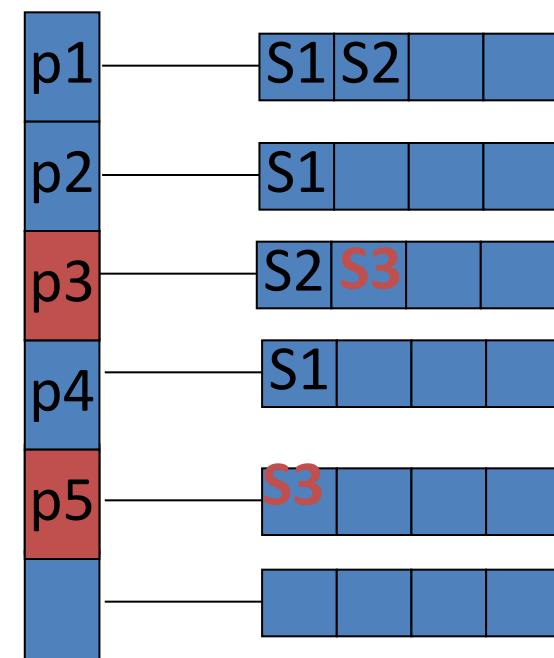
preds-per-sub    hit count

S1	3
S2	2
S3	2

S1	0
S2	0

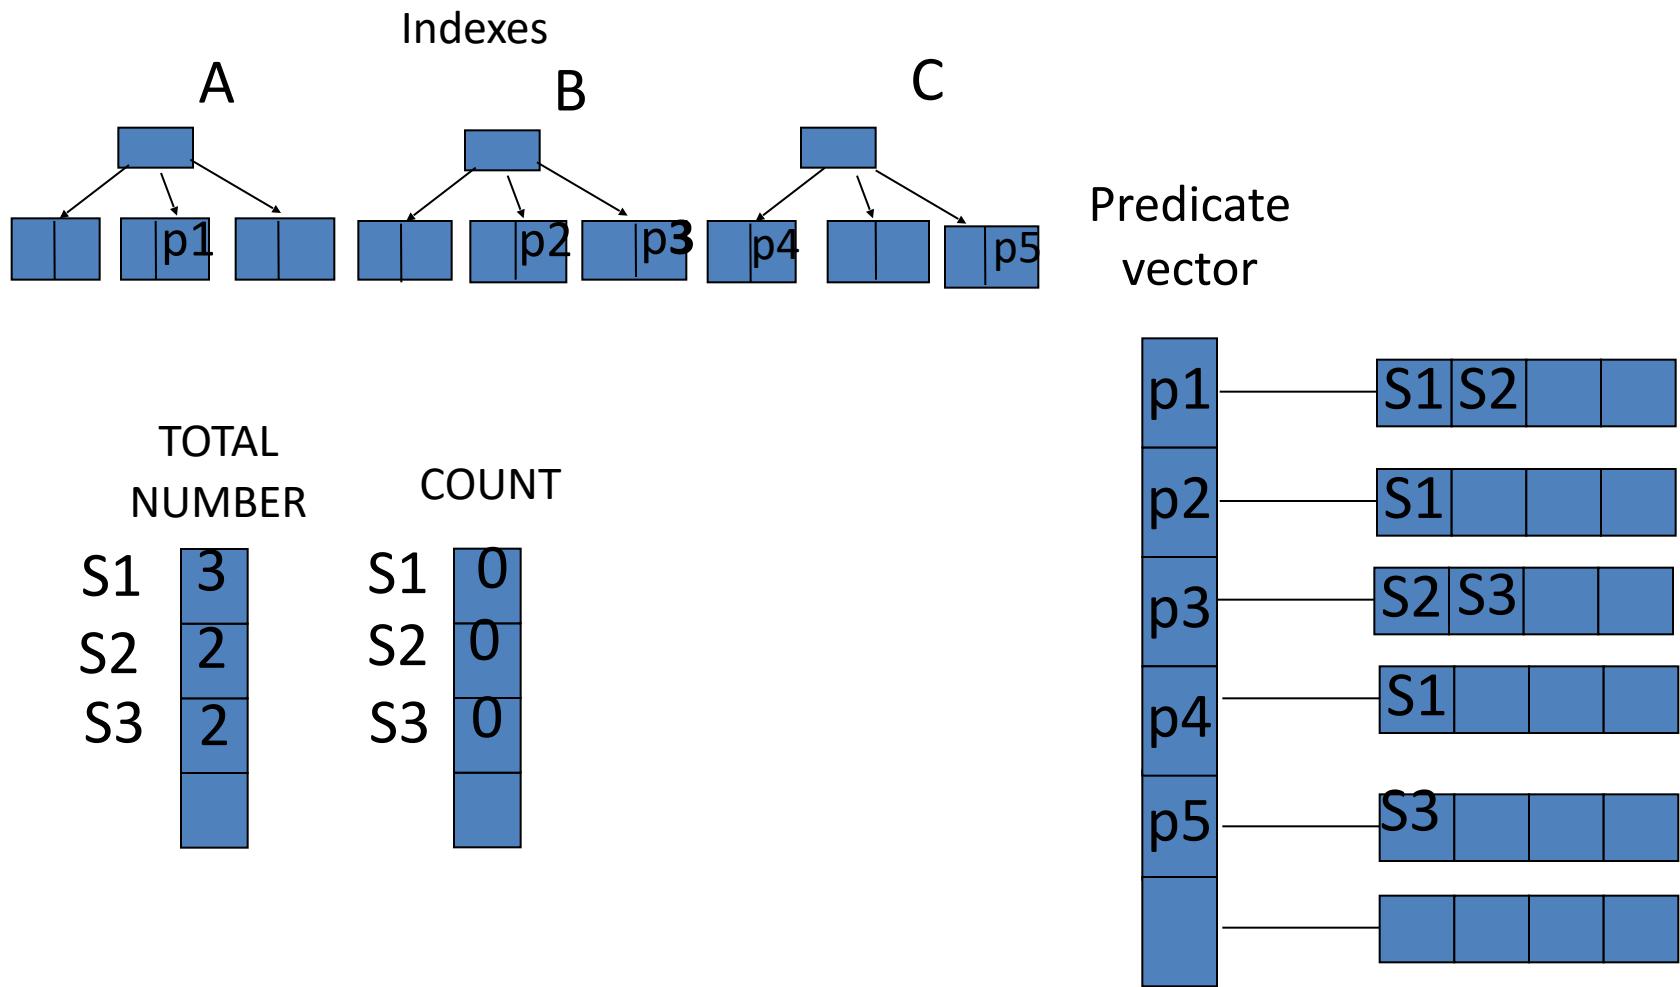
S1: p1, p2, p4  
S2: p1, p3  
S3: p3, p5

Predicate vector

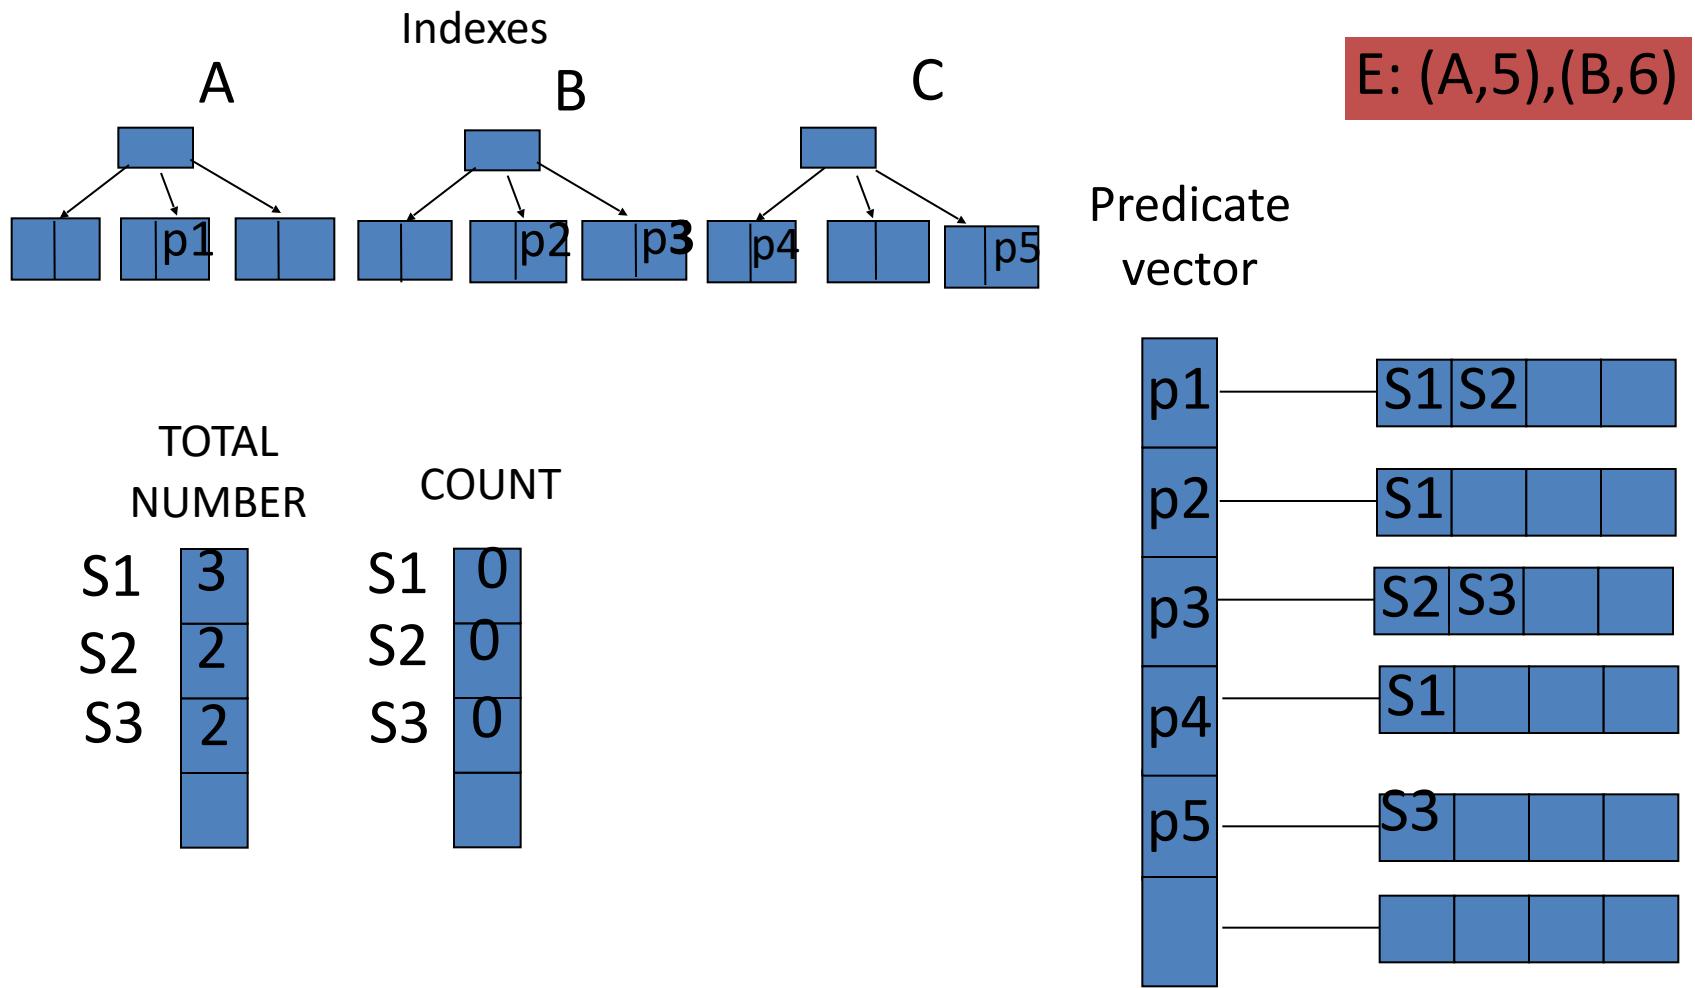


predicate-to-subscription association

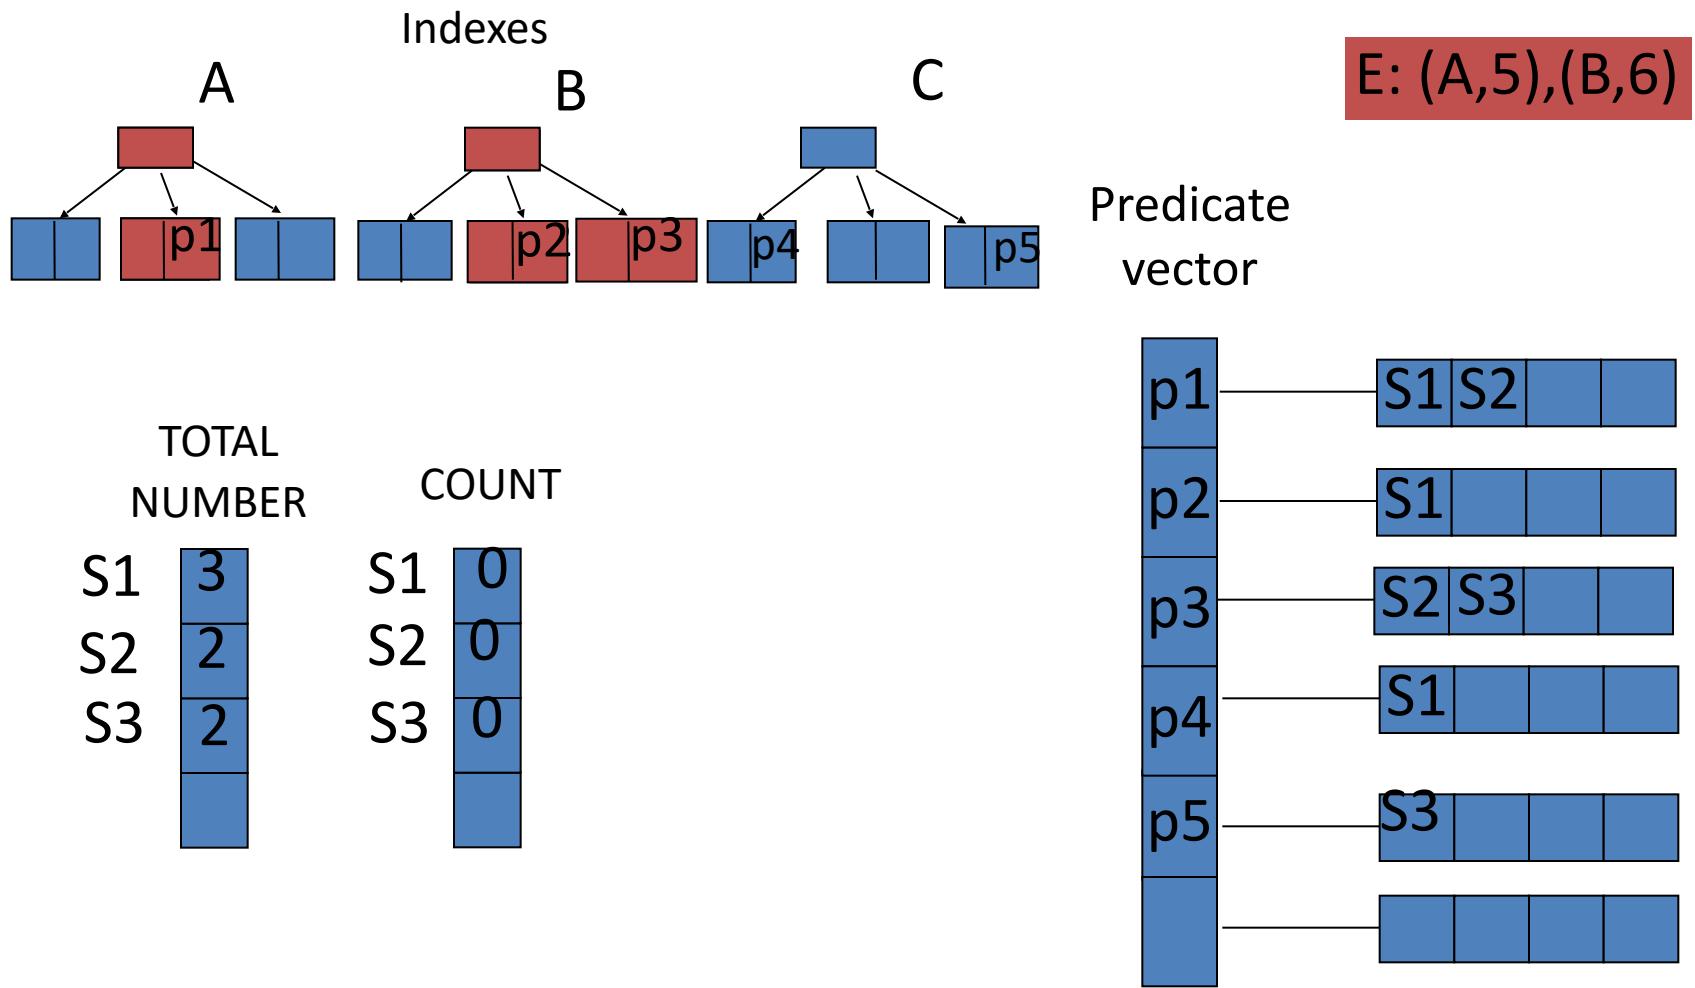
# Counting Algorithm: Event Matching



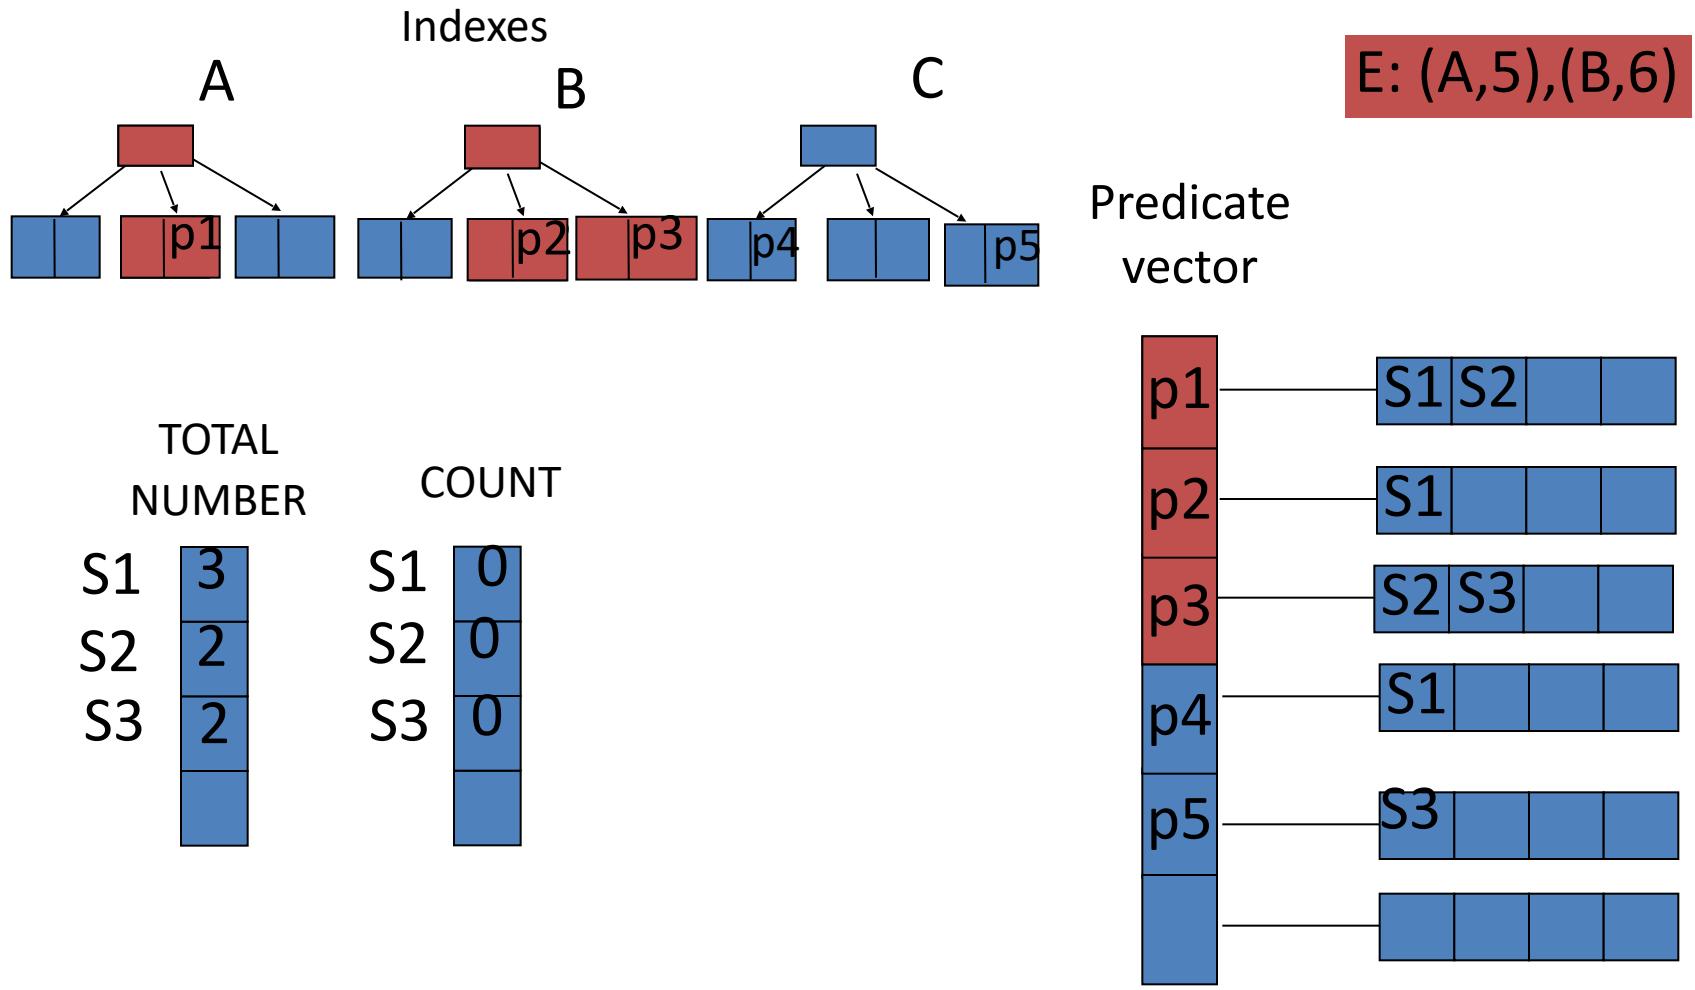
# Counting Algorithm: Event Matching



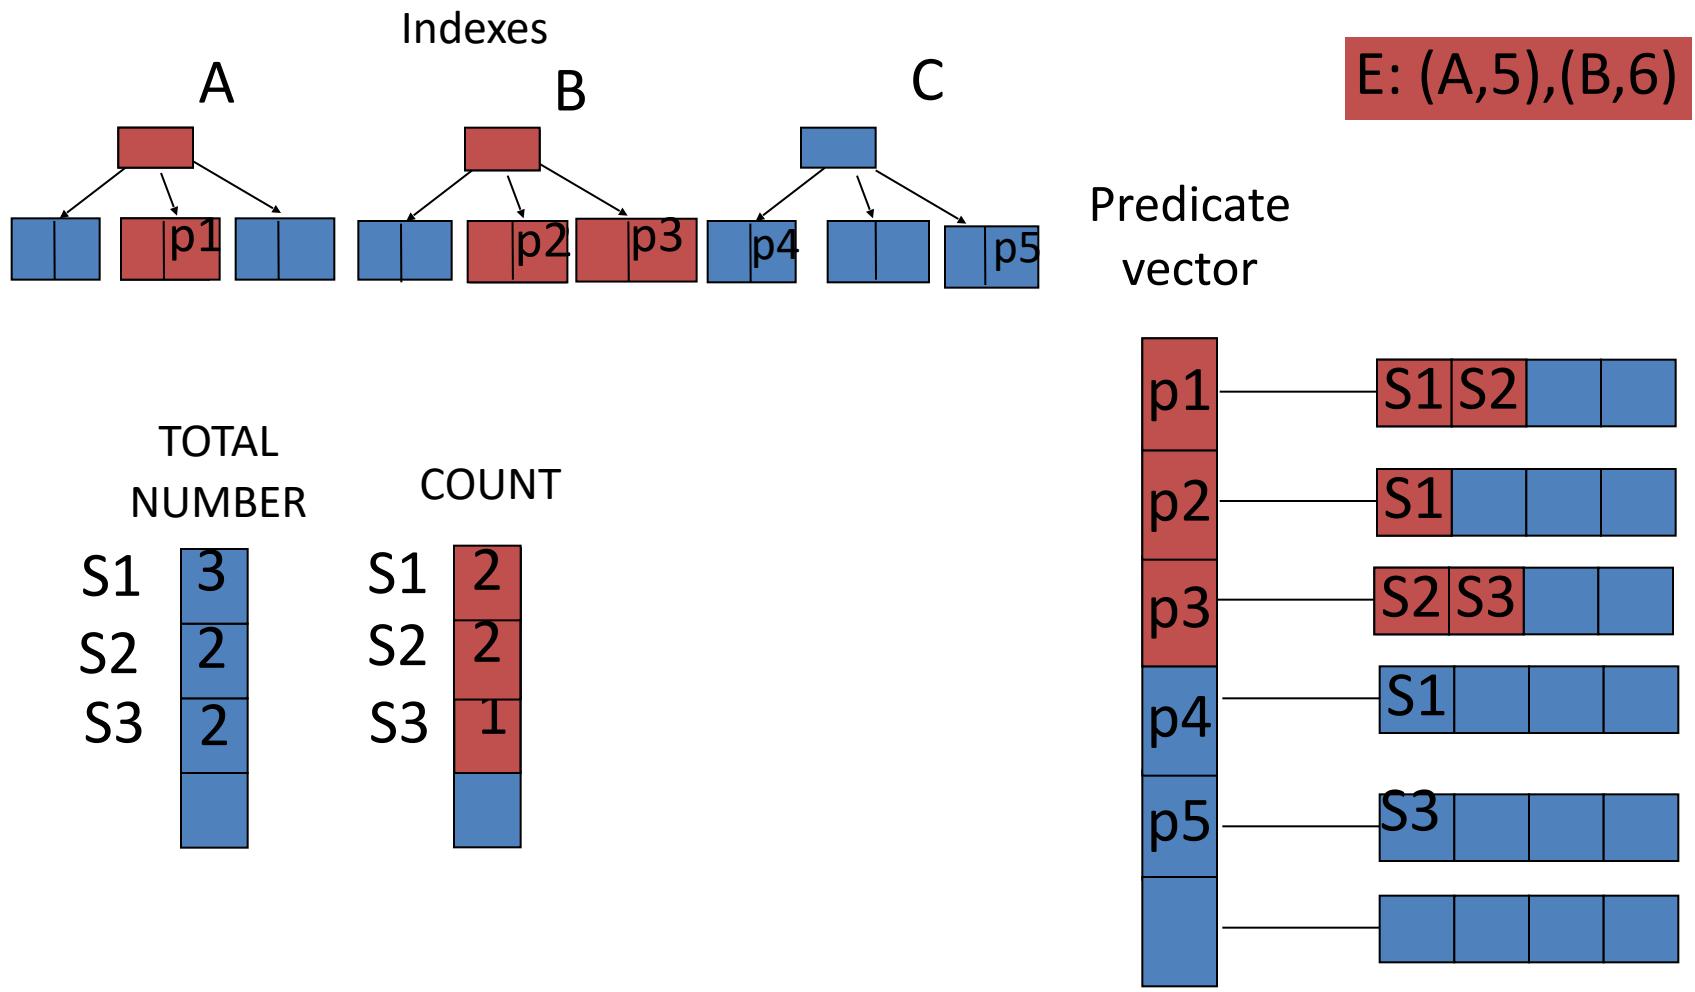
# Counting Algorithm: Event Matching



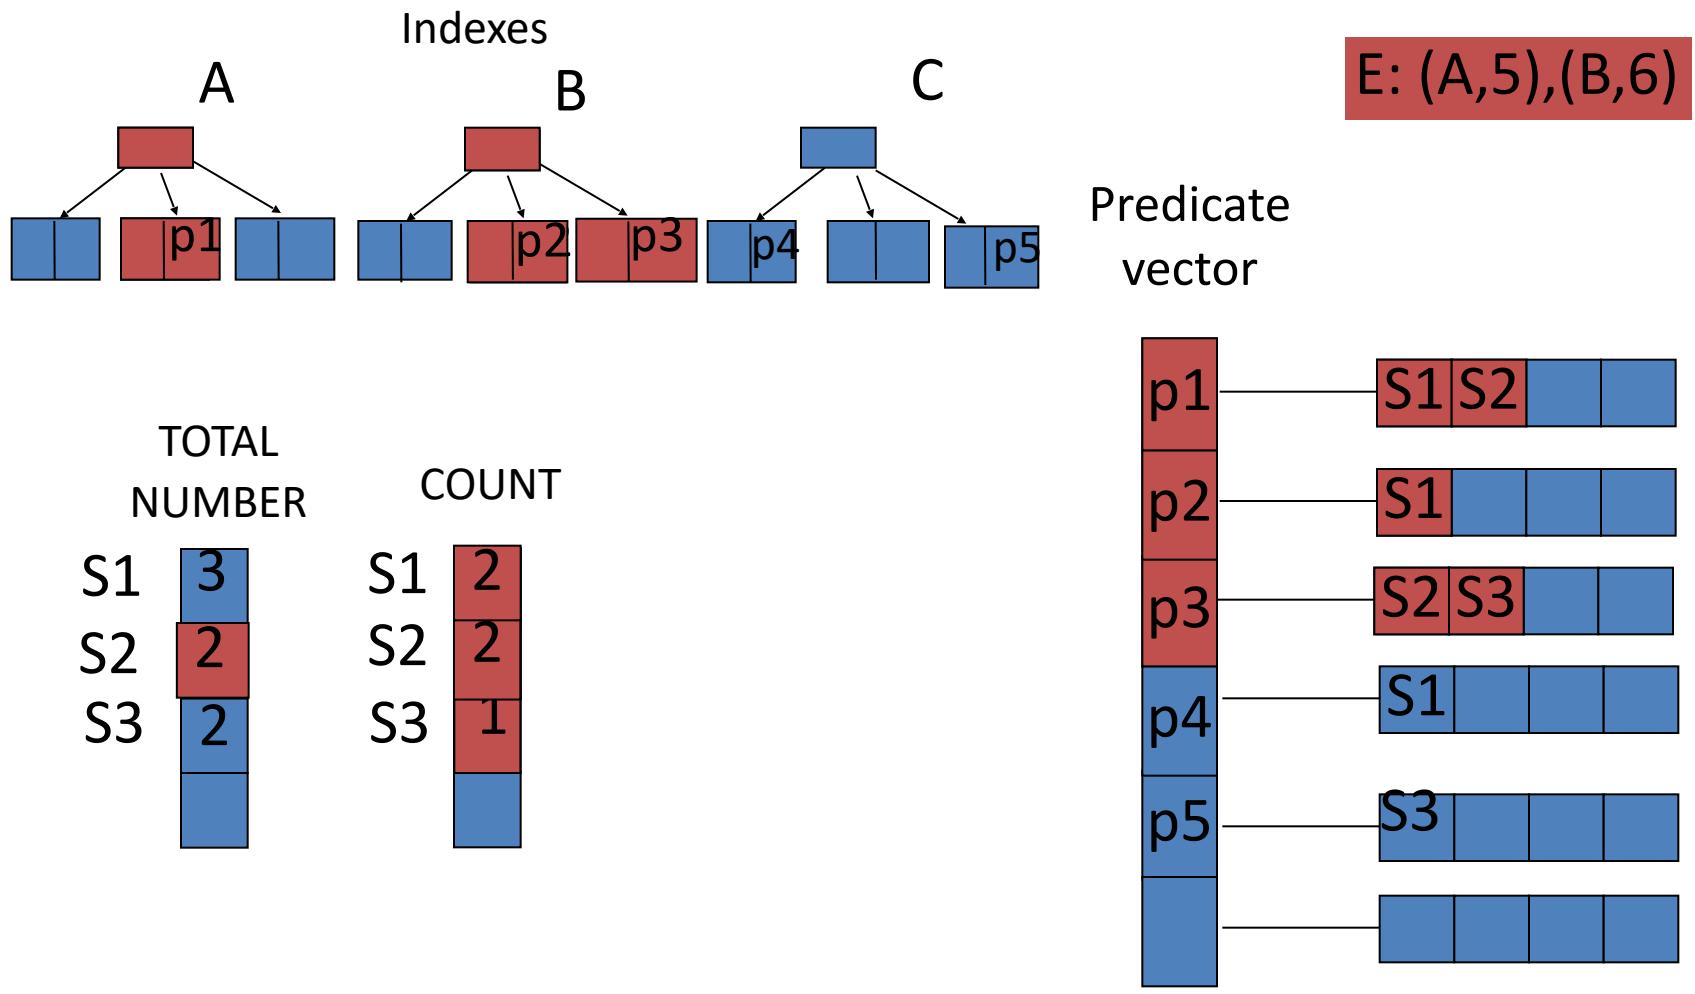
# Counting Algorithm: Event Matching



# Counting Algorithm: Event Matching



# Counting Algorithm: Event Matching



# Matching Models Summary

- **Matching model:** The language used to express pubs and subs in order to match them.
- **Matching problem:** Given a publication,  $p$ , and a set of subscriptions,  $S$ , determine all subscriptions,  $s \in S$ , that match  $e$ .
- **Topic-based** vs. **content-based** matching
- **Tradeoff:** computation vs. communication
- **Two-phase matching** is an efficient technique for quickly finding matches in content-based model
- **Counting algorithm**



Pixabay.com

# **PUBLISH/SUBSCRIBE ROUTING MODELS**

# Pub/Sub Architecture

- Pub/Sub Middleware
  - Comprised of a dedicated **pub/sub middleware**
  - Receives all operations
  - Stores subscriptions
  - Performs **publication matching and dissemination**
- Infrastructure-less
  - No dedicated middleware
  - Clients (publishers and subscribers) perform the tasks
  - **Peer-to-peer model**

# Centralized vs. Distributed

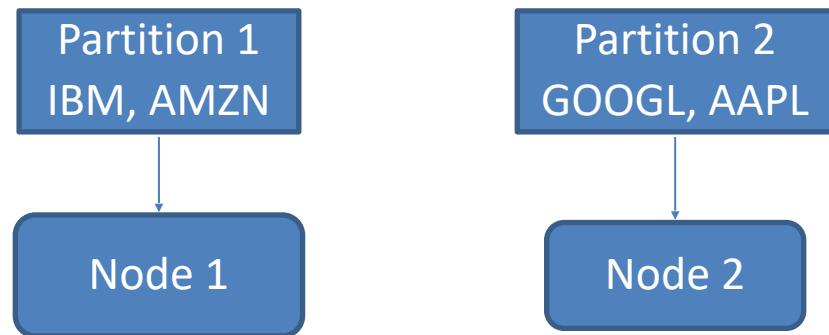
- **Centralized**
  - A single pub/sub entity
  - Client-server interactions
  - CORBA, JMS, ...
- **Decentralized**
  - Matching and dissemination tasks **distributed**
  - Either, multiple pub/sub entities (called **brokers**)
  - Or, P2P model without brokers (e.g., DHT)

# **RENDEZVOUS-BASED ROUTING**

# Rendezvous-based Routing

- The “publication space” is partitioned
- Each node is in charge of a partition
- Publications and subscriptions belonging to a partition are sent to the corresponding node

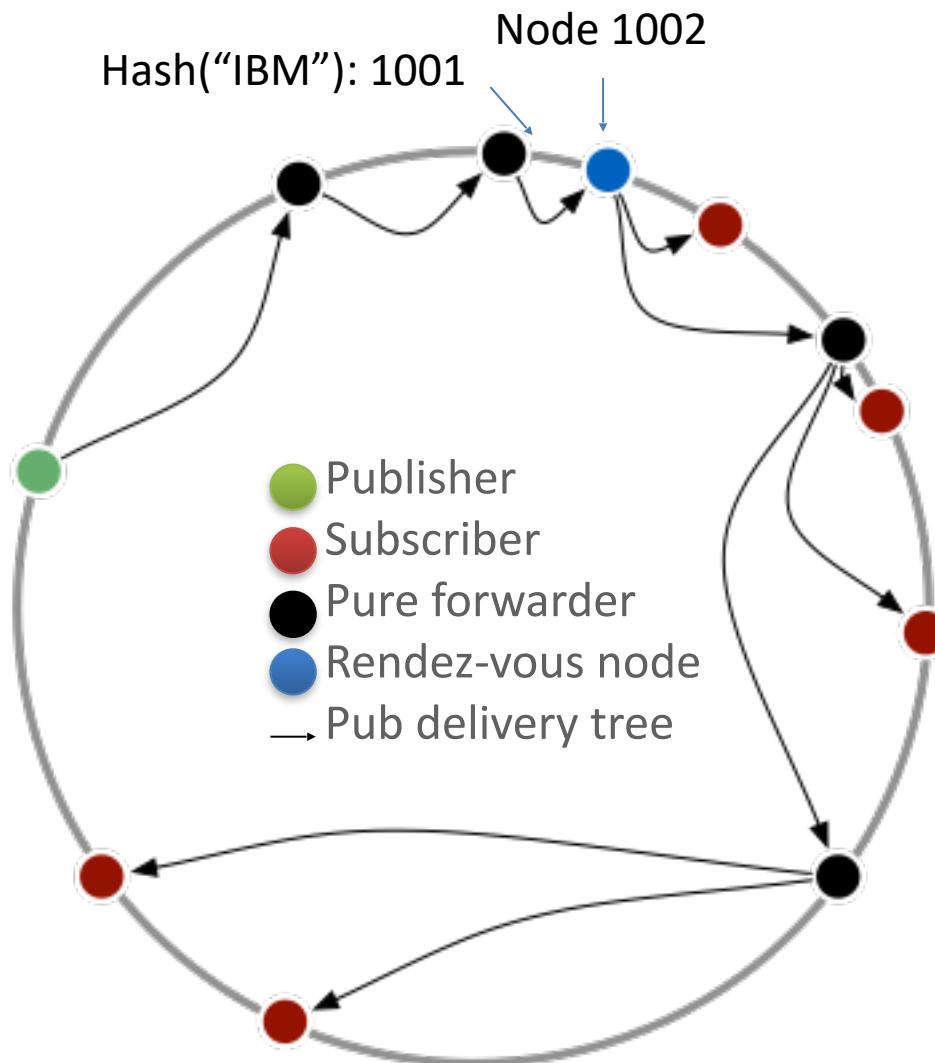
Example for topic-based:



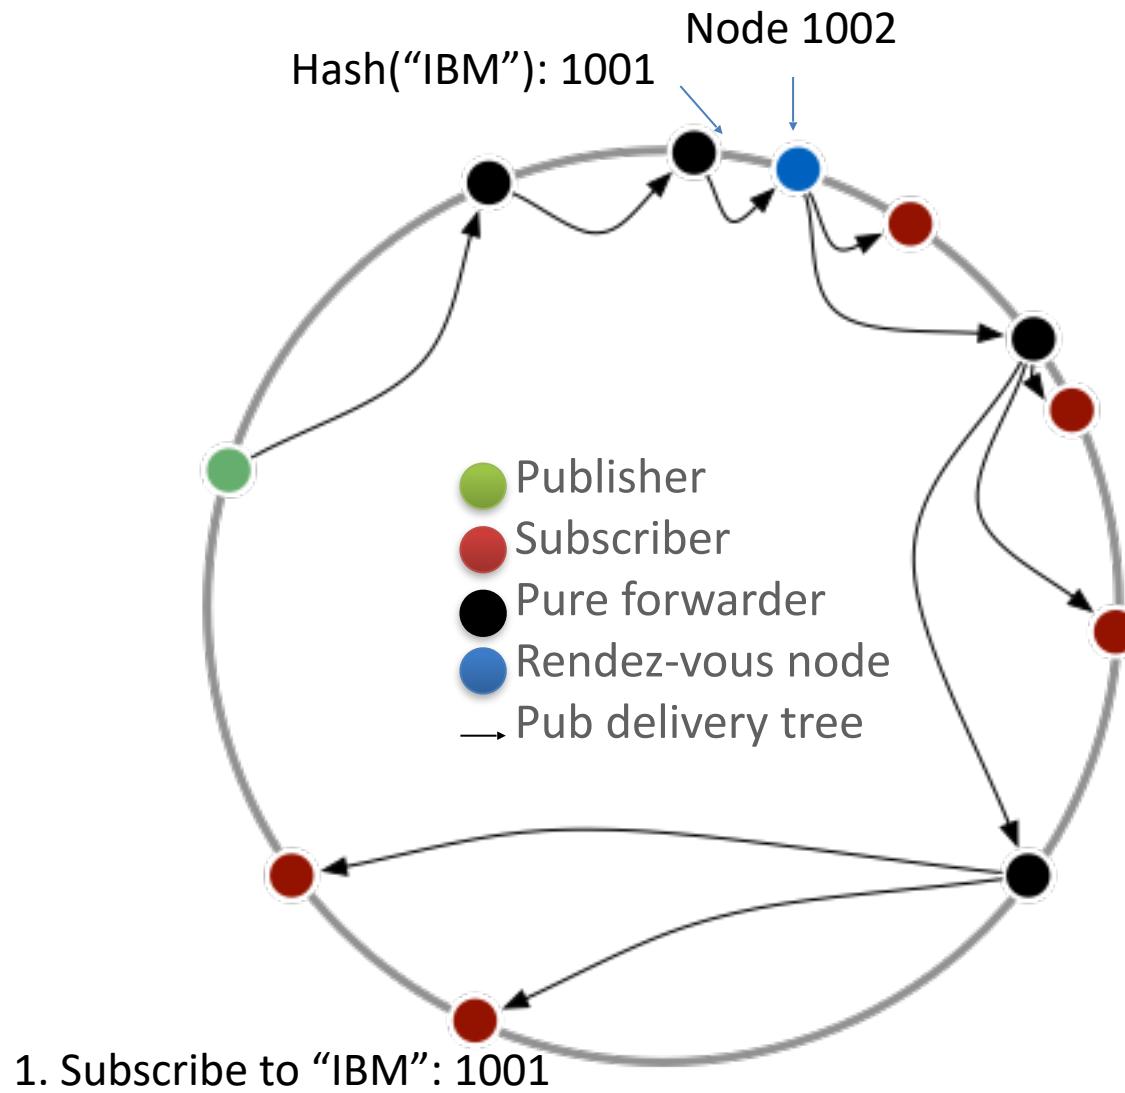
# P2P Pub/Sub: Scribe (2001)

- Infrastructure-less, distributed, pub/sub network built on top of a P2P network
- Relies on a structured P2P DHT (Pastry)
- Supports topic-based pub/sub API on top of Pastry
- Topics are hashed in the DHT

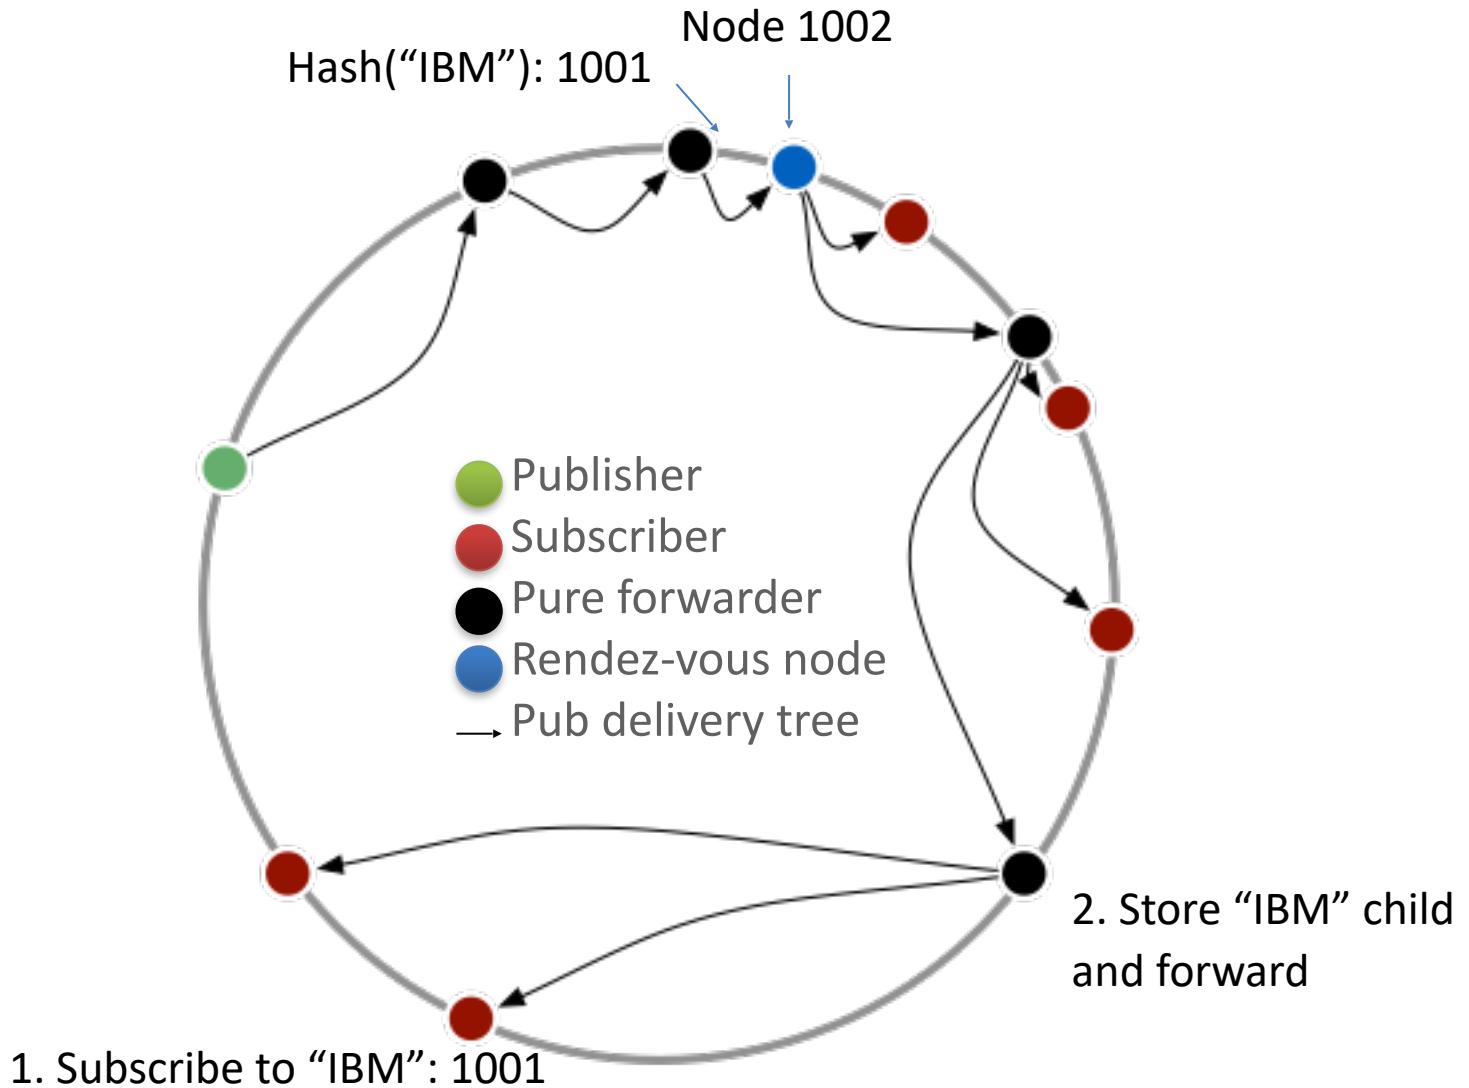
# Scribe Routing



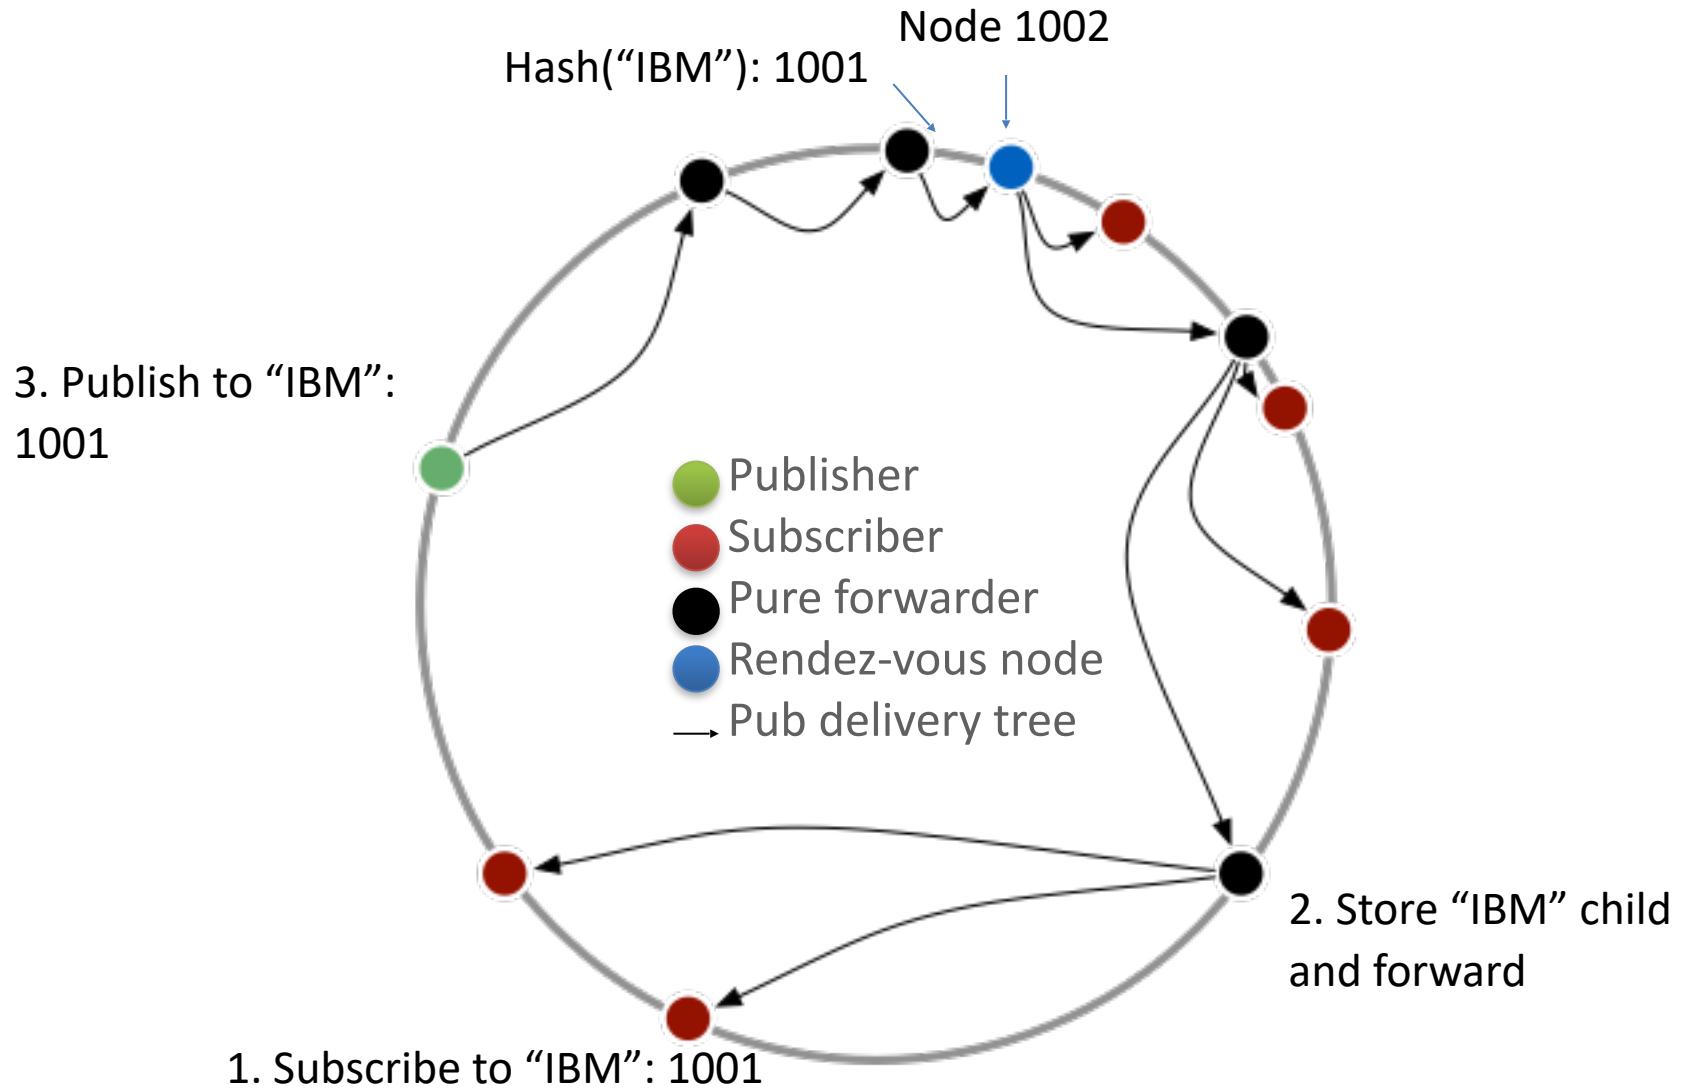
# Scribe Routing



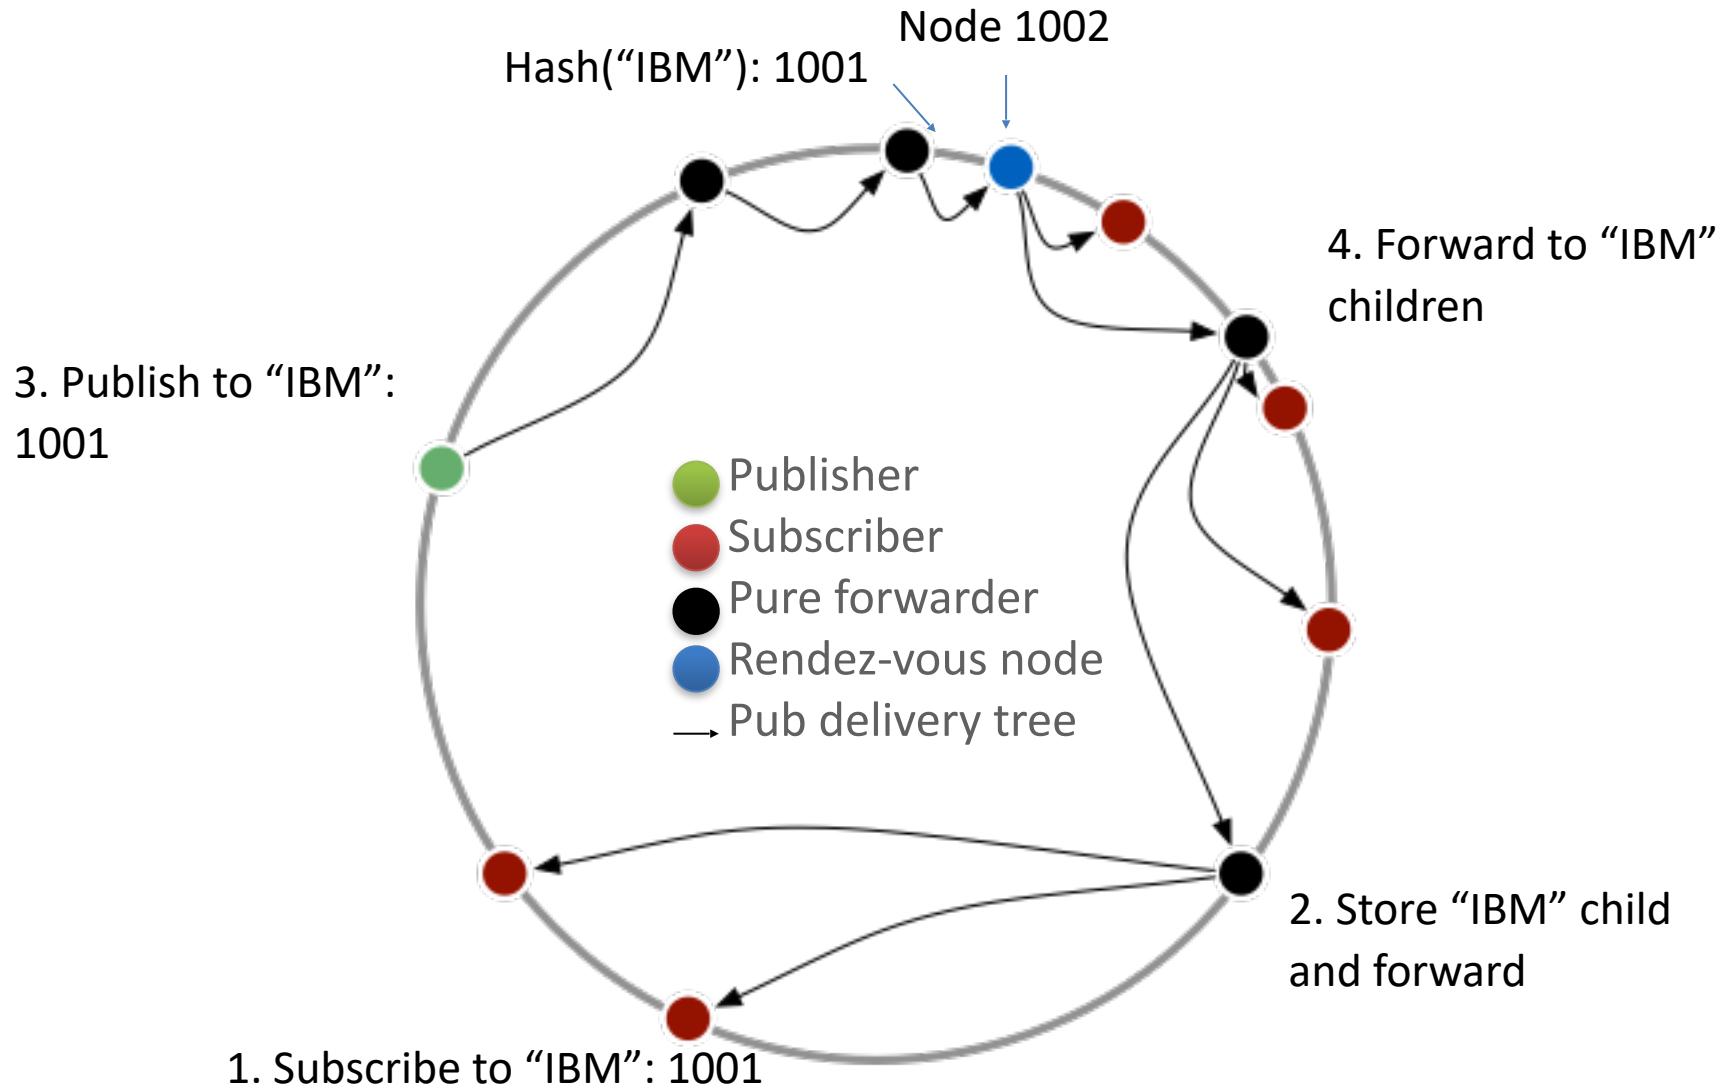
# Scribe Routing



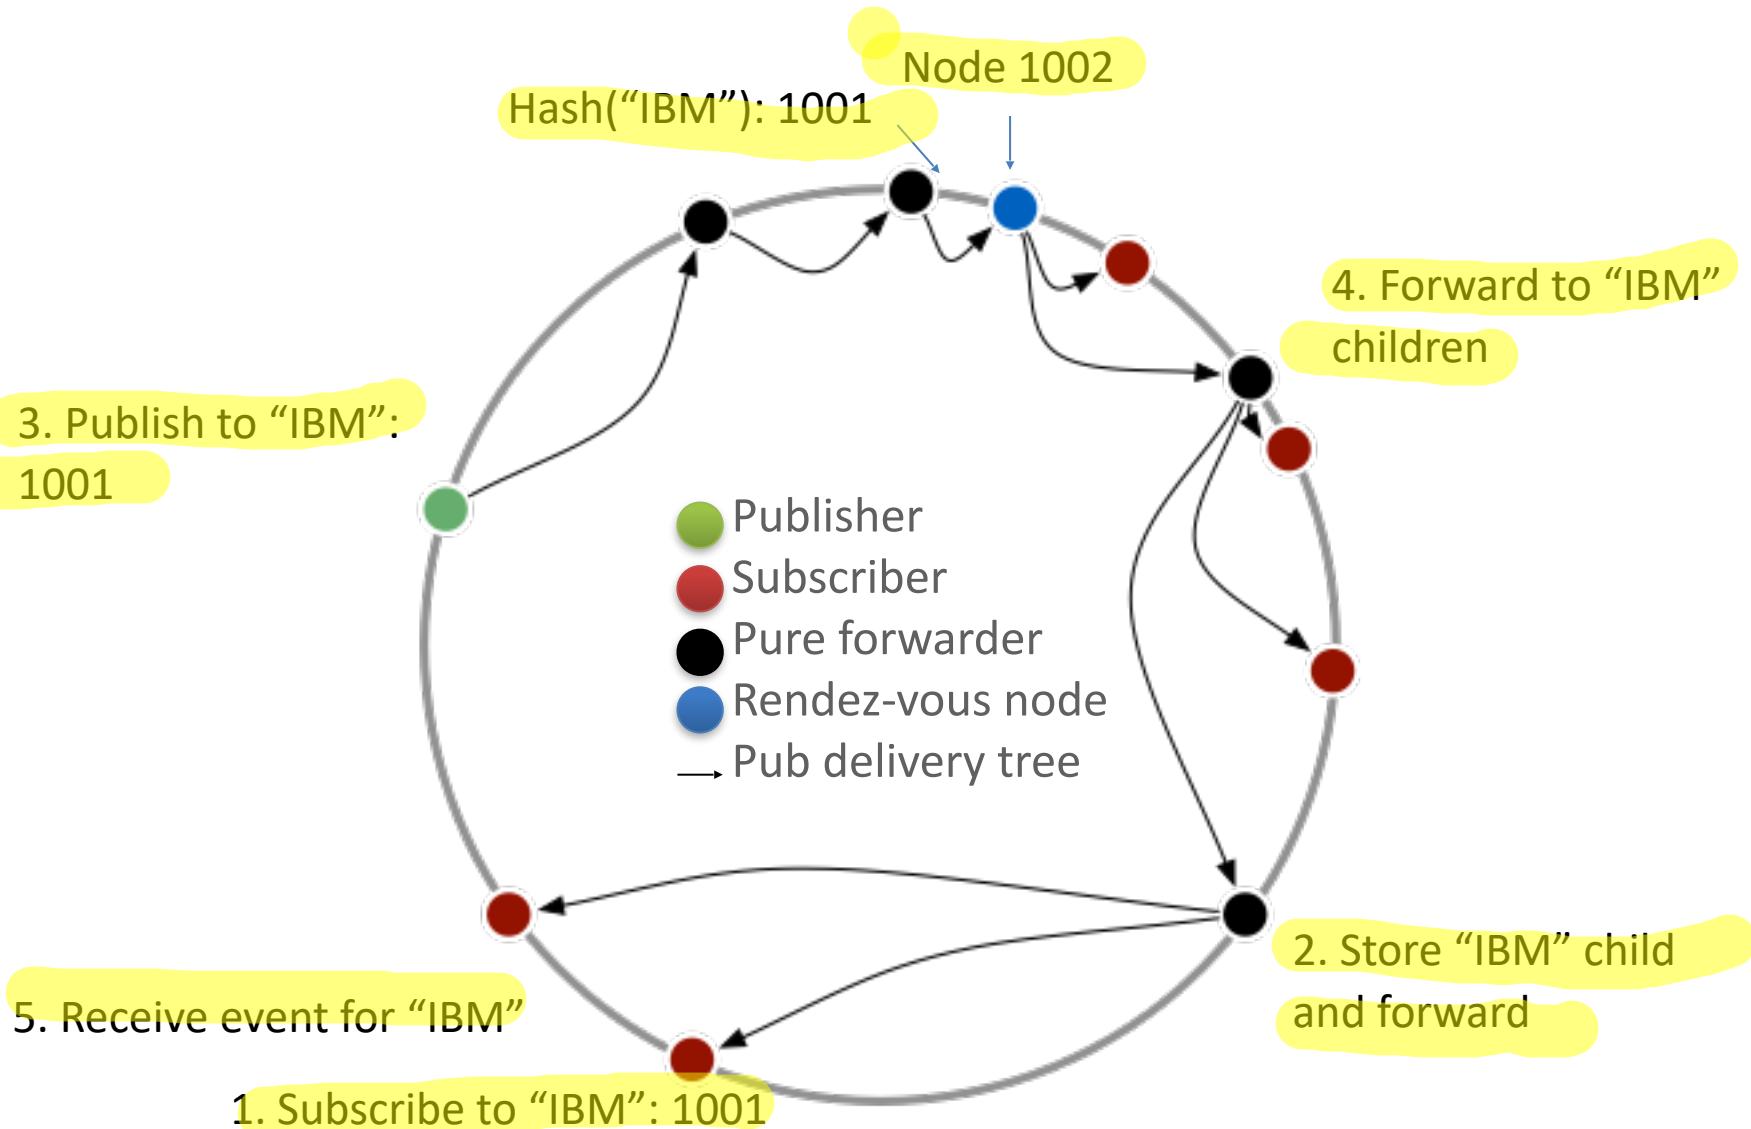
# Scribe Routing



# Scribe Routing



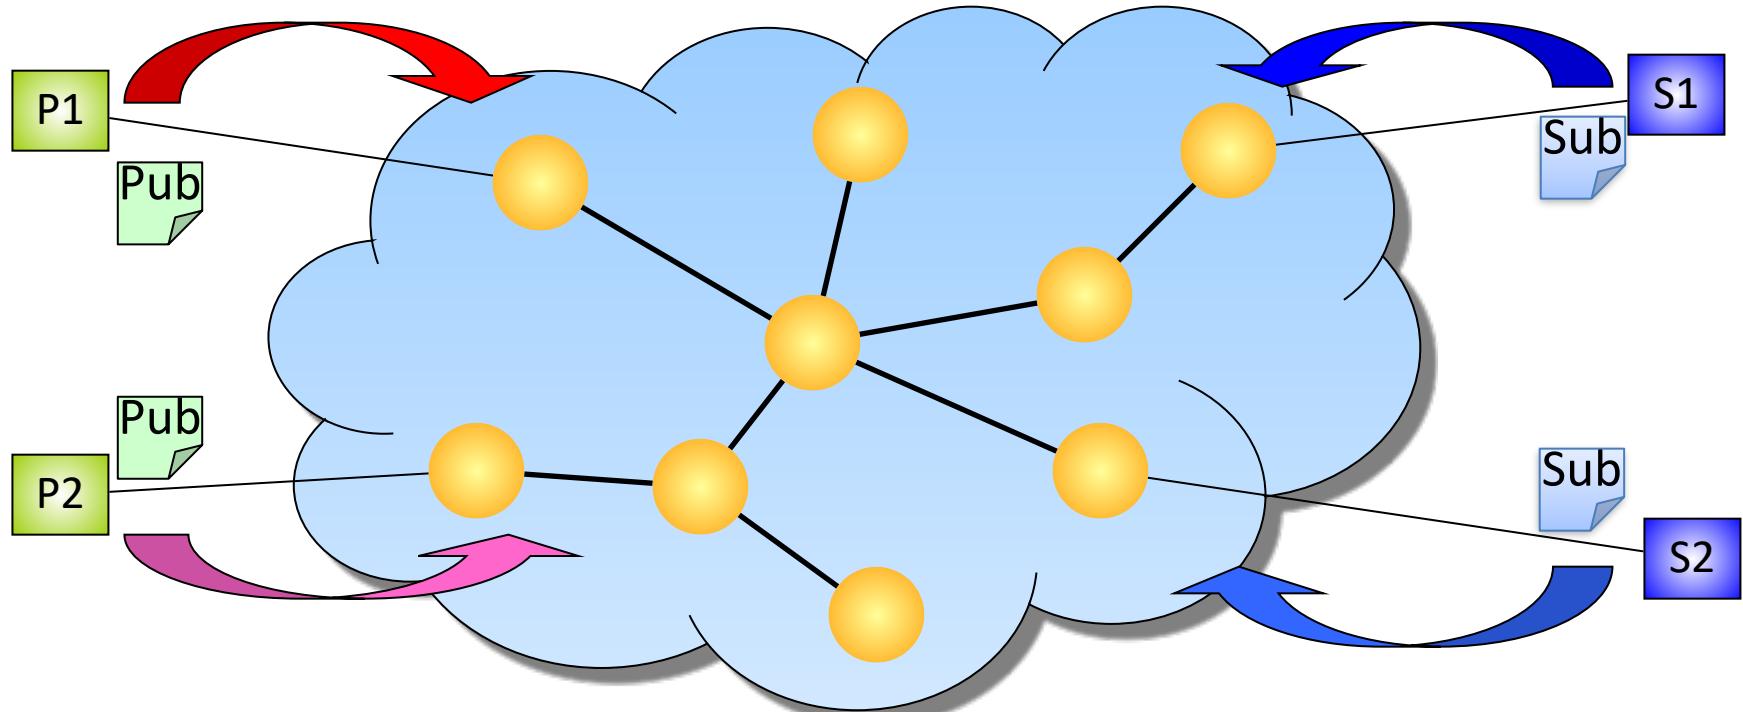
# Scribe Routing



# **OVERLAY-BASED ROUTING**

# Overlay Broker Networks

- Messages flow based on overlay links
- Clients connect to any edge broker

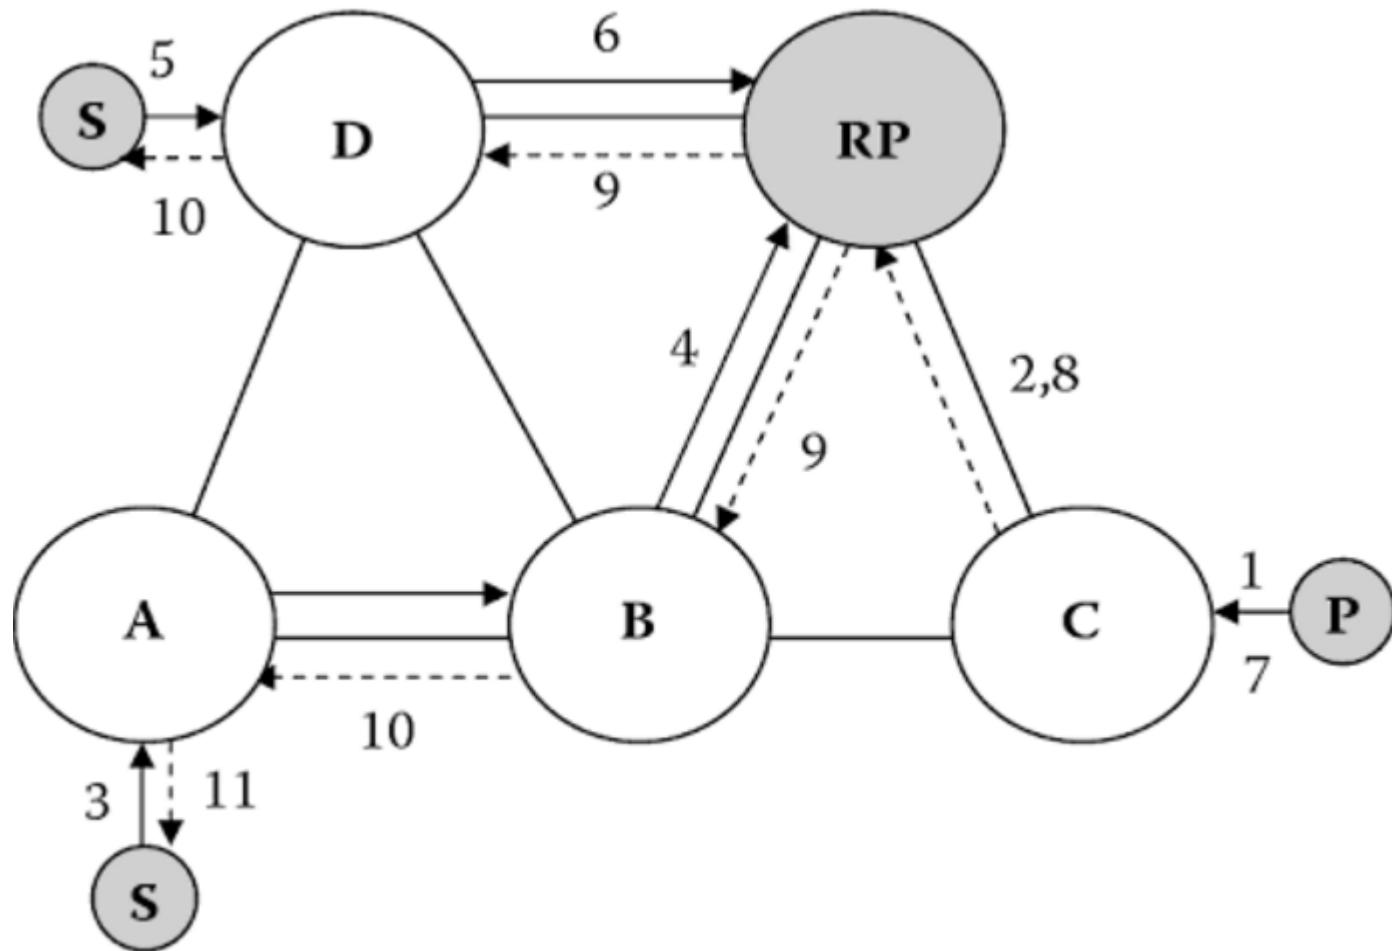


Matching and dissemination are performed inside the overlay

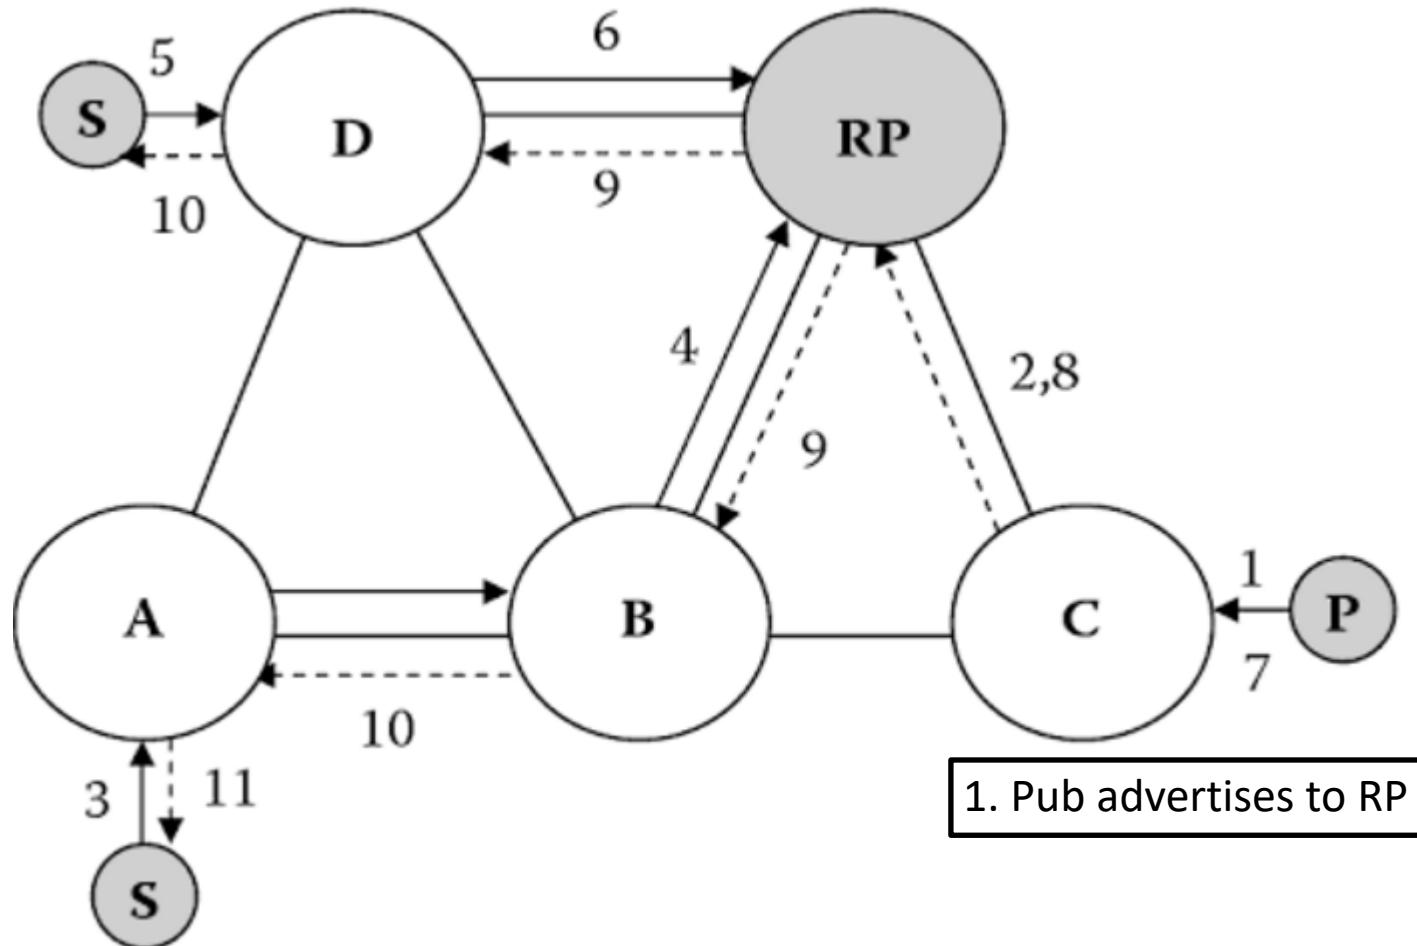
# **RDV-based Overlay: Hermes (2004)**

- Supports **content-based** matching with advertisements
- One broker is the **rendezvous point**
- Each broker knows how to reach the rendezvous point
- Rendezvous point performs **matching** and **subscriber notification**

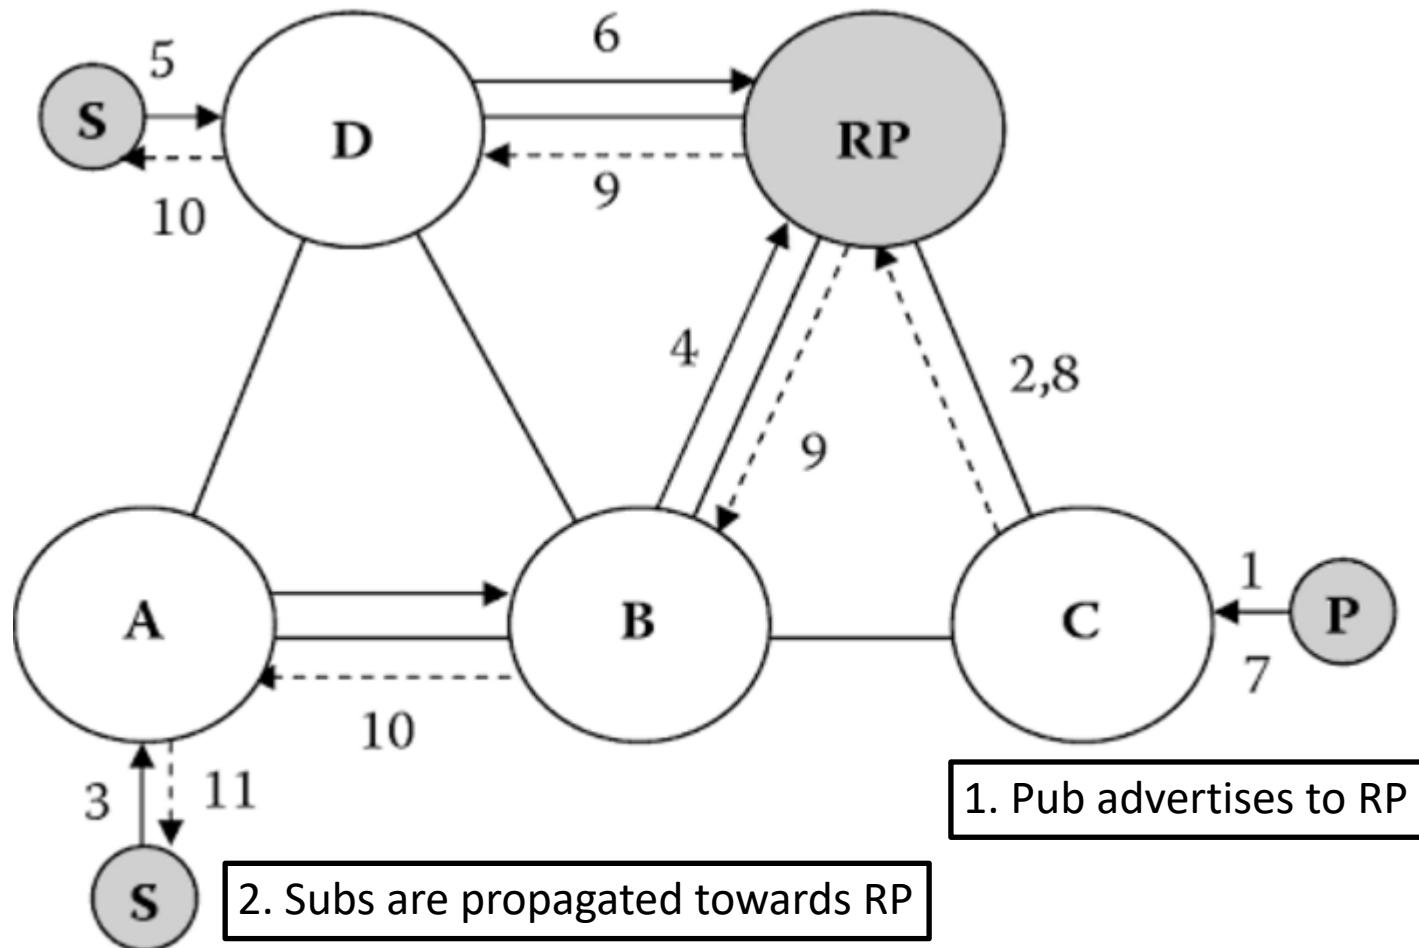
# Hermes Routing



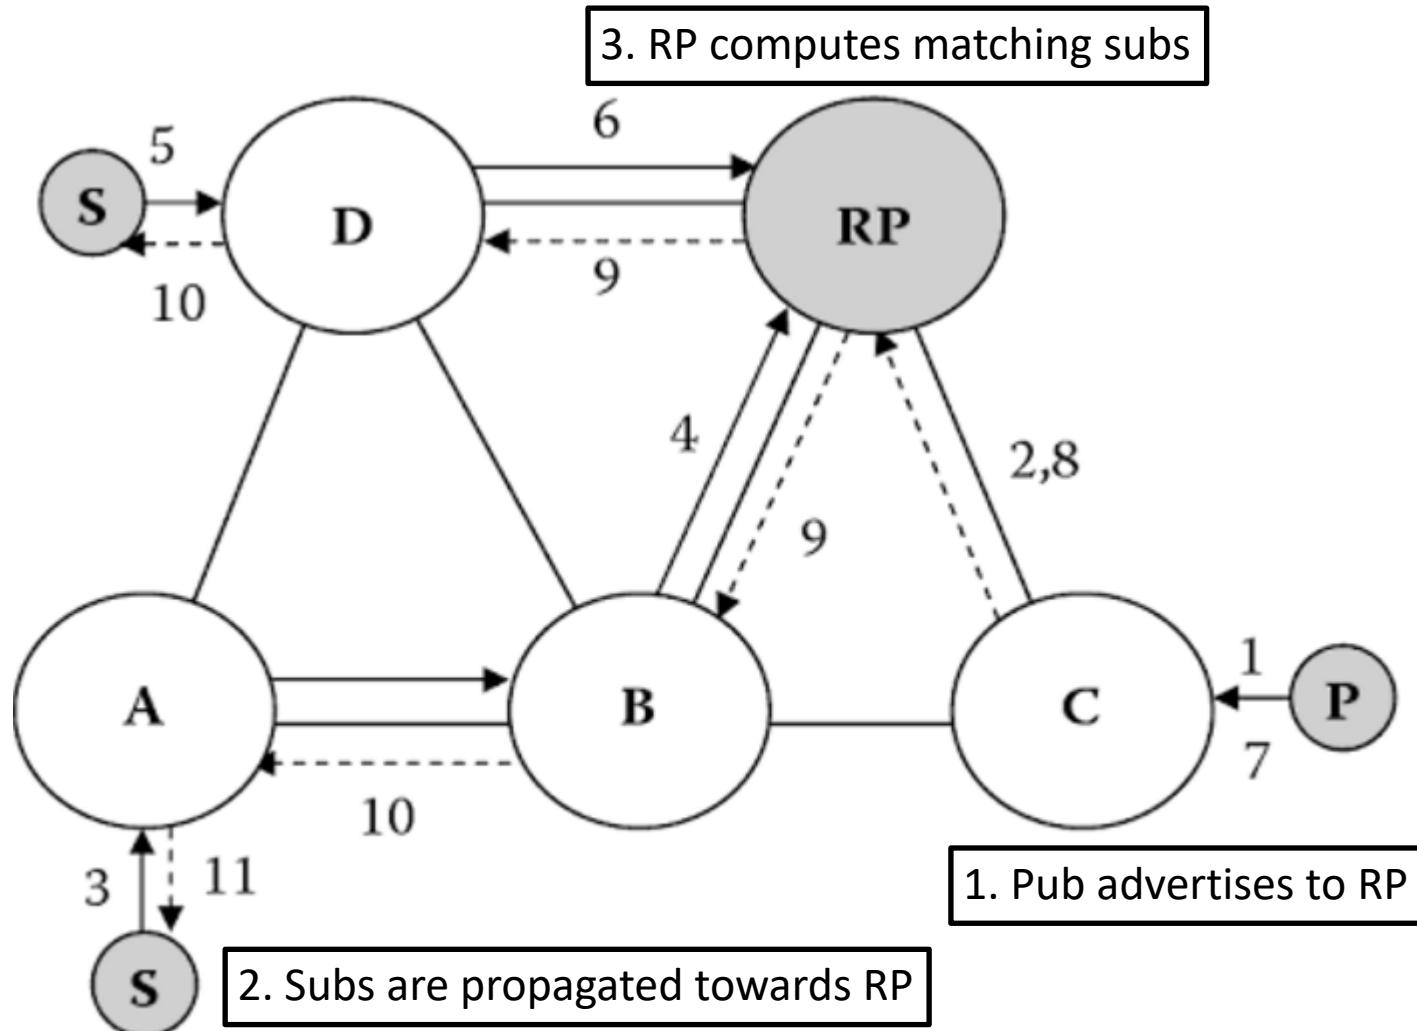
# Hermes Routing



# Hermes Routing



# Hermes Routing



# Possible Optimizations

- More than one rendezvous point
  - Use space partitioning
  - Increases state complexity of forwarding brokers
- Bloom filters using Link IDs
  - Reduces state complexity of forwarding brokers
  - Introduces false positive dissemination
  - Increases processing time at the RP

# Bloom Filter

- Efficient data structure for representing whether elements are in a set
- Represented as a **bit vector** with **fixed size**
- Elements can be inserted into the Bloom filter (the set)
- Element can be tested to identify whether it is in the set
- Cannot obtain a list of all elements in the Bloom filter
- **False positives** (elements that would show as members but are not) are possible – i.e., can incorrectly identify an element as part of the set
- **False negatives** (elements that are in the set but don't show as such) **are not possible**

Based on  $k$   
hash functions.

## Bloom Filter Operations

Start with an  $m$  bit array, filled with 0s.

$B$

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Hash each item  $x_j$  in set  $S$ ,  $k$  times. If  $H_i(x_j) = a$ , set  $B[a] = 1$

$B$

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

To check if  $y$  is in  $S$ , check  $B$  at  $H_i(y)$ . All  $k H_i(y)$  must be 1

$B$

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Possible to have a *false positive*; all  $k$  values are 1, but  $y$  is not in  $S$ !

$B$

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

# Example of Bloom Filter

Start with a **16** bit array, filled with 0s.

**B**

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

$\text{Hash1}(\text{"id1"}) = 7, \text{Hash2}(\text{"id1"}) = 2, \dots, \text{Hash8}(\text{"id1"}) = 5$

**B**

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

$\text{Hash1}(\text{"id2"}) = 7, \text{Hash2}(\text{"id2"}) = 11, \text{Hash3}(\text{"id2"}) = 4$

**B**

0	1	0	1	1	0	1	0	0	1	1	1	0	1	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

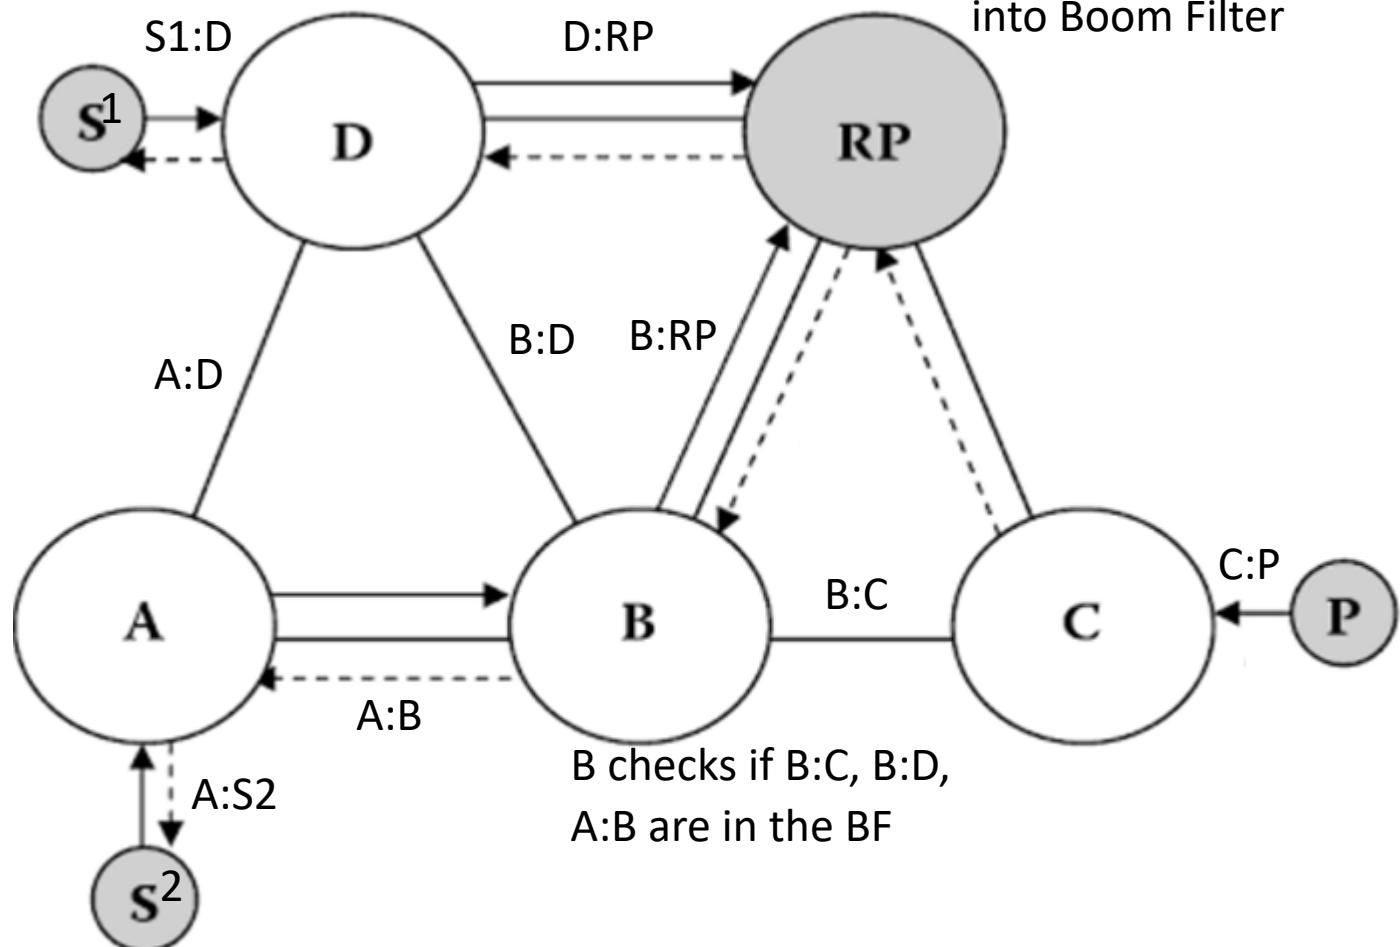
$\text{Hash1}(\text{"id3"}) = 5, \text{Hash2}(\text{"id3"}) = 5, \dots, \text{Hash8}(\text{"id3"}) = 7$

**B**

0	1	0	0	1	0	1	0	0	1	1	1	0	1	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

**8** hash functions

# Link IDs encoding in Bloom Filters (LIPSIN 2009)



# Filtering-Based Routing: PADRES (2003)

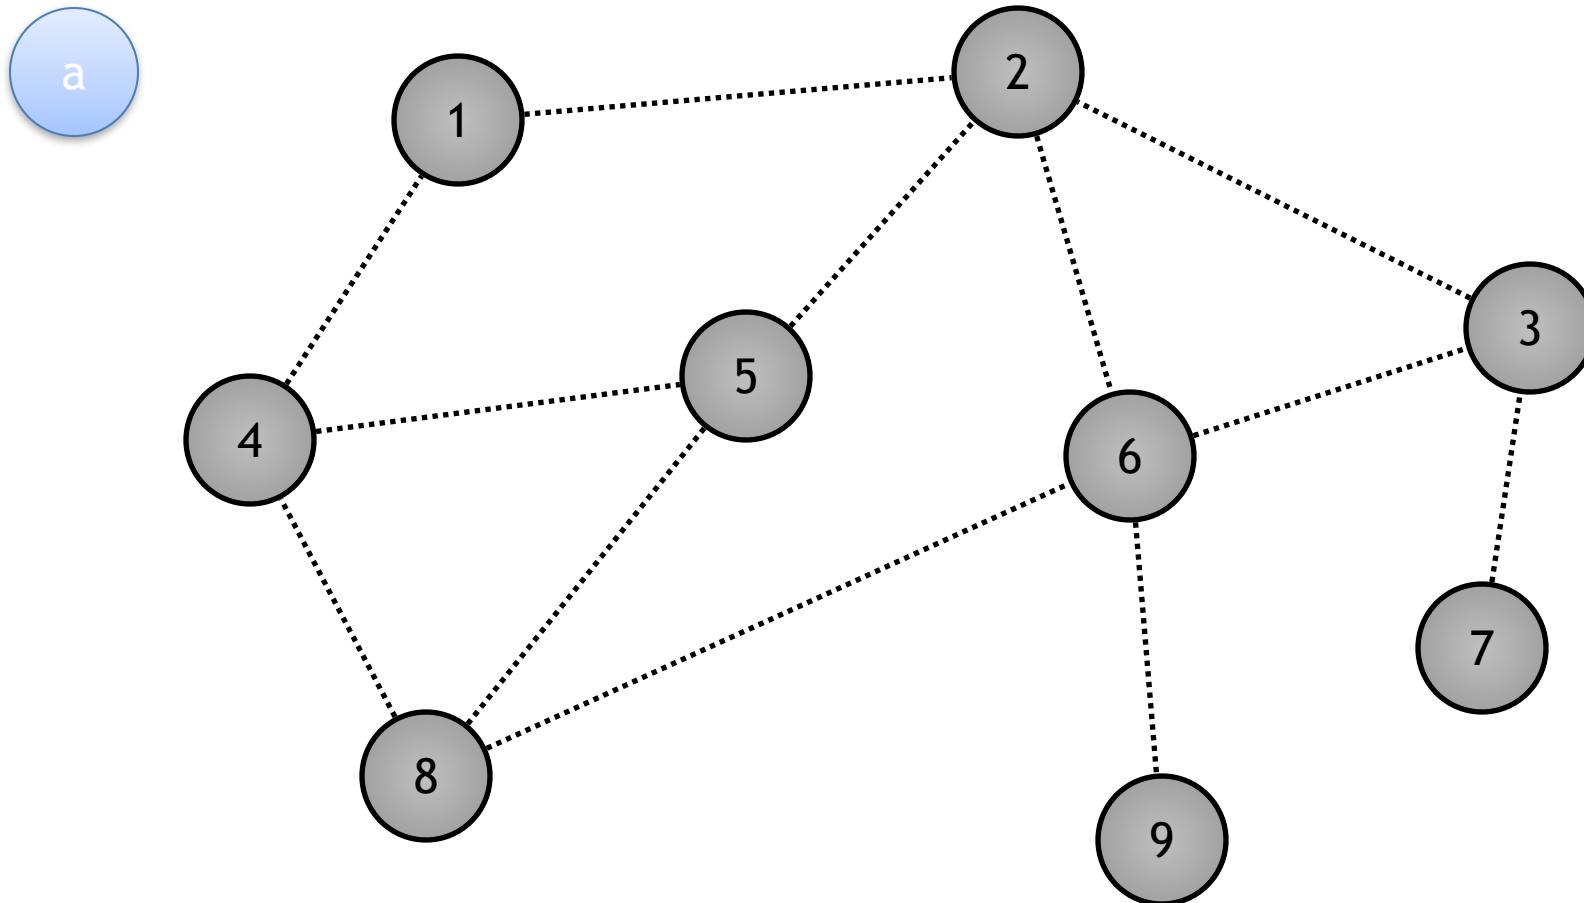
- Each broker performs filtering
- Based on the filtering, each broker disseminates publications towards matching subs
- The next hops repeat the process until subs are reached
- Supports any matching model (e.g. content-based)
- Subs are flooded to initialize routing paths
- Optimization:
  - Subscription covering
  - Advertising-based routing



<http://padres.msrg.org>

# Filtering-Based Routing

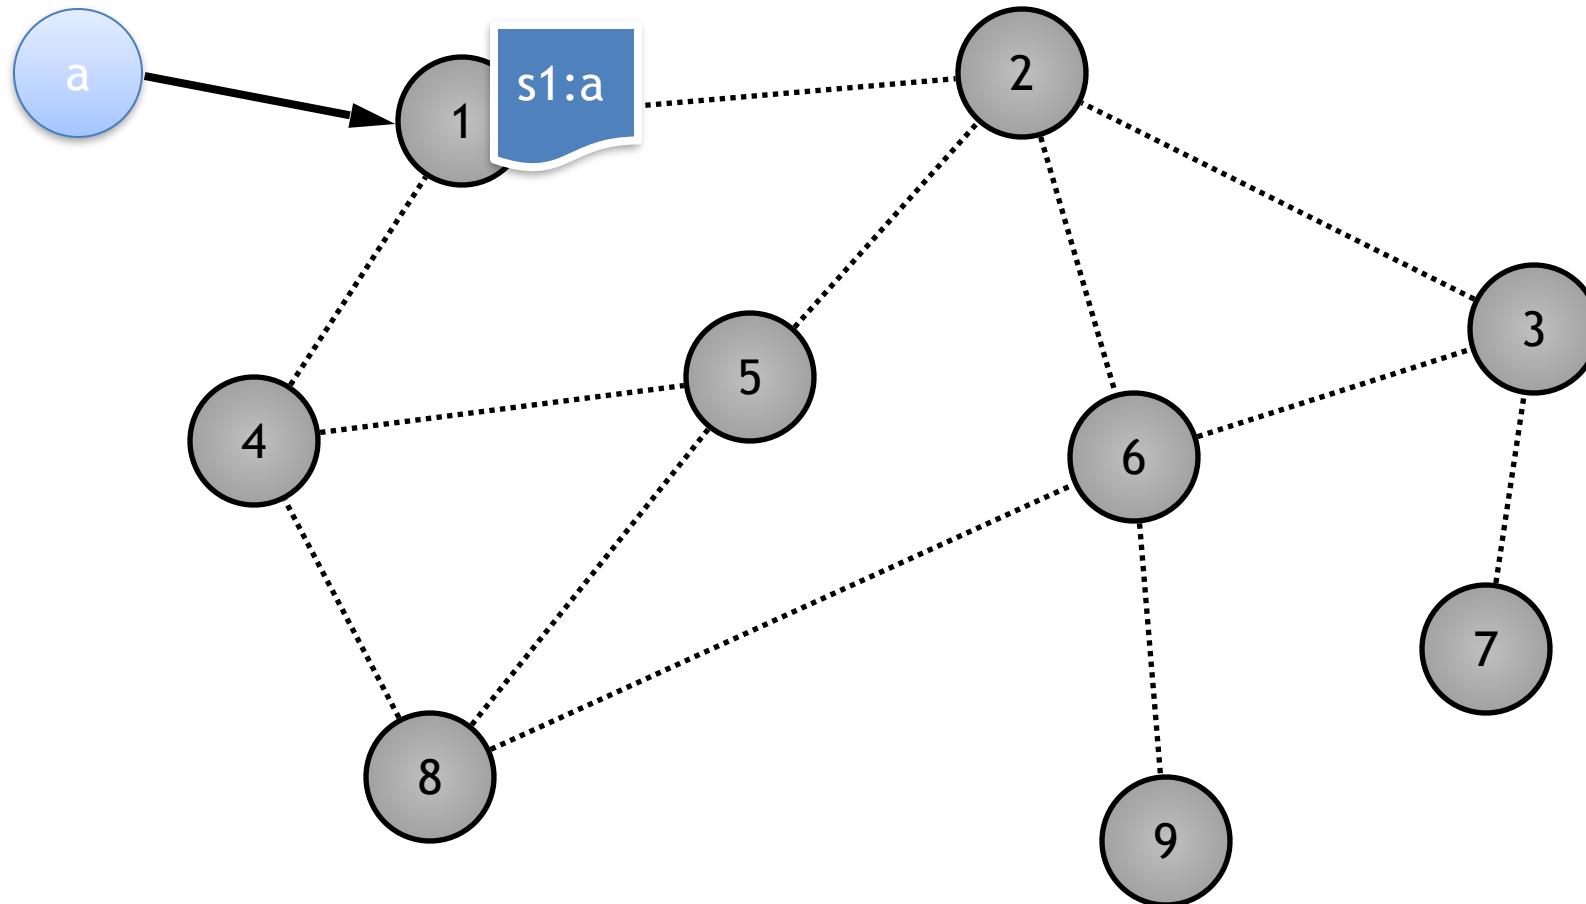
## Subscription Flooding



# Filtering-Based Routing

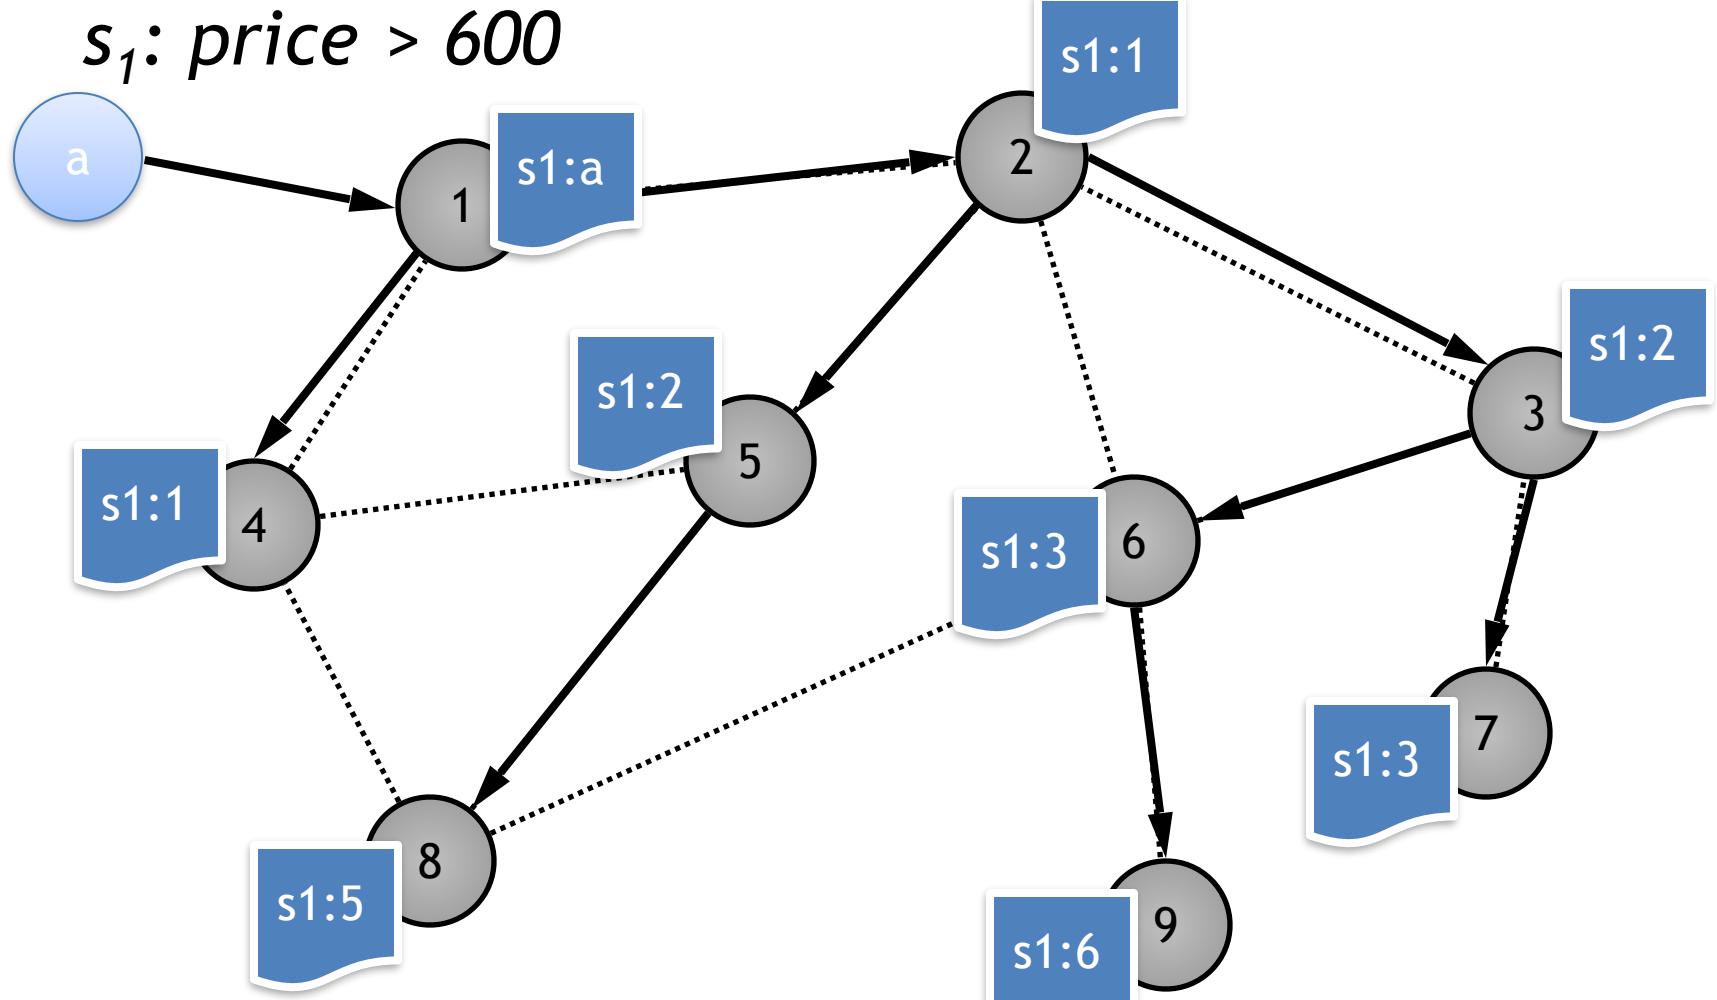
## Subscription Flooding

$s_1: \text{price} > 600$



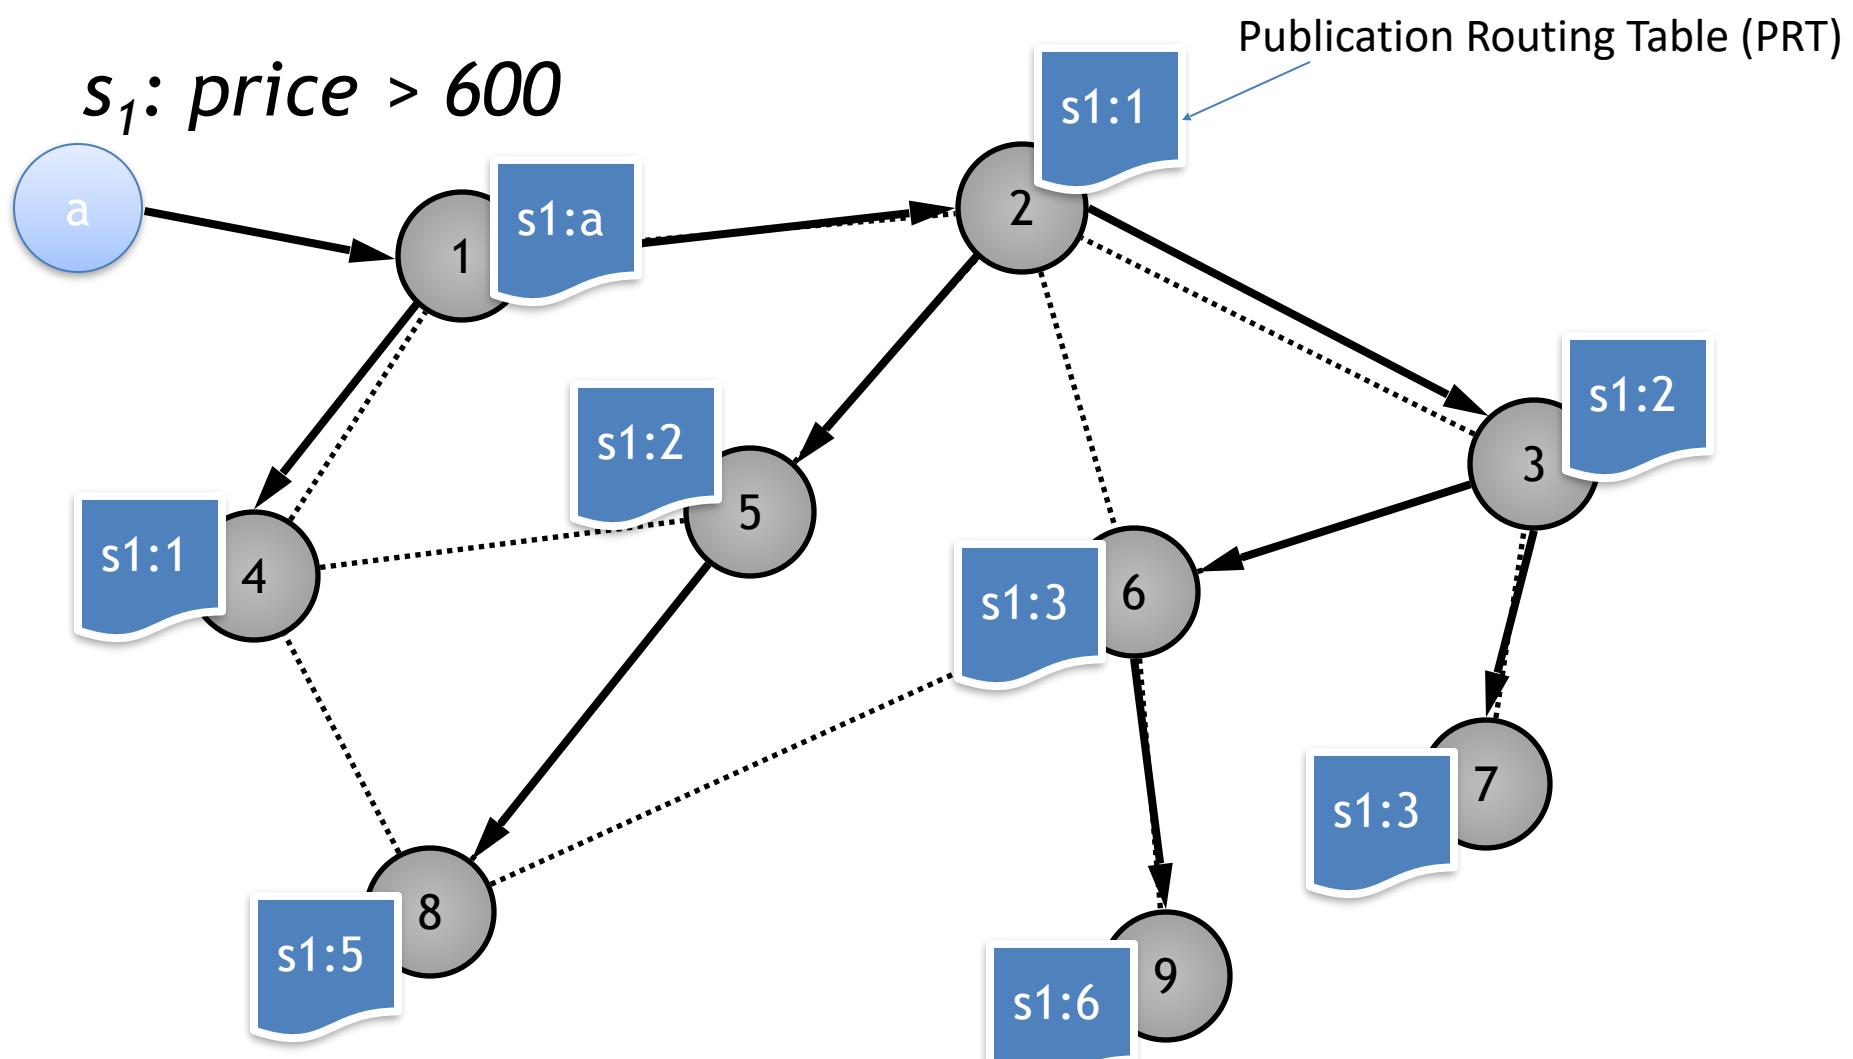
# Filtering-Based Routing

## Subscription Flooding



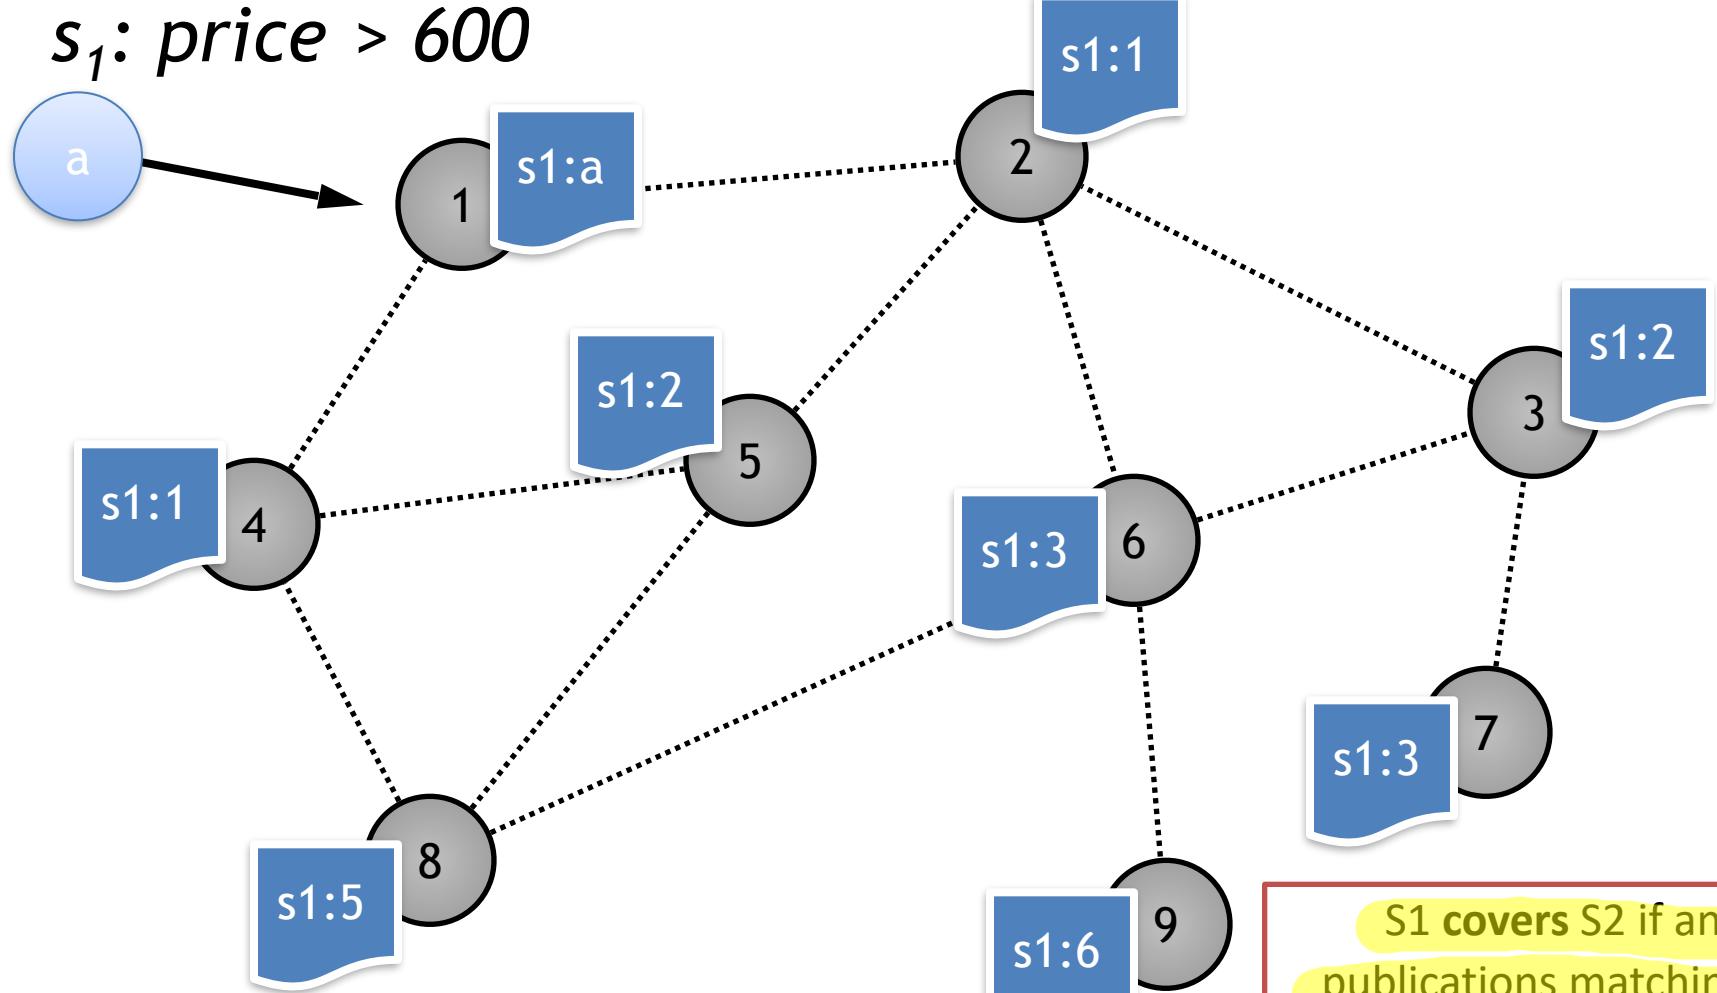
# Filtering-Based Routing

## Subscription Flooding



# Filtering-Based Routing

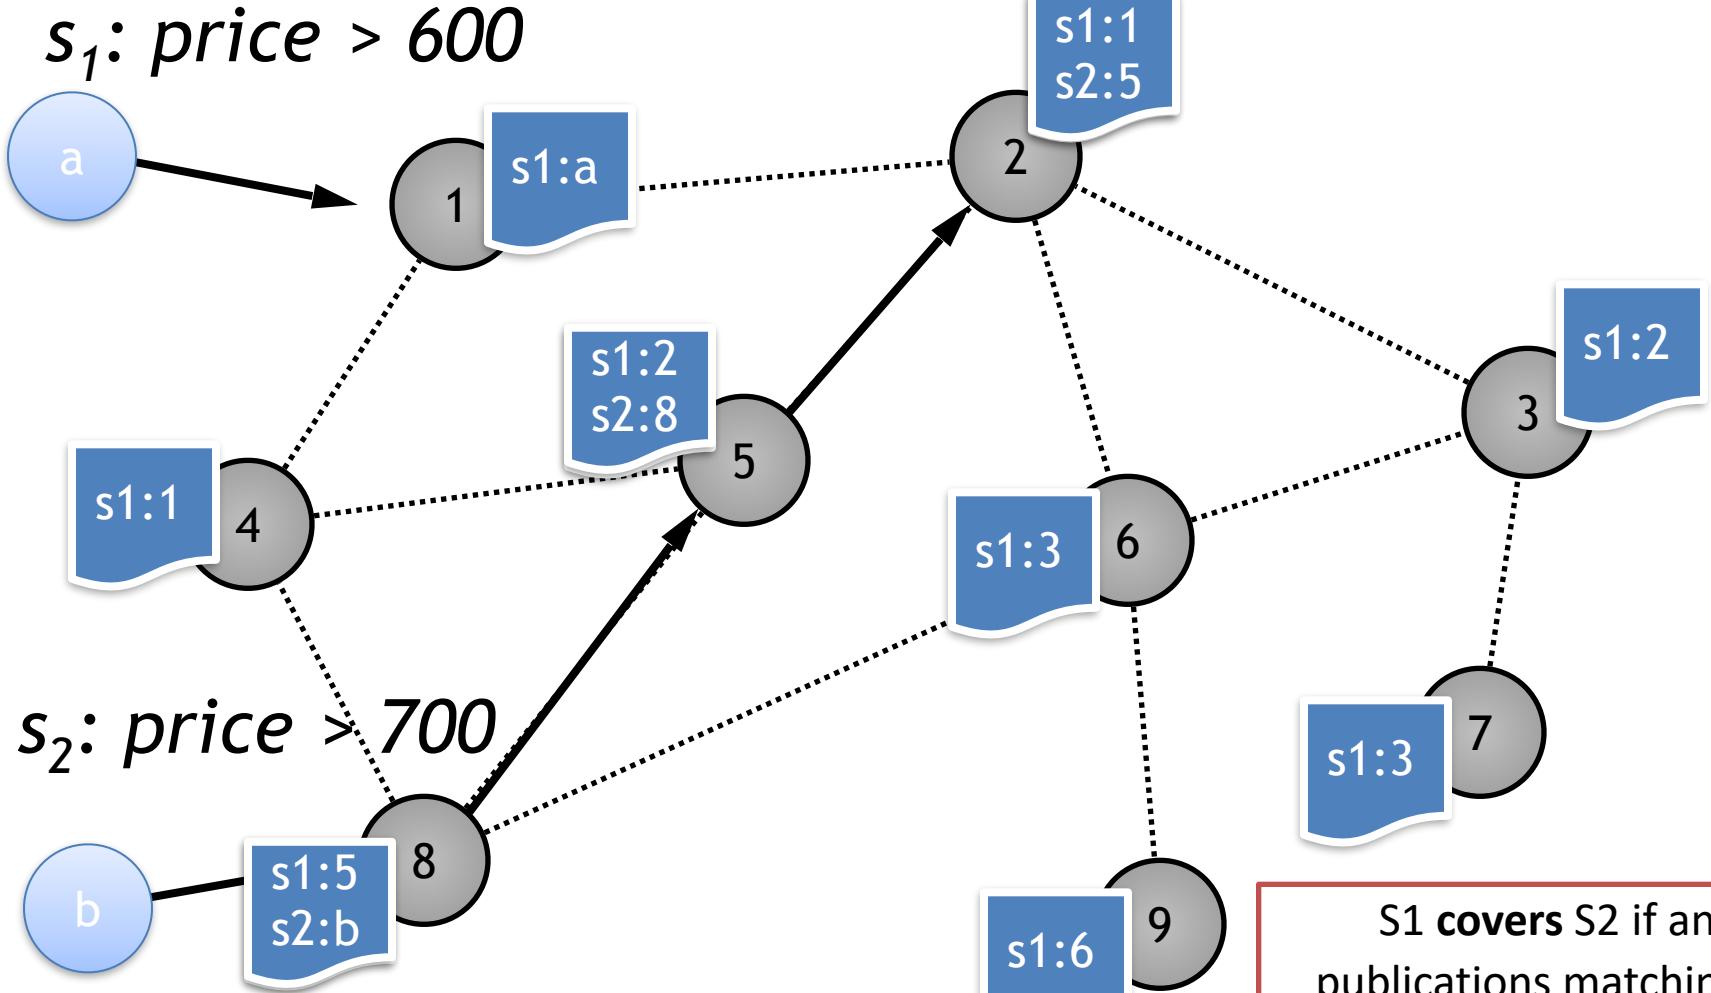
## Subscription Covering



S1 covers S2 if any publications matching S2 also matches S1

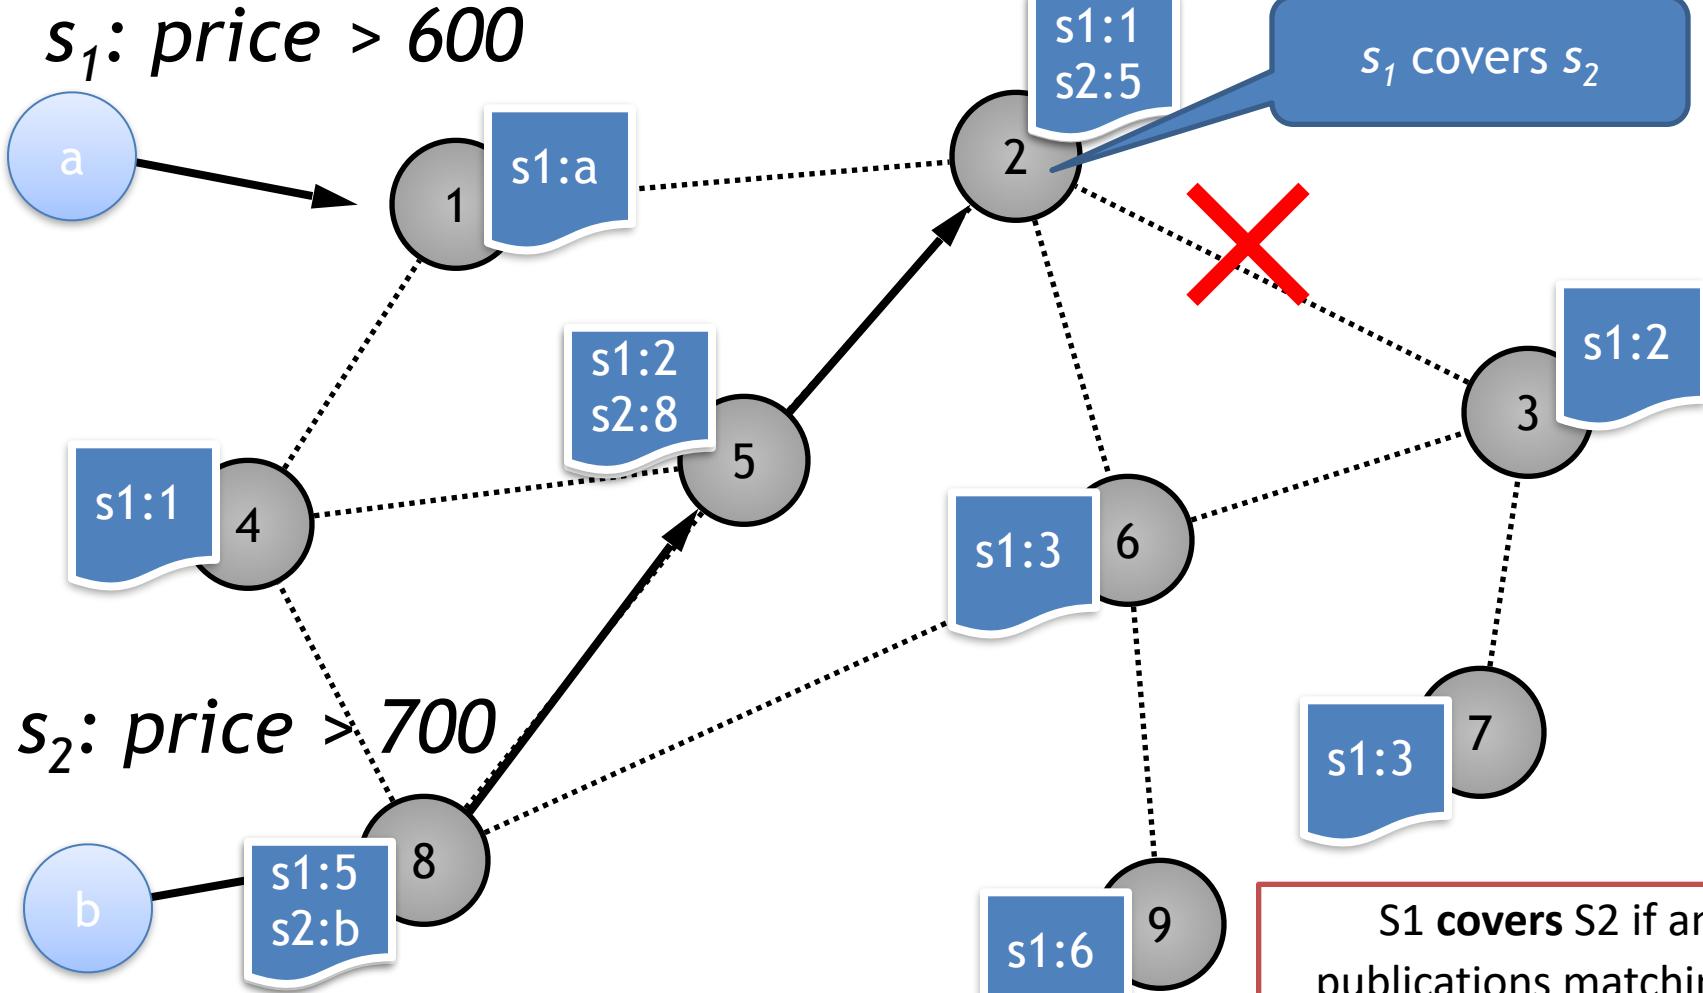
# Filtering-Based Routing

## Subscription Covering



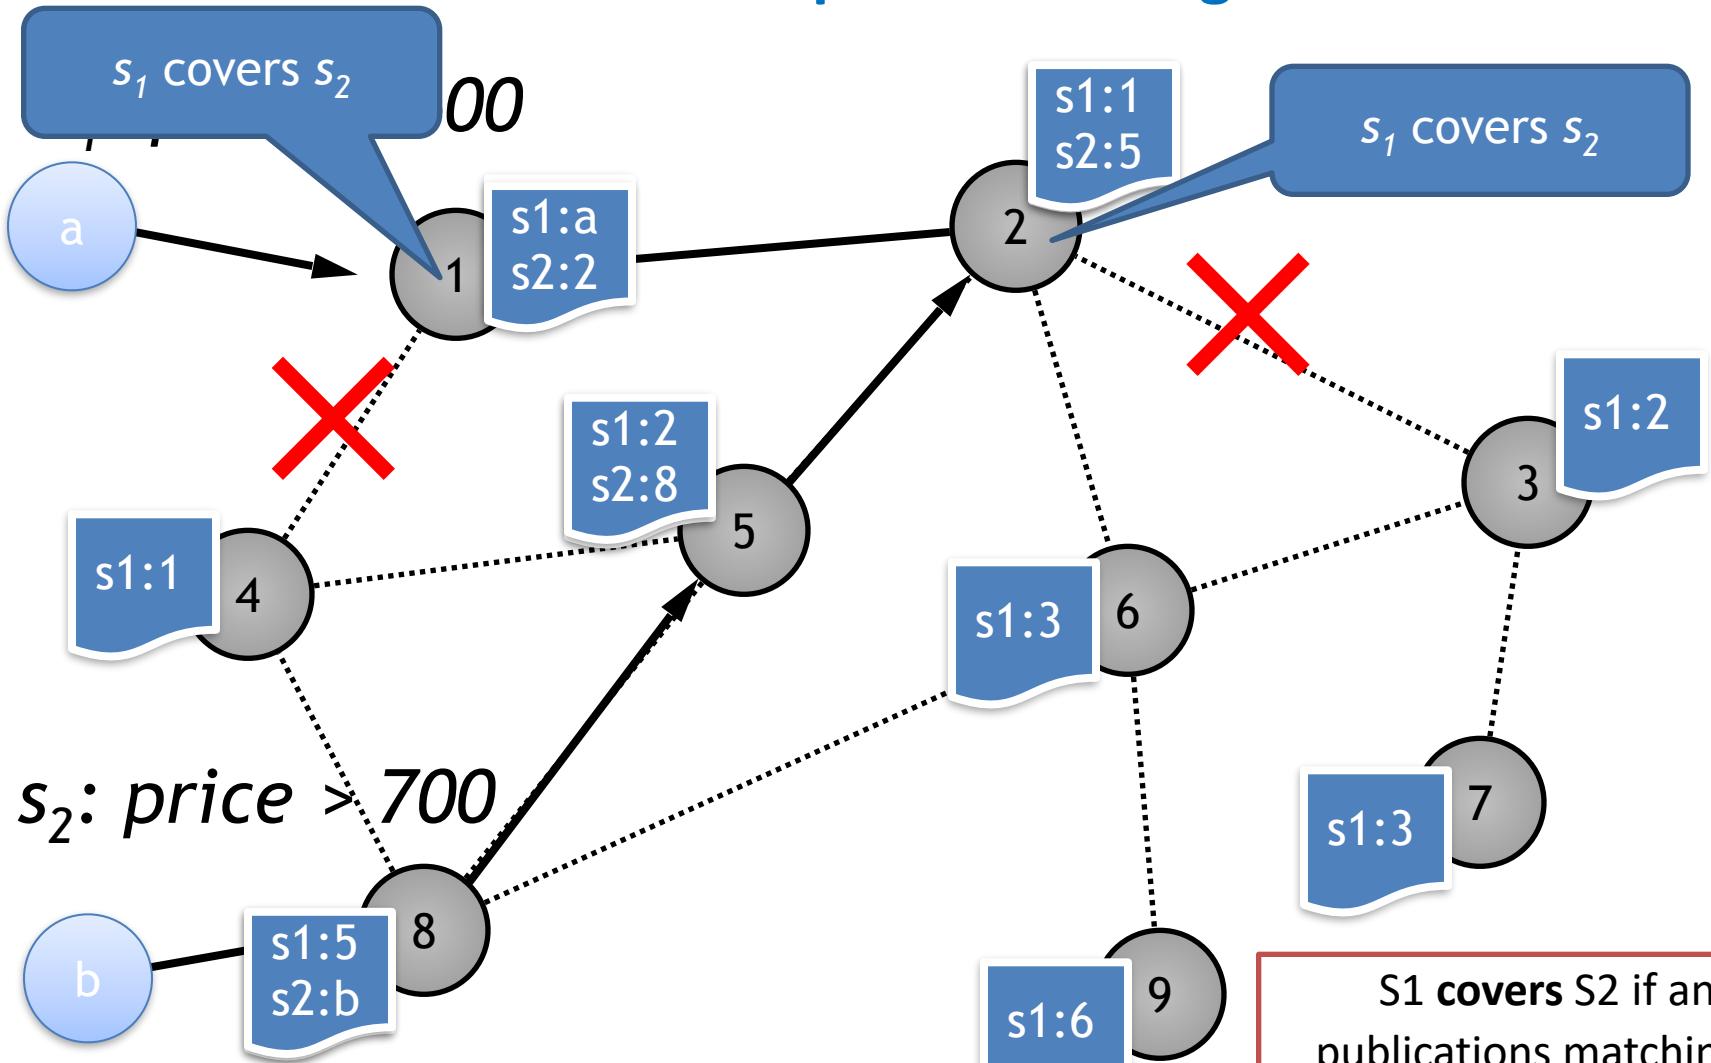
# Filtering-Based Routing

## Subscription Covering



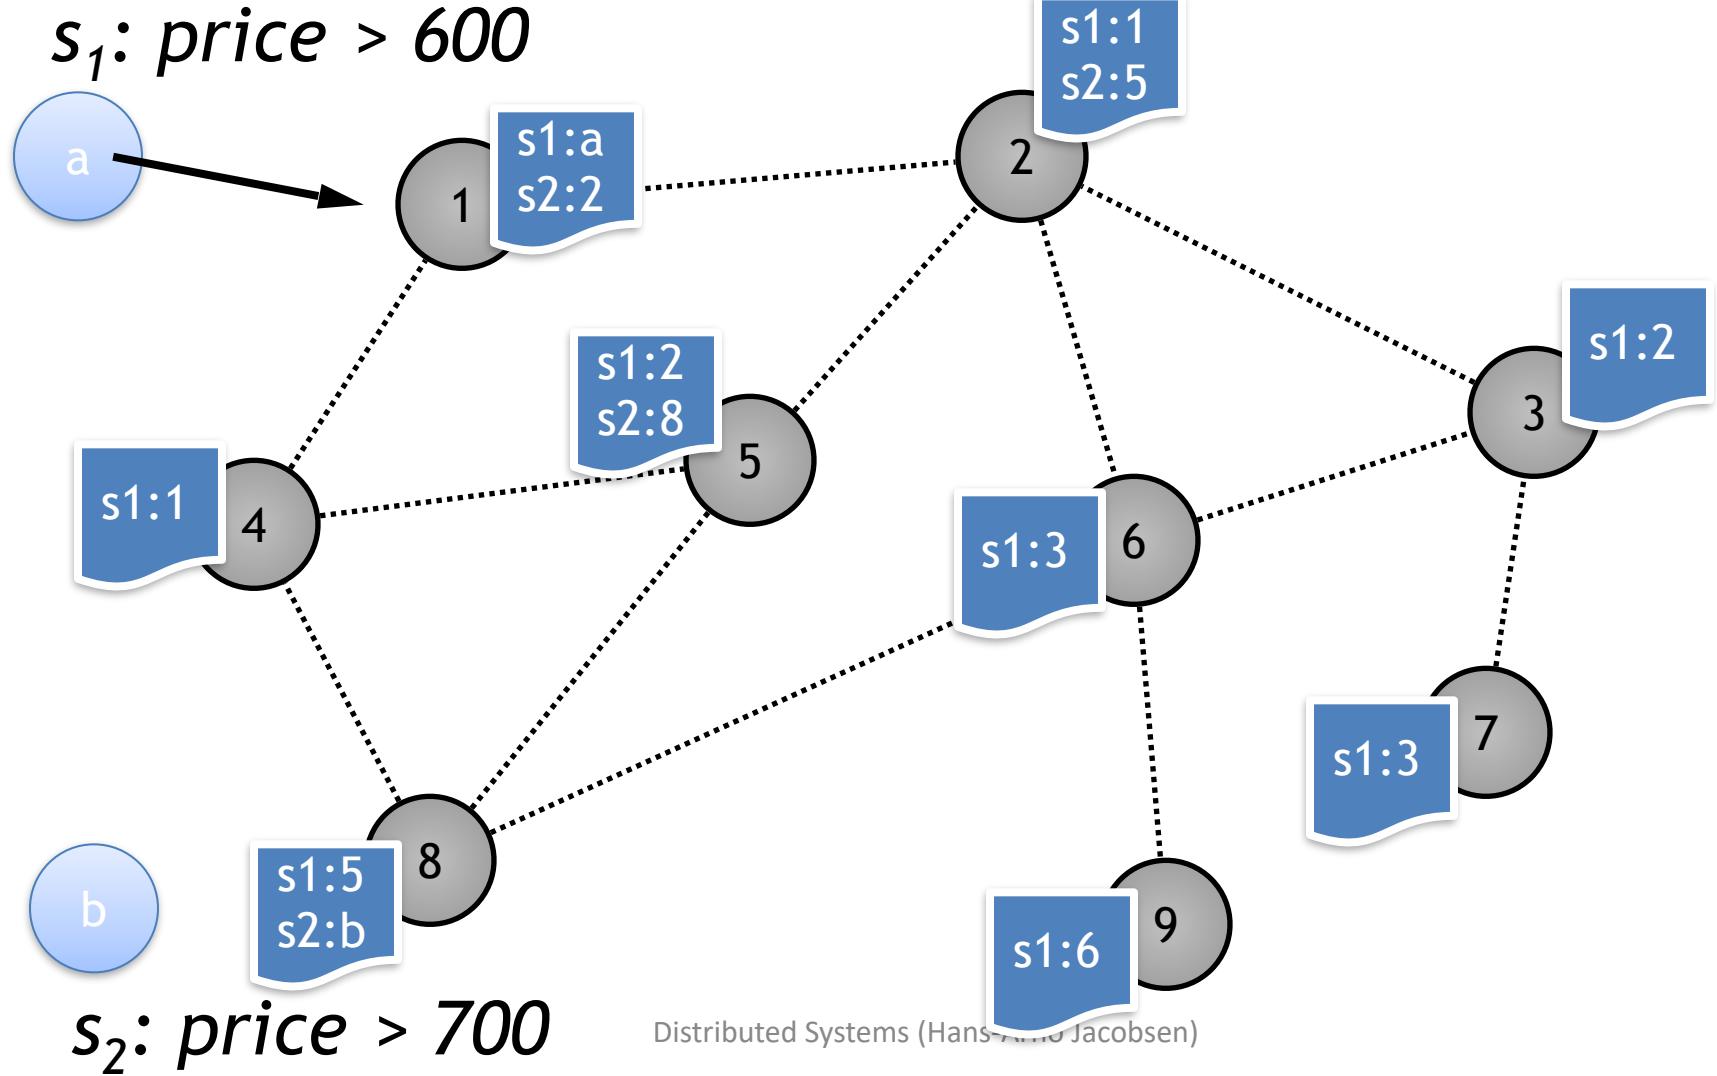
# Filtering-Based Routing

## Subscription Covering



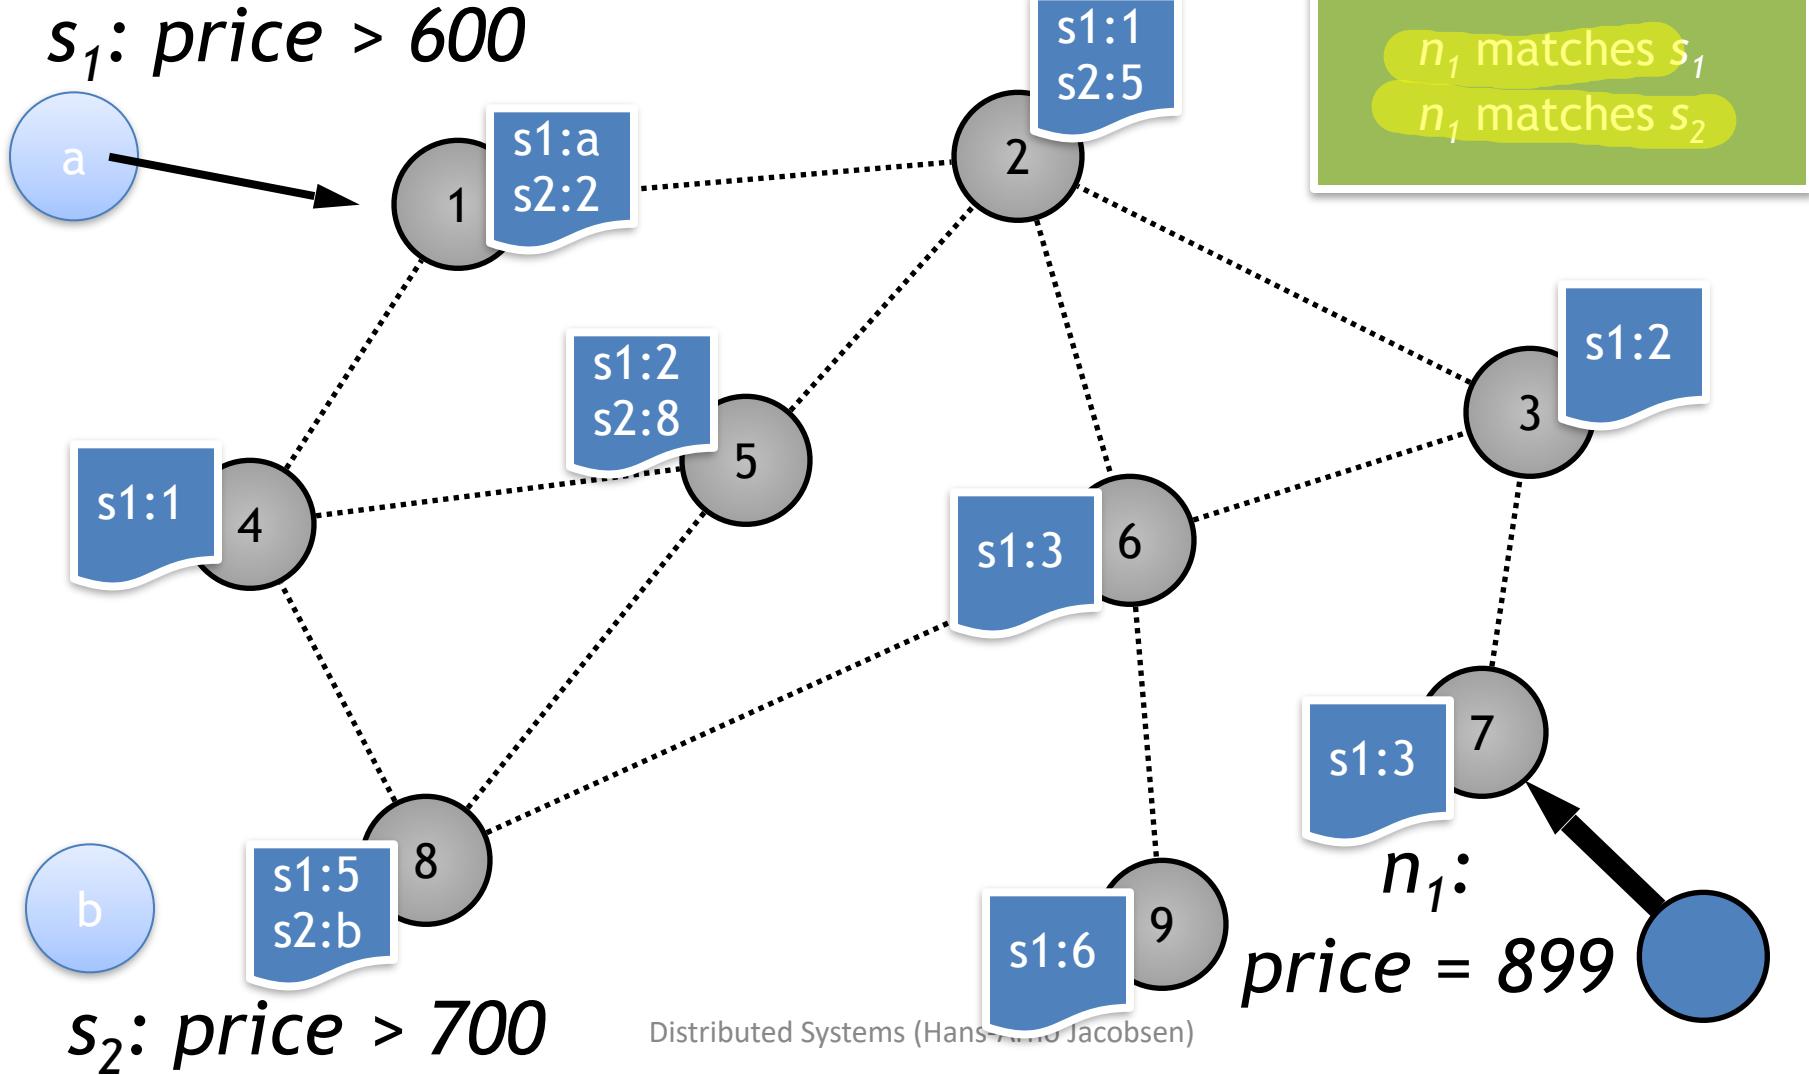
# Filtering-Based Routing

## Notification Delivery



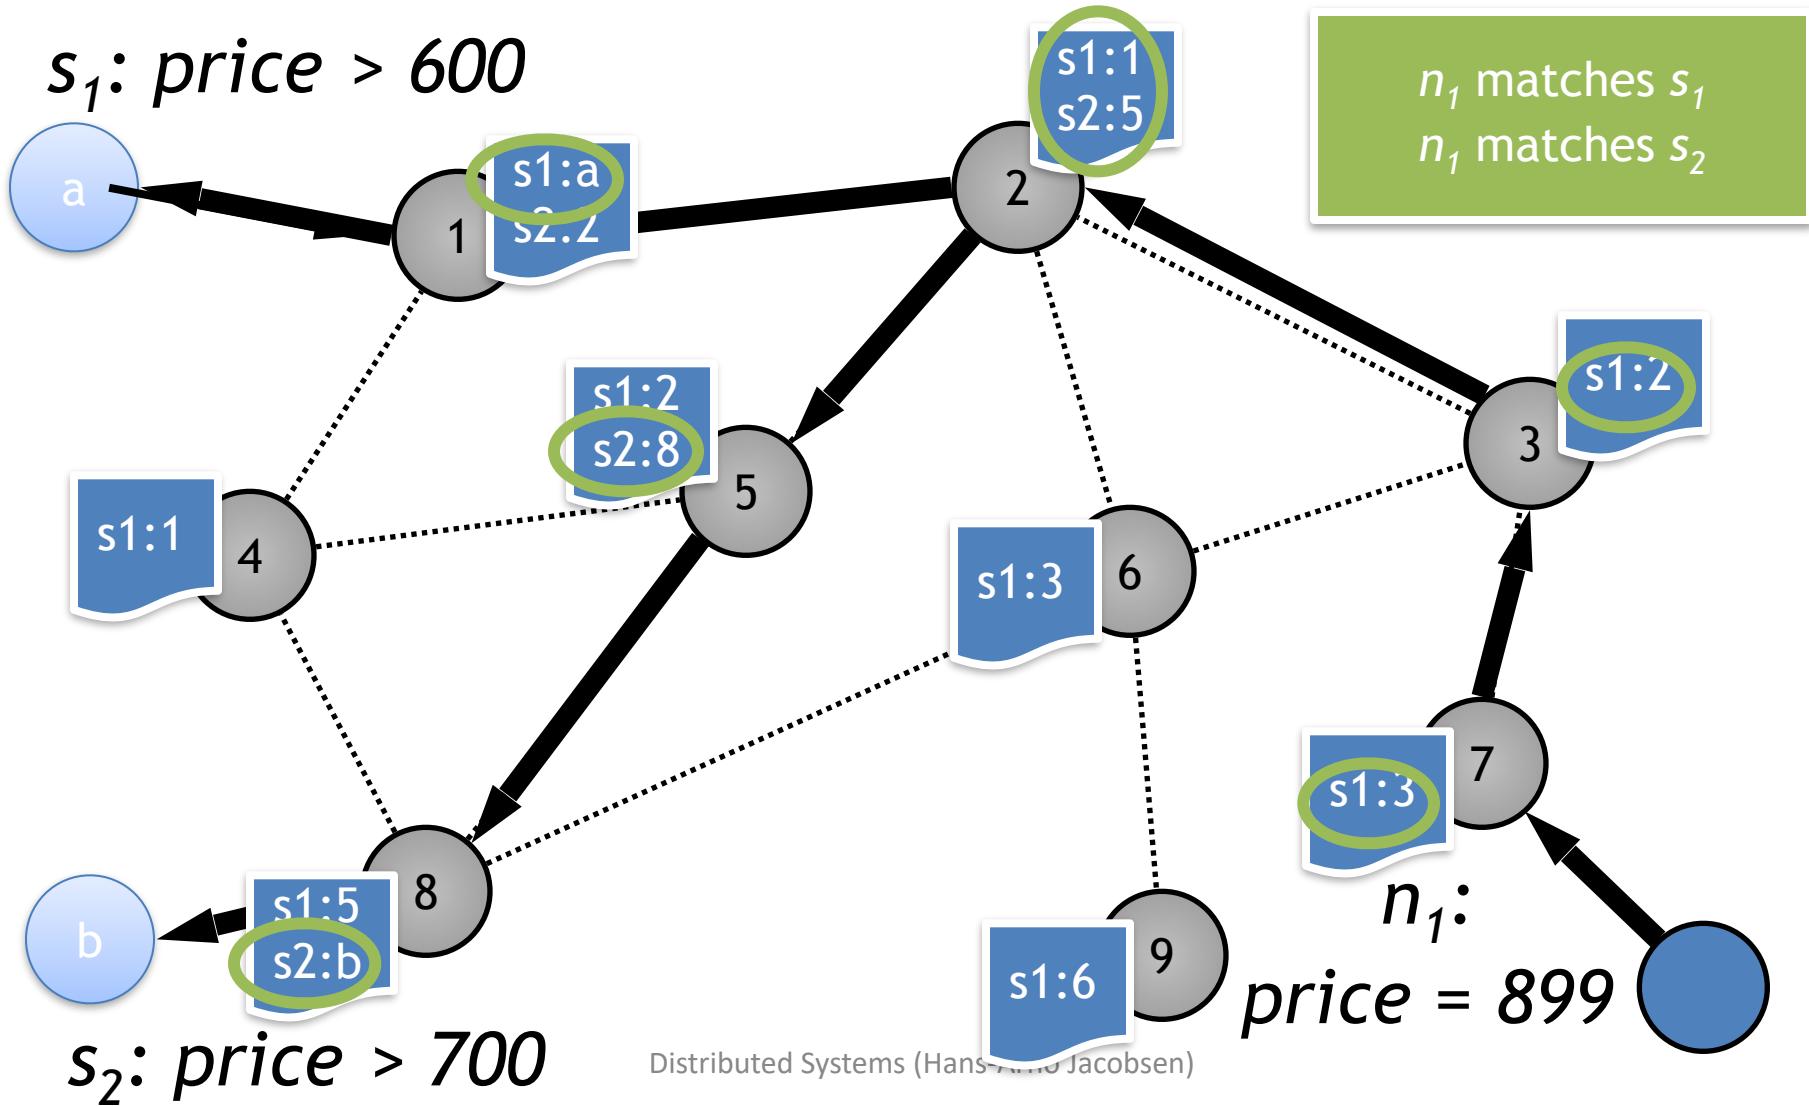
# Filtering-Based Routing

## Notification Delivery



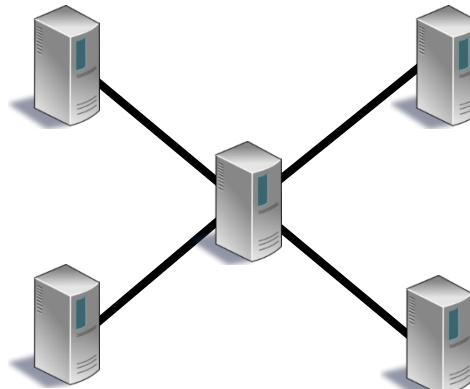
# Filtering-Based Routing

## Notification Delivery

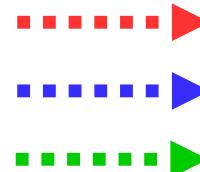


# Advertisement-Based Routing

Publishers must first *advertise* the type of publications they are going to make.  
Advertisements are stored at each broker in a *Subscription Routing Table (SRT)*.



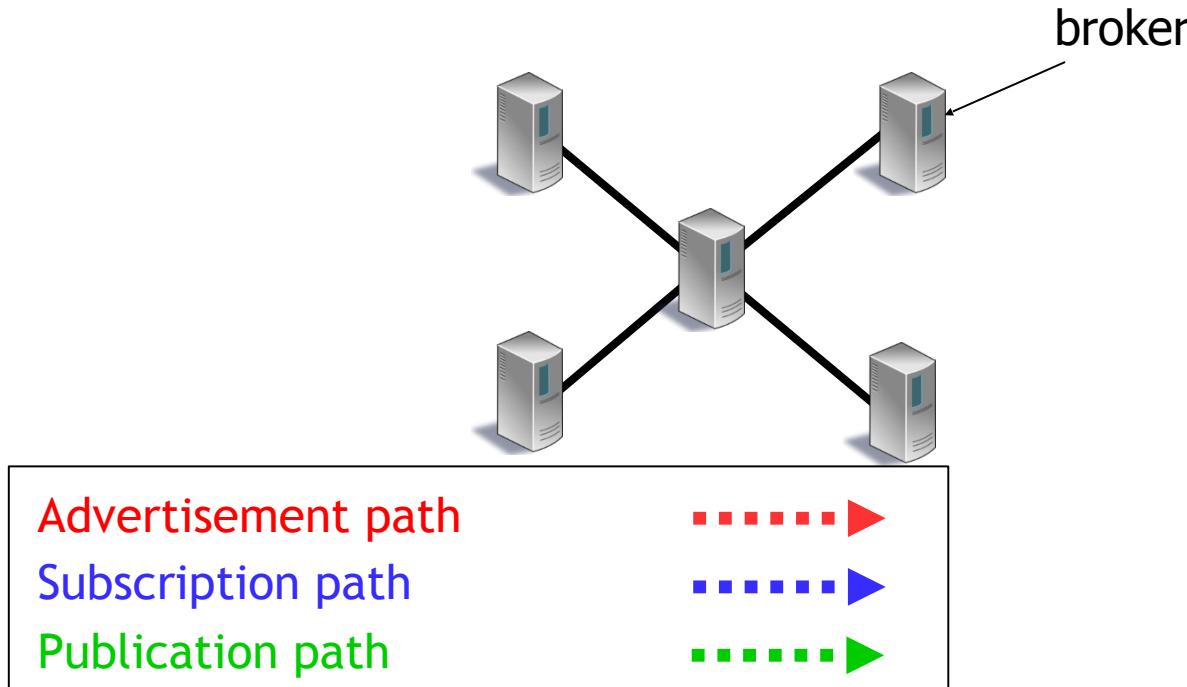
Advertisement path  
Subscription path  
Publication path



Advertisements are flooded, but not subscriptions

# Advertisement-Based Routing

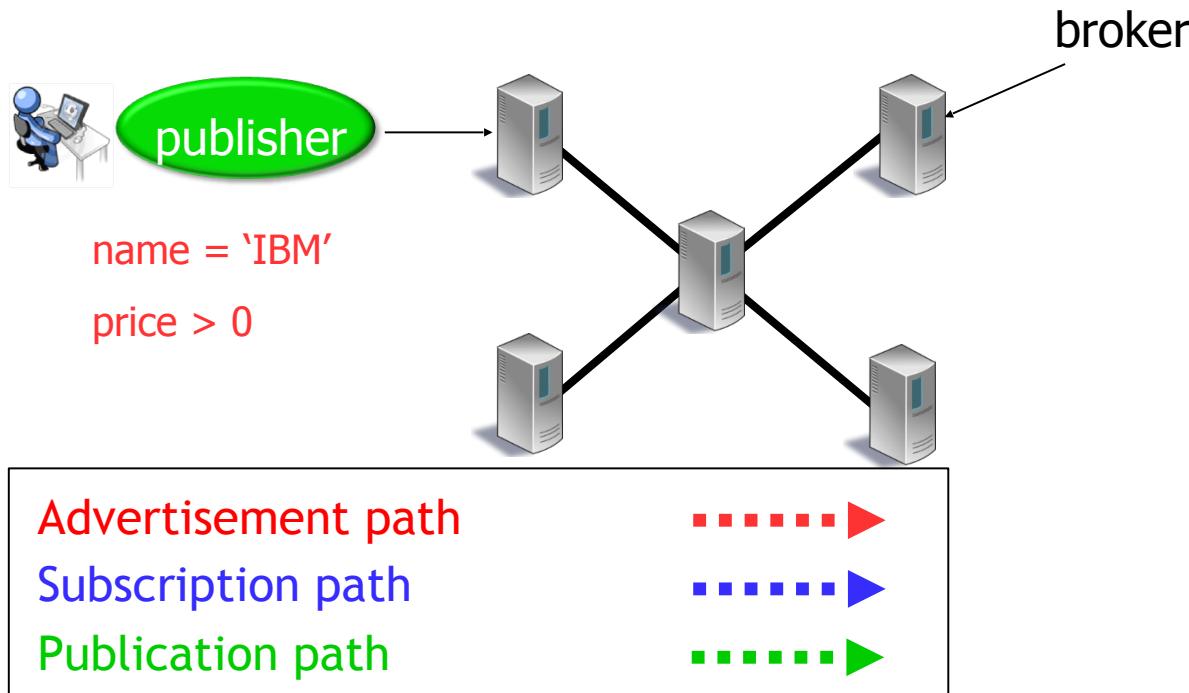
Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



Advertisements are flooded, but not subscriptions

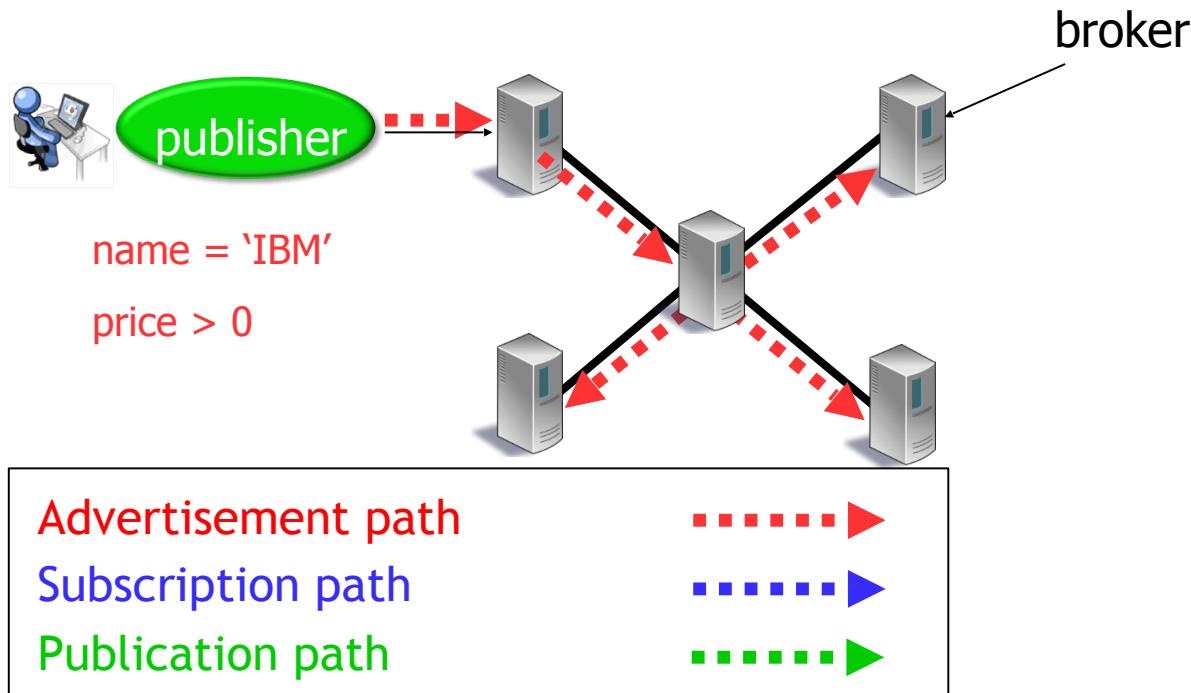
# Advertisement-Based Routing

Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



# Advertisement-Based Routing

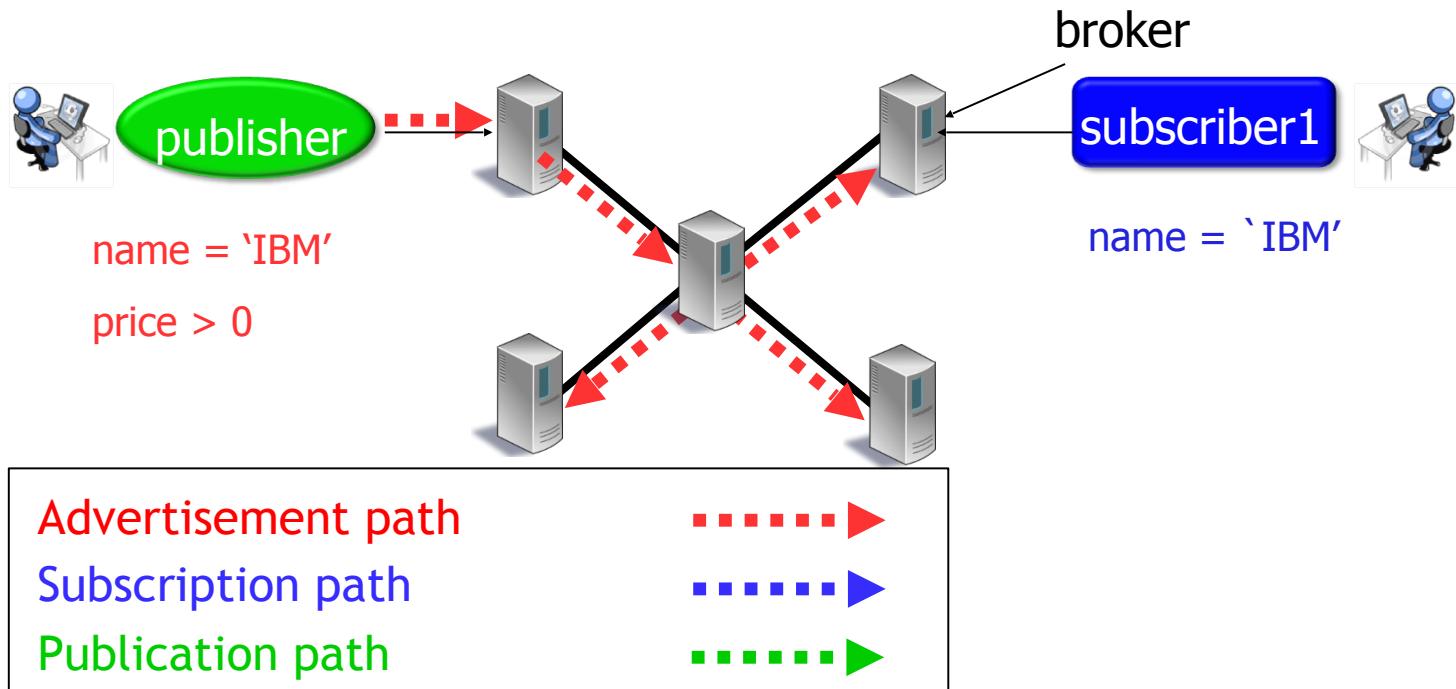
Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



Advertisements are flooded, but not subscriptions

# Advertisement-Based Routing

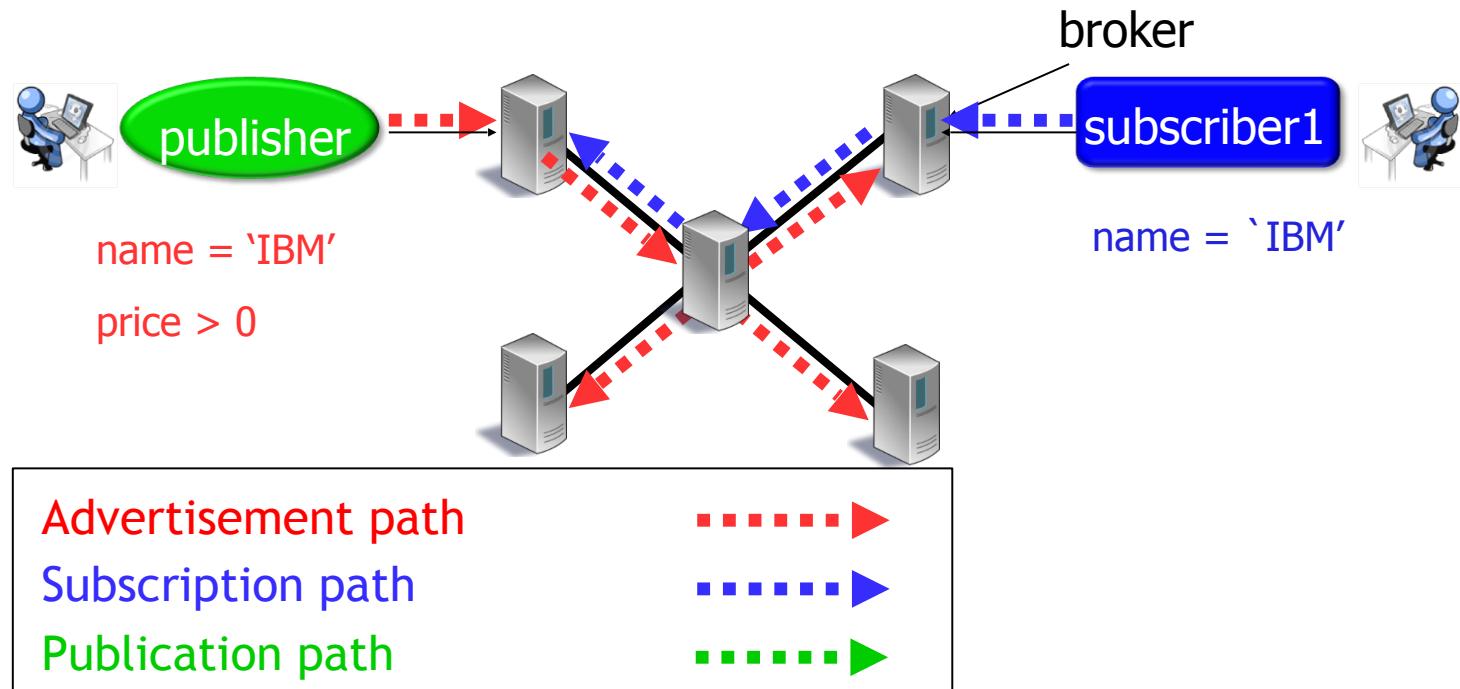
Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



Advertisements are flooded, but not subscriptions

# Advertisement-Based Routing

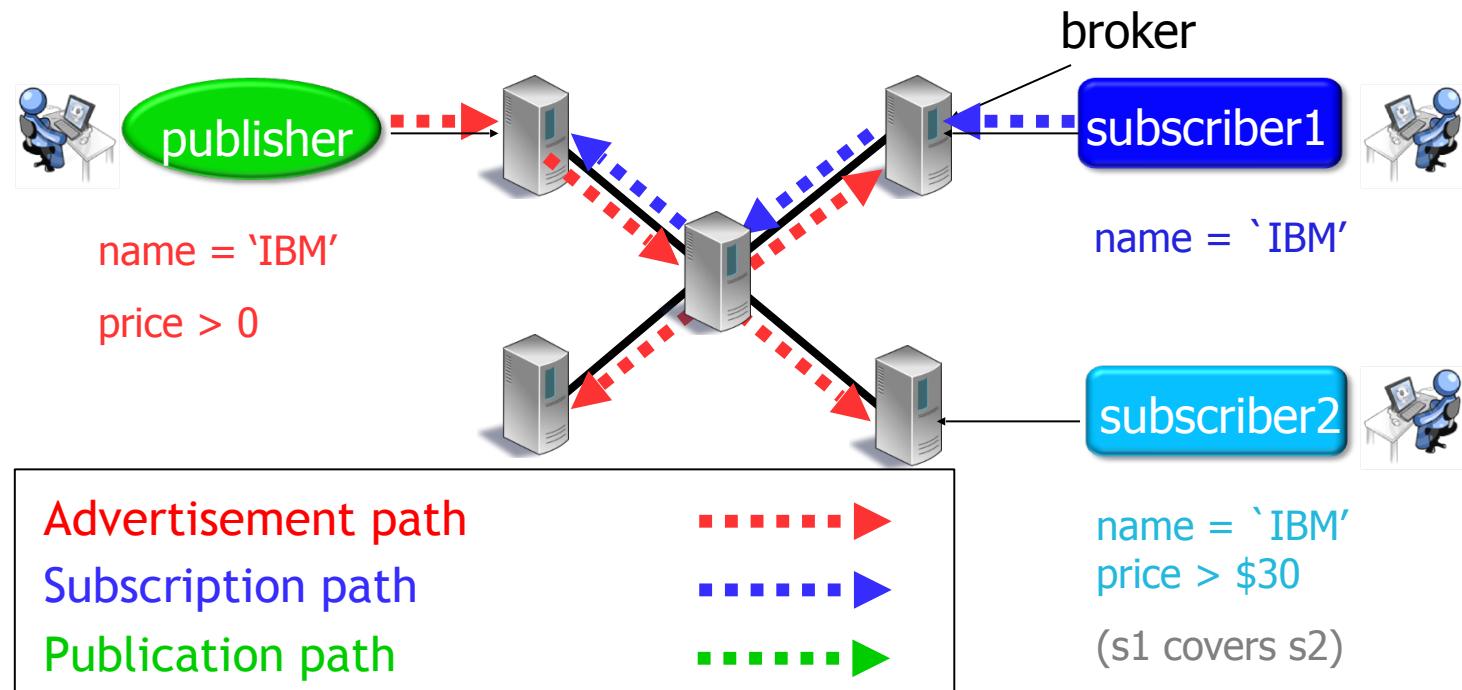
Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



Advertisements are flooded, but not subscriptions

# Advertisement-Based Routing

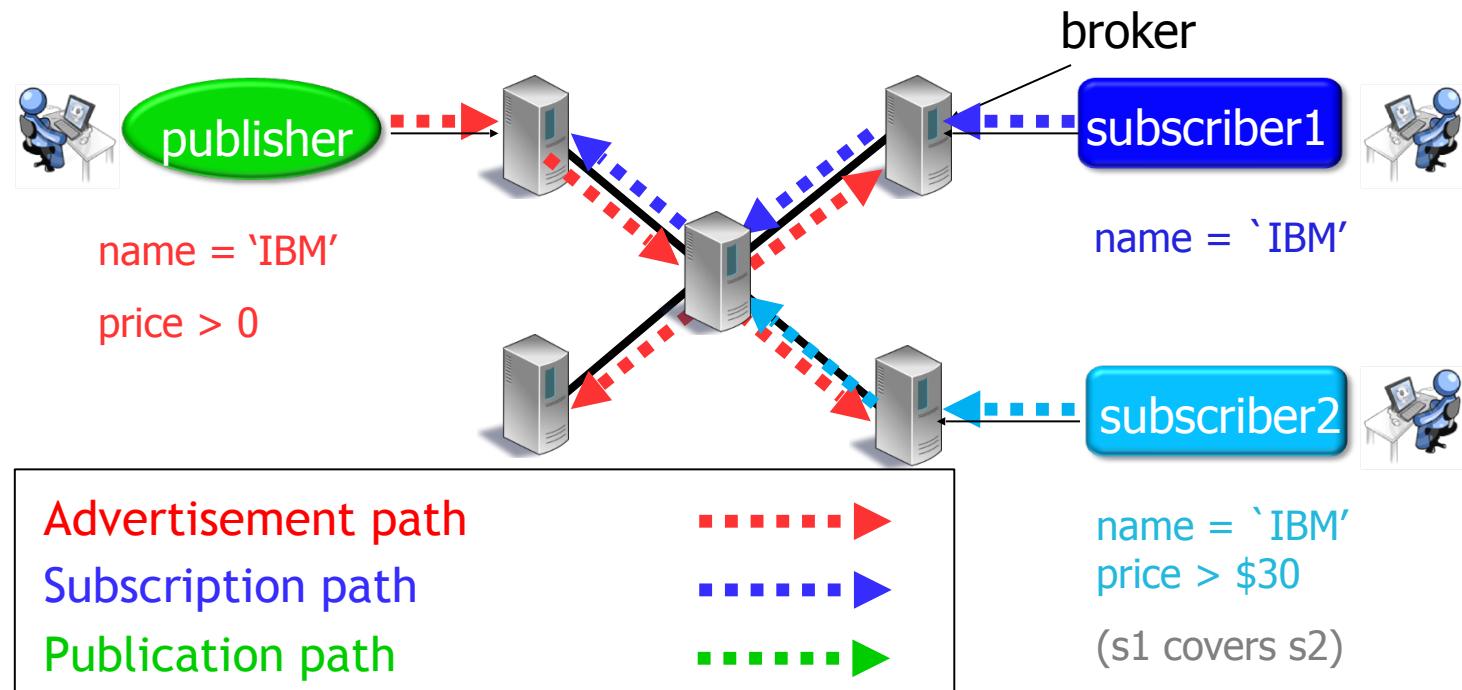
Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



Advertisements are flooded, but not subscriptions

# Advertisement-Based Routing

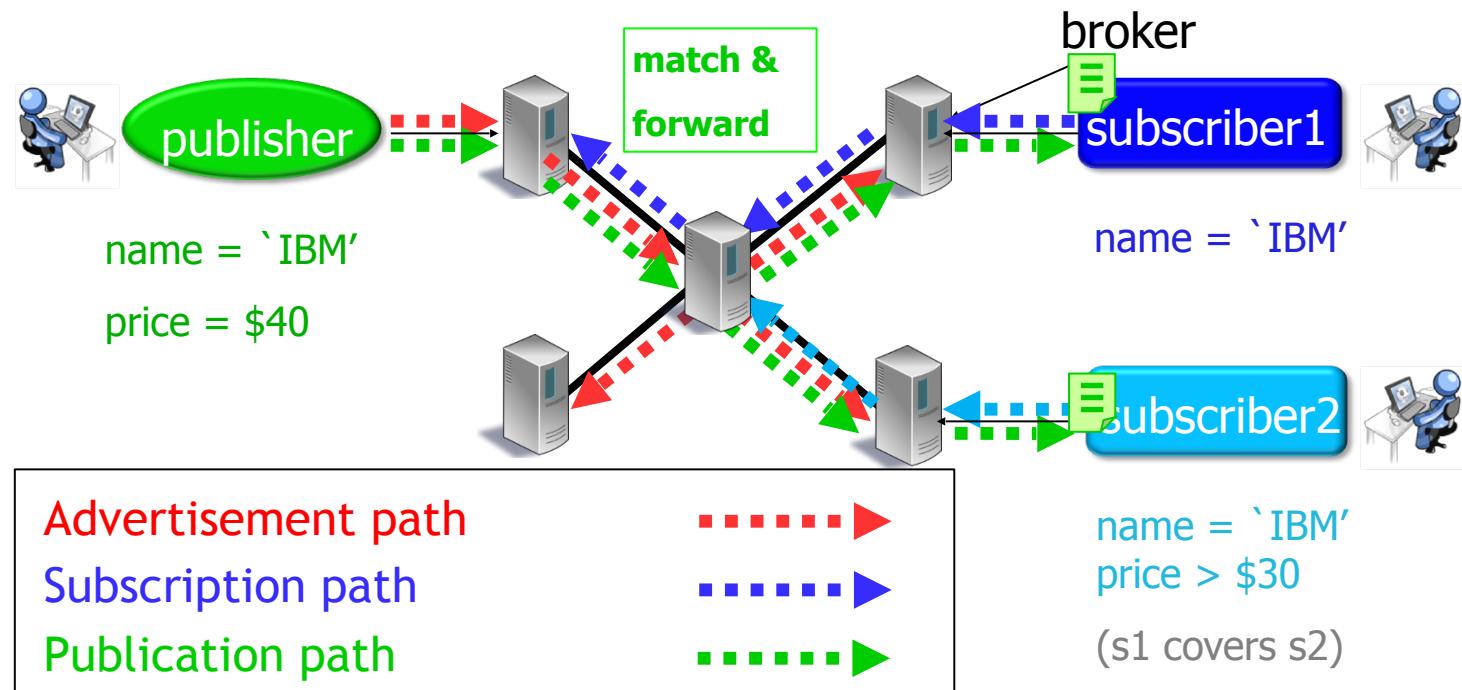
Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



Advertisements are flooded, but not subscriptions

# Advertisement-Based Routing

Publishers must first *advertise* the type of publications they are going to make. Advertisements are stored at each broker in a Subscription Routing Table (SRT).



Advertisements are flooded, but not subscriptions

# Routing Algorithms Summary

- Rendezvous-based
  - Partition-based: Simple, requires master
  - DHT-based:  $\log(n)$  hops, structured network
- Overlay-based
  - Rendezvous-based: Efficient matching, inefficient communication, RPs are bottlenecks
  - Filtering-based: Computation intensive, efficient communication
  - Advertisement flooding: Useful if the number of publishers is much smaller than subscribers

